

# ClassificationChurnLatestFeb26

March 3, 2025

## 0.1 CLASSIFICATION ANALYSIS

DATASET: CUSTOMER CHURN FROM Kaggle.com

```
[3]: # Necessary imports for preprocessing and machine learning

import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, MinMaxScaler
# if you want to use one Hot encoder then you need to import the below one
from sklearn.preprocessing import OneHotEncoder
from sklearn.impute import SimpleImputer    # interpolation for missing values
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import RocCurveDisplay, confusion_matrix, accuracy_score,
                           precision_score, recall_score, f1_score, classification_report
from imblearn.over_sampling import SMOTE
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[4]: # load the dataset

df=pd.read_csv(r"Customer_Churn_Dataset.csv")
```

```
[5]: df.head(5)
```

```
[5]:   CustomerID  Age  Subscription_Length_Months  Monthly_Bill  Total_Usage_GB \
0           1  56.0                               38      112.48       61.87
1           2  69.0                               28      174.51       310.65
2           3  46.0                               48      35.11       310.02
3           4  32.0                               38      123.14       434.93
4           5  60.0                               31      95.38       286.54

  Support_Calls  Contract_Type Payment_Method  Has_Additional_Services \
0              8        Two Year        Debit Card                  1
```

```

1           6      Two Year      Debit Card      1
2           1  Month-to-Month      UPI          1
3           7      One Year      PayPal         0
4           1  Month-to-Month      Credit Card     0

Churn
0      0
1      0
2      0
3      0
4      1

```

[6]: df.dtypes

```

[6]: CustomerID          int64
      Age              float64
      Subscription_Length_Months  int64
      Monthly_Bill        float64
      Total_Usage_GB      float64
      Support_Calls       int64
      Contract_Type       object
      Payment_Method      object
      Has_Additional_Services  int64
      Churn             int64
      dtype: object

```

[ ]:

Descriptive Statistics of Data

[8]: df.describe().T

	count	mean	std	min	\
CustomerID	10000.0	5000.500000	2886.895680	1.00	
Age	9998.0	46.181136	16.551075	2.00	
Subscription_Length_Months	10000.0	29.901700	16.996680	1.00	
Monthly_Bill	10000.0	104.676996	54.670610	10.00	
Total_Usage_GB	10000.0	252.664805	145.052928	0.51	
Support_Calls	10000.0	4.538300	2.876658	0.00	
Has_Additional_Services	10000.0	0.506500	0.499983	0.00	
Churn	10000.0	0.210700	0.407826	0.00	
	25%	50%	75%	max	
CustomerID	2500.750	5000.50	7500.2500	10000.00	
Age	32.000	46.00	61.0000	250.00	
Subscription_Length_Months	15.000	29.00	45.0000	59.00	
Monthly_Bill	57.575	104.54	151.2475	199.98	
Total_Usage_GB	129.045	253.18	377.5475	2000.11	

```
Support_Calls           2.000      5.00    7.0000    9.00
Has_Additional_Services 0.000      1.00    1.0000    1.00
Churn                   0.000      0.00    0.0000    1.00
```

```
[9]: df.describe(include=['object']).T
```

```
count unique          top freq
Contract_Type    9998     3 Month-to-Month 3352
Payment_Method   10000    4 Debit Card    2540
```

```
[ ]:
```

Checking for Null Values

```
[11]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 10 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   CustomerID        10000 non-null   int64  
 1   Age               9998 non-null    float64 
 2   Subscription_Length_Months 10000 non-null   int64  
 3   Monthly_Bill       10000 non-null    float64 
 4   Total_Usage_GB     10000 non-null    float64 
 5   Support_Calls      10000 non-null    int64  
 6   Contract_Type      9998 non-null    object  
 7   Payment_Method      10000 non-null    object  
 8   Has_Additional_Services 10000 non-null   int64  
 9   Churn              10000 non-null    int64  
dtypes: float64(3), int64(5), object(2)
memory usage: 781.4+ KB
```

```
[12]: # find the total number of missing values for each variable
```

```
df.isnull().sum()
```

```
CustomerID          0
Age                 2
Subscription_Length_Months 0
Monthly_Bill         0
Total_Usage_GB       0
Support_Calls        0
Contract_Type        2
Payment_Method        0
Has_Additional_Services 0
Churn                0
```

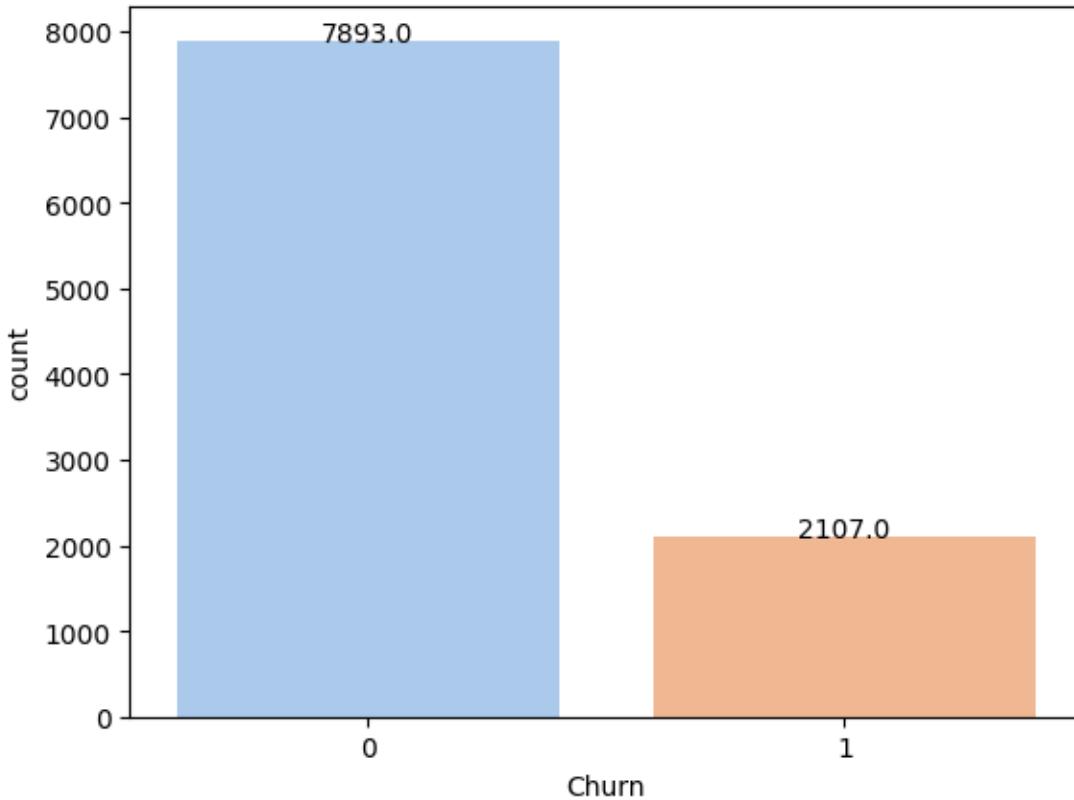
```
dtype: int64
```

```
[ ]:
```

Is the dataset balanced?

```
[14]: # Create the countplot
#ax = sns.countplot(data=df, x='Churn')
ax = sns.countplot(data=df, x='Churn', hue='Churn', palette='pastel', legend=False)
# Add count labels on top of the bars
for p in ax.patches:
    ax.annotate(f'{p.get_height()}', (p.get_x() + p.get_width() / 2., p.get_height()),
                ha='center', va='baseline')

# Display the plot
plt.show()
```



The dataset is not balanced. We balance it using SMOTE

```
[16]: # count the number of people who have customer churn yes against number of people who have customer churn no to see if the label/response is balanced or not
#customer churn yes = 1

# Number of people with customer churn yes
(df[['Churn']]==1).sum()      # you can also use      #df.Churn.
                               ↪value_counts()[1]
```

```
[16]: Churn    2107
      dtype: int64
```

```
[17]: # Number of people with customer churn no
#Customer churn no = 0
(df[['Churn']]==0).sum()  # you can also use  # df.Churn.value_counts()[0]
```

```
[17]: Churn    7893
      dtype: int64
```

```
[18]: # calculate probability of customer churn yes
      (df[['Churn']]==1).sum() / len(df['Churn'])
```

```
[18]: Churn    0.2107
      dtype: float64
```

```
[19]: # calculate probability of customer churn no
      (df[['Churn']]==0).sum() / len(df['Churn'])
```

```
[19]: Churn    0.7893
      dtype: float64
```

```
[ ]:
```

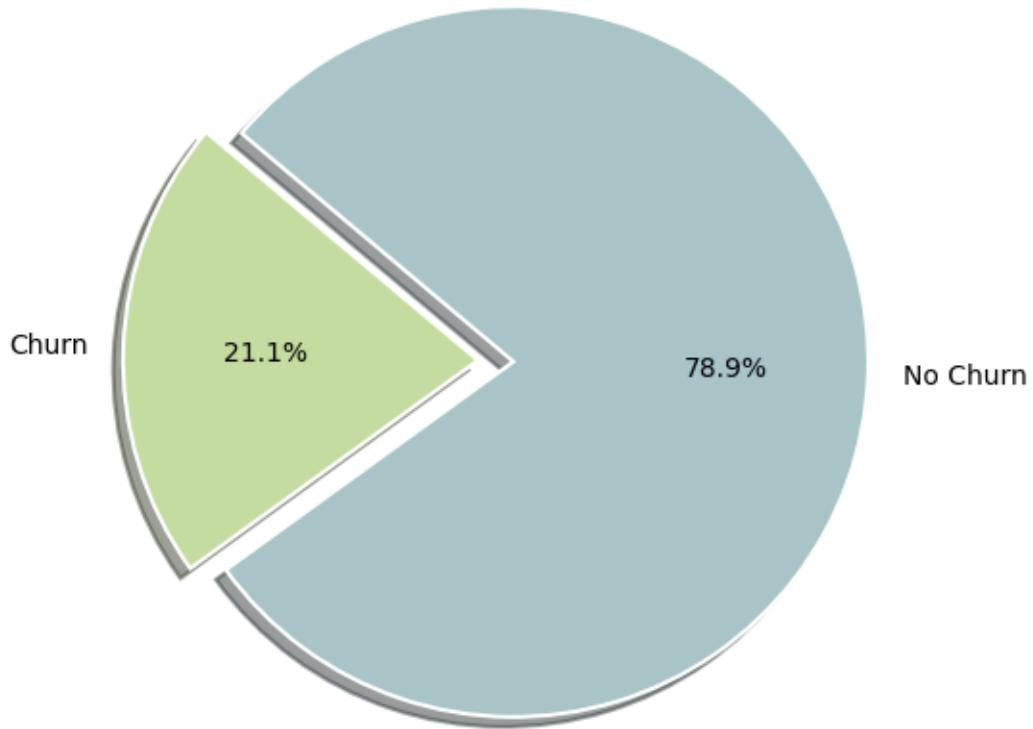
Pie Chart to show data imbalance

```
[21]: l = ['Churn', 'No Churn']
s = [ (df['Churn']==1).sum(), (df['Churn']==0).sum()]
c = ['#C5DCA0', '#A8C4C7']
```

```
[22]: # Plot pie chart
plt.figure(figsize=(6,6))
plt.pie(s, labels=l, colors=c, explode=(0.1, 0), shadow=True, autopct='%.1f%%',
         startangle=140, wedgeprops={'edgecolor': 'white', 'linewidth': 1.5})

plt.title("Class Distribution Before Balancing")
plt.show()
```

Class Distribution Before Balancing



[ ]:

Handling Missing values

[24]: # How to get the list of all columns in my dataset

```
columns_list=df.columns.tolist()
```

[25]: columns\_list

```
[25]: ['CustomerID',
       'Age',
       'Subscription_Length_Months',
       'Monthly_Bill',
       'Total_Usage_GB',
       'Support_Calls',
       'Contract_Type',
       'Payment_Method',
```

```

'Has_Additional_Services',
'Churn']

[26]: # We need to create a variable that stores the name of categorical columns
# as well a varibale that stores the name of numerical variable.
# because we need to handle missing values for categorical using mode technique
# we also need to replace the missing values of numerical variable using mean. ↴
    ↪Can also use linear interpocation

# List of categorical variables (dtype == object or category)
categorical_cols = df.select_dtypes(include=['object', 'category']).columns.
    ↪tolist()

# List of numerical variables (dtype == int or float)
numerical_cols = df.select_dtypes(include=['number']).columns.tolist()

# Display the lists
print("Categorical columns:", categorical_cols)
print("Numerical columns:", numerical_cols)

```

Categorical columns: ['Contract\_Type', 'Payment\_Method']  
 Numerical columns: ['CustomerID', 'Age', 'Subscription\_Length\_Months',  
 'Monthly\_Bill', 'Total\_Usage\_GB', 'Support\_Calls', 'Has\_Additional\_Services',  
 'Churn']

```

[27]: # Input categorical variable with mode (replace missing values of a categorical ↴
    ↪variable with mode)

for col in categorical_cols:
    df[col] = df[col].fillna(df[col].mode()[0])
    #for col in categorical_cols:
    #    df[col].fillna(df[col].mode()[0], inplace=True)

# note that we may have multiple values for mode. mode()[0] returns the first ↴
    ↪most frequent category.
# mode()[1] returns the second mode. This is because the mode() function can ↴
    ↪potentially return multiple values if there
# is more than one mode (i.e., multiple values that occur with the same highest ↴
    ↪frequency).
# If you don't use [0] and the mode() returns more than one value, the .
    ↪fillna() function will not know which value to use
# If you remove it, you might get an error if there is more than one mode,
# because mode() can return multiple values and .fillna() requires only one ↴
    ↪value.

# If you remove inplace=True, the original DataFrame remains unchanged unless ↴
    ↪you explicitly reassign the result.

```

```
[28]: df.isnull().sum() # all missing values under categorical variables have been  
↪replaced with their mode
```

```
[28]: CustomerID          0  
Age                  2  
Subscription_Length_Months 0  
Monthly_Bill           0  
Total_Usage_GB         0  
Support_Calls          0  
Contract_Type          0  
Payment_Method         0  
Has_Additional_Services 0  
Churn                 0  
dtype: int64
```

```
[29]: # handling missing values of numerical variables  
# Input numerical column (variable), which is Age, with mean  
df['Age'] = df['Age'].fillna(df['Age'].mean())  
#df['age'].fillna(df['age'].mean(), inplace=True)
```

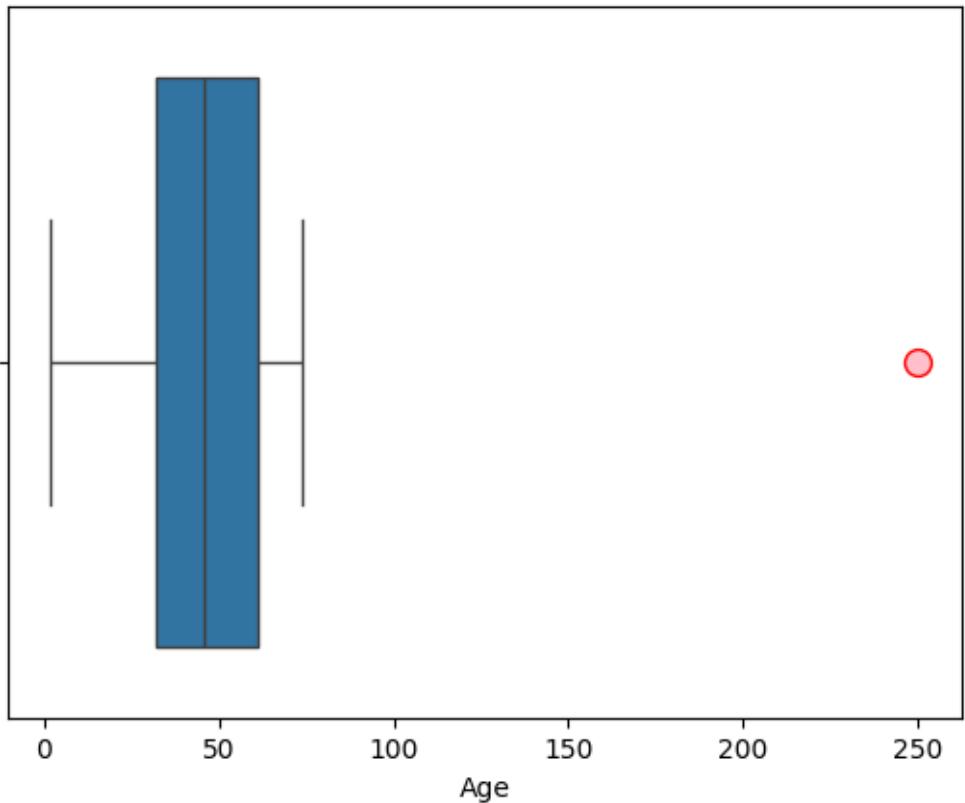
```
[30]: df.isnull().sum() # zero missing values for numerical variable (age)
```

```
[30]: CustomerID          0  
Age                  0  
Subscription_Length_Months 0  
Monthly_Bill           0  
Total_Usage_GB         0  
Support_Calls          0  
Contract_Type          0  
Payment_Method         0  
Has_Additional_Services 0  
Churn                 0  
dtype: int64
```

```
[ ]:
```

Find and remove outliers from numerical variable

```
[32]: # Create a boxplot for the 'age' column to detect outliers  
sns.boxplot(data=df, x='Age');  
# Customizing outlier appearance  
sns.boxplot(data=df, x='Age',  
            flierprops={"marker": "o", "markerfacecolor": "pink",  
↪"markeredgecolor": "red", "markersize": 10});
```



```
[33]: # Calculate the IQR for the 'age' column
Q1 = df['Age'].quantile(0.25) # 25th percentile
Q3 = df['Age'].quantile(0.75) # 75th percentile
IQR = Q3 - Q1 # Interquartile range

# Define the lower and upper bounds for outliers
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

# Filter out the outliers
df1 = df[(df['Age'] >= lower_bound) & (df['Age'] <= upper_bound)]

# df1 is a new datafram after removing outliers

# Display the results (rows without outliers)
print("Cleaned DataFrame (Without Age Outliers):")
print(df1)
```

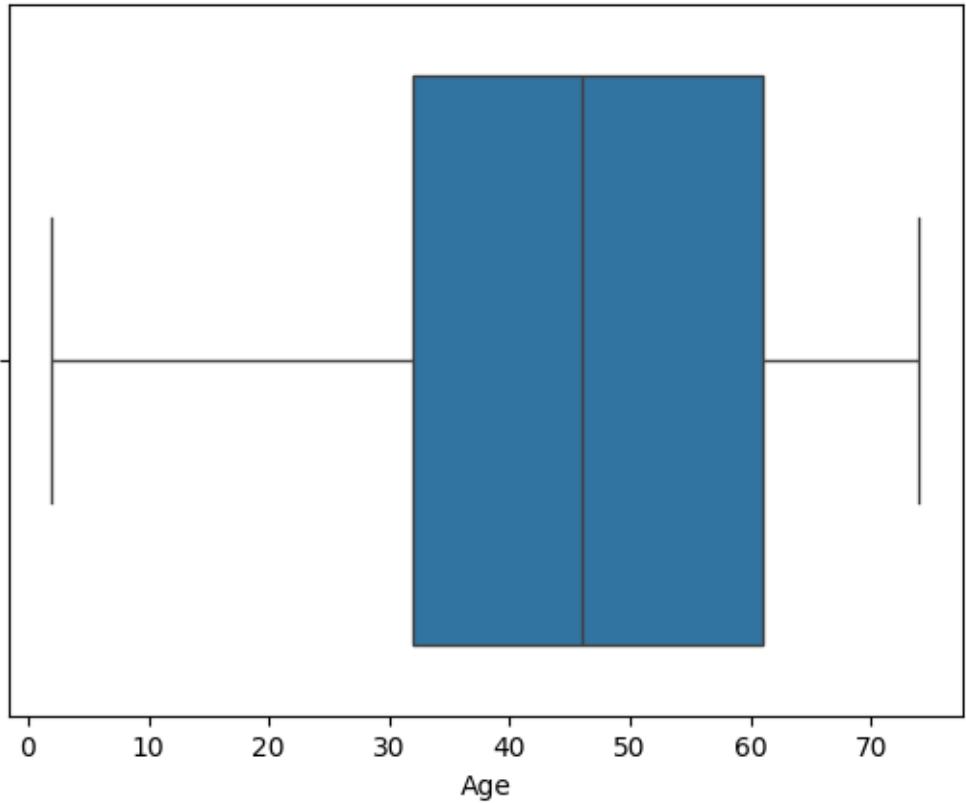
Cleaned DataFrame (Without Age Outliers):

CustomerID	Age	Subscription_Length_Months	Monthly_Bill	
0	1	56.0	38	112.48

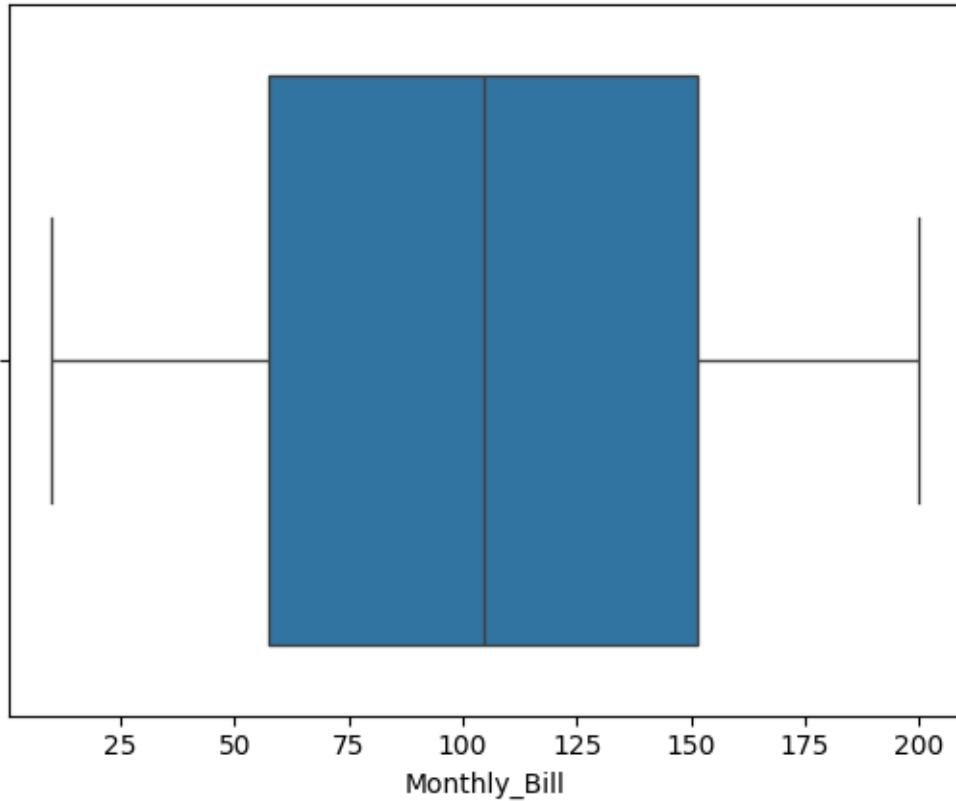
1	2	69.0	28	174.51
2	3	46.0	48	35.11
3	4	32.0	38	123.14
4	5	60.0	31	95.38
...	...	...	...	...
9995	9996	40.0	8	173.06
9996	9997	58.0	33	175.46
9997	9998	57.0	18	13.80
9998	9999	46.0	45	20.59
9999	10000	33.0	29	184.48
0	Total_Usage_GB	Support_Calls	Contract_Type	Payment_Method \
1	61.87	8	Two Year	Debit Card
2	310.65	6	Two Year	Debit Card
3	310.02	1	Month-to-Month	UPI
4	434.93	7	One Year	PayPal
...	...	...	...	...
9995	127.42	2	Two Year	Credit Card
9996	385.53	1	Two Year	Debit Card
9997	115.07	3	Month-to-Month	UPI
9998	260.86	7	One Year	PayPal
9999	55.79	4	Month-to-Month	UPI
0	Has_Additional_Services	Churn		
1	1	0		
2	1	0		
3	0	0		
4	0	1		
...	...	...		
9995	1	0		
9996	0	0		
9997	1	0		
9998	1	0		
9999	1	0		

[9999 rows x 10 columns]

```
[34]: sns.boxplot(data=df1, x='Age'); # outliers have been removed
```



```
[35]: # Find and remove outlier from numerical variable  
  
# Create a boxplot for the 'Monthly_Bill' column to detect outliers  
sns.boxplot(data=df, x='Monthly_Bill');
```



```
[36]: # Calculate the IQR for the 'age' column
Q1 = df1['Monthly_Bill'].quantile(0.25) # 25th percentile
Q3 = df1['Monthly_Bill'].quantile(0.75) # 75th percentile
IQR = Q3 - Q1 # Interquartile range

# Define the lower and upper bounds for outliers
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

# Filter out the outliers
df2 = df1[(df1['Monthly_Bill'] >= lower_bound) & (df1['Monthly_Bill'] <= upper_bound)]

# df1 is a new dataframe after removing Age outliers and df2 is the new
# dataframe after removing Age and MonthlyBill outliers

# Display the results (rows without outliers)
print("Cleaned DataFrame (Without Monthly_Bill Outliers):")
print(df2)
```

Cleaned DataFrame (Without Monthly\_Bill Outliers):

```

CustomerID    Age   Subscription_Length_Months  Monthly_Bill \
0             1     56.0                           38      112.48
1             2     69.0                           28      174.51
2             3     46.0                           48      35.11
3             4     32.0                           38      123.14
4             5     60.0                           31      95.38
...
9995          ...   ...                           ...      ...
9996          9996  40.0                           8       173.06
9996          9997  58.0                           33      175.46
9997          9998  57.0                           18      13.80
9998          9999  46.0                           45      20.59
9999         10000 33.0                           29      184.48

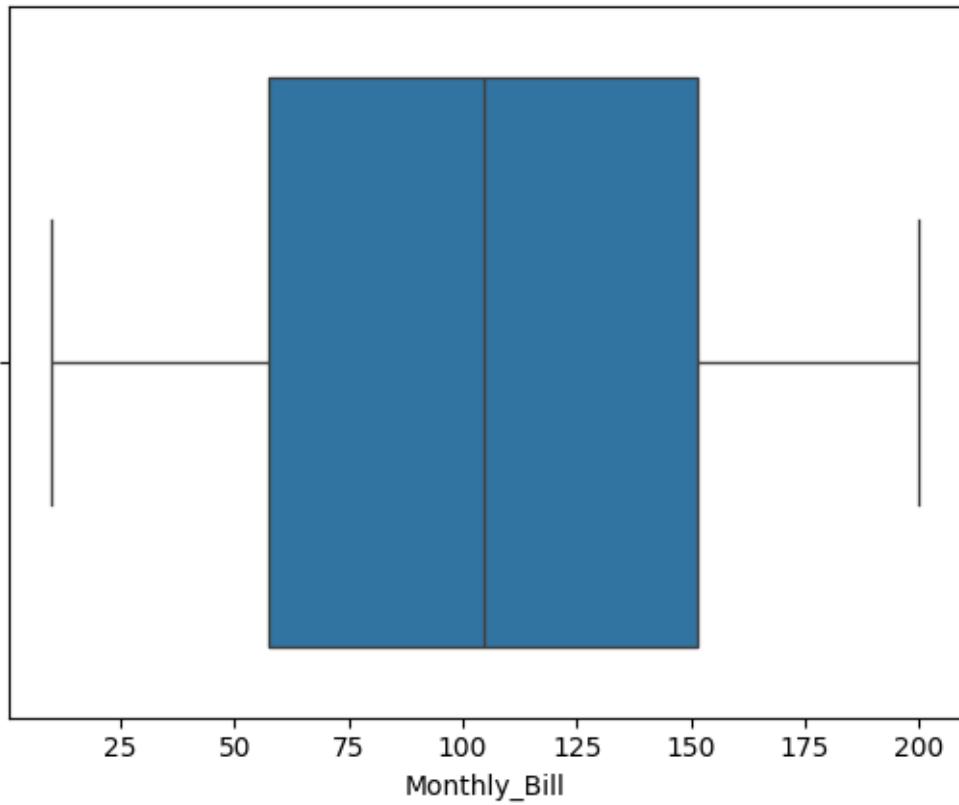
Total_Usage_GB  Support_Calls  Contract_Type Payment_Method \
0           61.87        8     Two Year     Debit Card
1           310.65       6     Two Year     Debit Card
2           310.02       1 Month-to-Month UPI
3           434.93       7     One Year     PayPal
4           286.54       1 Month-to-Month Credit Card
...
9995         127.42       2     Two Year     Credit Card
9996         385.53       1     Two Year     Debit Card
9997         115.07       3 Month-to-Month UPI
9998         260.86       7     One Year     PayPal
9999         55.79       4 Month-to-Month UPI

Has_Additional_Services  Churn
0                      1     0
1                      1     0
2                      1     0
3                      0     0
4                      0     1
...
9995          ...   ...
9996          ...   ...
9997          ...   ...
9998          ...   ...
9999          ...   ...

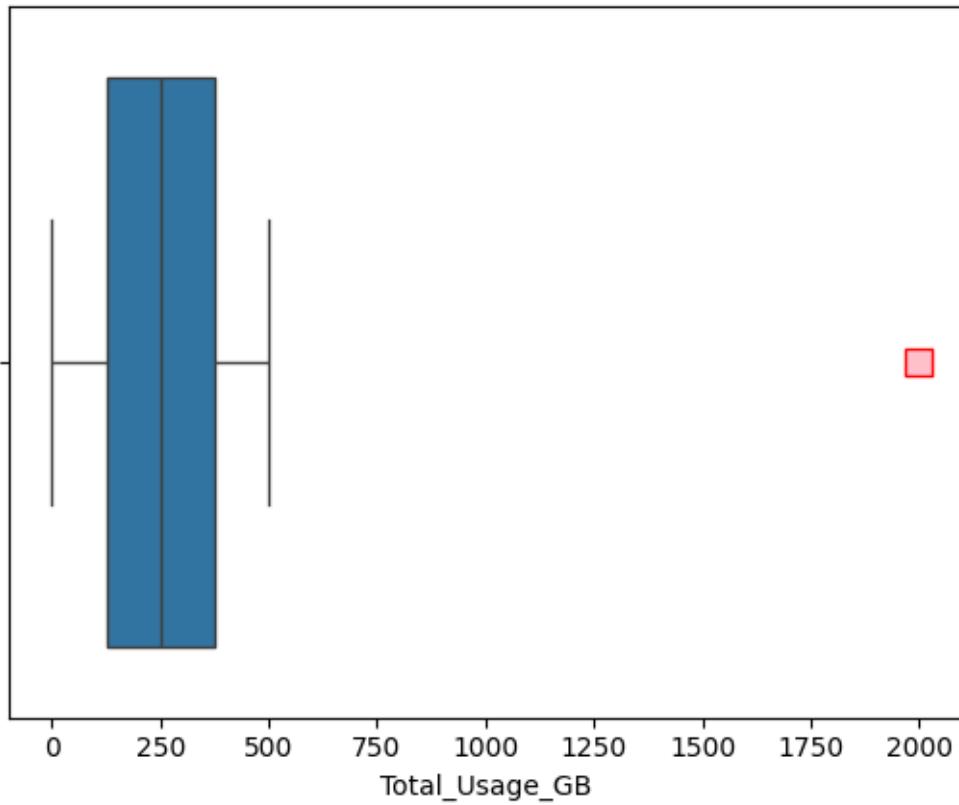
```

[9999 rows x 10 columns]

```
[37]: sns.boxplot(data=df2, x='Monthly_Bill'); # outliers have been removed
```



```
[38]: sns.boxplot(data=df2, x='Total_Usage_GB'); # with outliers
# Customizing outlier appearance
sns.boxplot(data=df2, x='Total_Usage_GB',
             flierprops={"marker": "s", "markerfacecolor": "pink", ↴
             "markeredgecolor": "red", "markersize": 10});
```



```
[39]: # Calculate the IQR for the 'age' column
Q1 = df2['Total_Usage_GB'].quantile(0.25) # 25th percentile
Q3 = df2['Total_Usage_GB'].quantile(0.75) # 75th percentile
IQR = Q3 - Q1 # Interquartile range

# Define the lower and upper bounds for outliers
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

# Filter out the outliers
df3 = df2[(df2['Total_Usage_GB'] >= lower_bound) & (df2['Total_Usage_GB'] <= upper_bound)]

# df3 is a new dataframe after removing Total Usage outliers.

# Display the results (rows without outliers)
print("Cleaned DataFrame (Without Total_Usage_GB Outliers):")
print(df3)
```

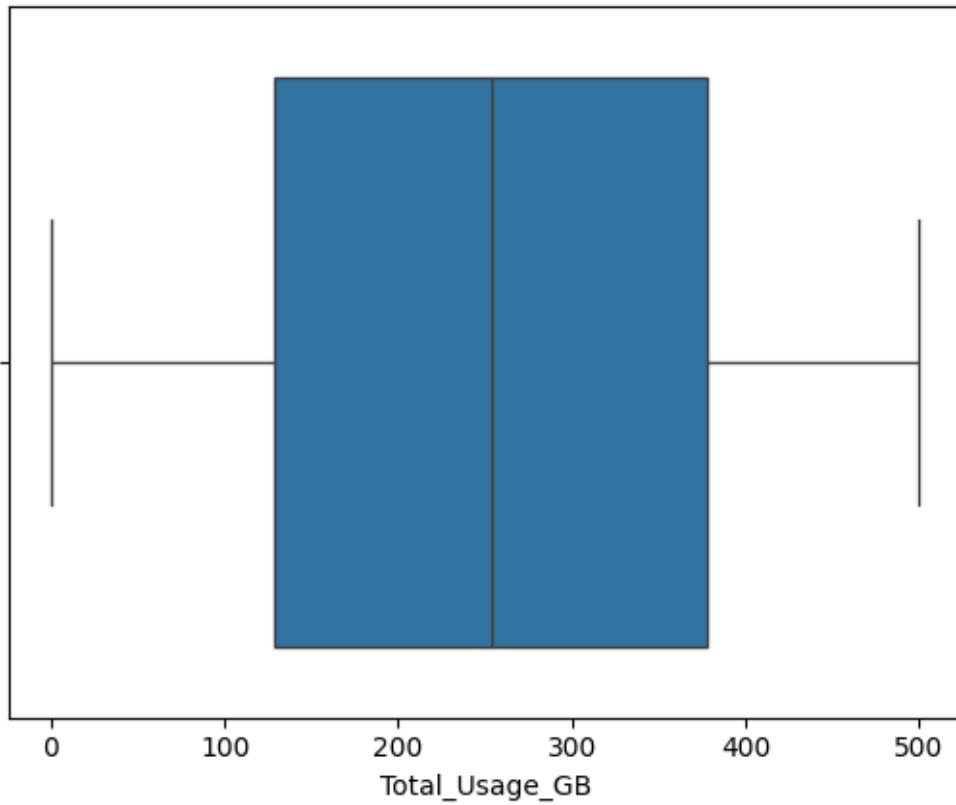
Cleaned DataFrame (Without Total\_Usage\_GB Outliers):

CustomerID	Age	Subscription_Length_Months	Monthly_Bill
1	22	12	30
2	22	12	30
3	22	12	30
4	22	12	30
5	22	12	30
6	22	12	30
7	22	12	30
8	22	12	30
9	22	12	30
10	22	12	30
11	22	12	30
12	22	12	30
13	22	12	30
14	22	12	30
15	22	12	30
16	22	12	30
17	22	12	30
18	22	12	30
19	22	12	30
20	22	12	30
21	22	12	30
22	22	12	30
23	22	12	30
24	22	12	30
25	22	12	30
26	22	12	30
27	22	12	30
28	22	12	30
29	22	12	30
30	22	12	30
31	22	12	30
32	22	12	30
33	22	12	30
34	22	12	30
35	22	12	30
36	22	12	30
37	22	12	30
38	22	12	30
39	22	12	30
40	22	12	30
41	22	12	30
42	22	12	30
43	22	12	30
44	22	12	30
45	22	12	30
46	22	12	30
47	22	12	30
48	22	12	30
49	22	12	30
50	22	12	30
51	22	12	30
52	22	12	30
53	22	12	30
54	22	12	30
55	22	12	30
56	22	12	30
57	22	12	30
58	22	12	30
59	22	12	30
60	22	12	30
61	22	12	30
62	22	12	30
63	22	12	30
64	22	12	30
65	22	12	30
66	22	12	30
67	22	12	30
68	22	12	30
69	22	12	30
70	22	12	30
71	22	12	30
72	22	12	30
73	22	12	30
74	22	12	30
75	22	12	30
76	22	12	30
77	22	12	30
78	22	12	30
79	22	12	30
80	22	12	30
81	22	12	30
82	22	12	30
83	22	12	30
84	22	12	30
85	22	12	30
86	22	12	30
87	22	12	30
88	22	12	30
89	22	12	30
90	22	12	30
91	22	12	30
92	22	12	30
93	22	12	30
94	22	12	30
95	22	12	30
96	22	12	30
97	22	12	30
98	22	12	30
99	22	12	30
100	22	12	30

0		1	56.0		38	112.48
1		2	69.0		28	174.51
2		3	46.0		48	35.11
3		4	32.0		38	123.14
4		5	60.0		31	95.38
...	...	...	...	...	...	...
9995		9996	40.0		8	173.06
9996		9997	58.0		33	175.46
9997		9998	57.0		18	13.80
9998		9999	46.0		45	20.59
9999		10000	33.0		29	184.48
	Total_Usage_GB	Support_Calls	Contract_Type	Payment_Method		\
0	61.87	8	Two Year	Debit Card		
1	310.65	6	Two Year	Debit Card		
2	310.02	1	Month-to-Month	UPI		
3	434.93	7	One Year	PayPal		
4	286.54	1	Month-to-Month	Credit Card		
...	...	...	...	...	...	...
9995	127.42	2	Two Year	Credit Card		
9996	385.53	1	Two Year	Debit Card		
9997	115.07	3	Month-to-Month	UPI		
9998	260.86	7	One Year	PayPal		
9999	55.79	4	Month-to-Month	UPI		
	Has_Additional_Services	Churn				
0		1	0			
1		1	0			
2		1	0			
3		0	0			
4		0	1			
...	...	...	...	...	...	...
9995		1	0			
9996		0	0			
9997		1	0			
9998		1	0			
9999		1	0			

[9998 rows x 10 columns]

[40]: `sns.boxplot(data=df3, x='Total_Usage_GB'); # outliers have been removed`



```
[41]: type(df3)
```

```
[41]: pandas.core.frame.DataFrame
```

```
[ ]:
```

Label Encoding for Categorical Features

```
[43]: # in this dataset, we have 2 categorical features (e.g., contract_type(monthly, year-to-year, ...), and a response/label variable)

# we use the following function to encode categorical variable

from sklearn.preprocessing import LabelEncoder

lb=LabelEncoder()

for col in categorical_cols:    # we define categorical_cols in line [73]
    df3.loc[:,col]=lb.fit_transform(df3[col])  # note that df3 is new dataframe without outliers
```

```
[44]: df3.head(5)
```

```
[44]:   CustomerID  Age  Subscription_Length_Months  Monthly_Bill  Total_Usage_GB \
0           1  56.0                           38      112.48       61.87
1           2  69.0                           28      174.51       310.65
2           3  46.0                           48      35.11       310.02
3           4  32.0                           38     123.14       434.93
4           5  60.0                           31      95.38       286.54

   Support_Calls  Contract_Type  Payment_Method  Has_Additional_Services  Churn
0             8            2                 1                  1         0
1             6            2                 1                  1         0
2             1            0                 3                  1         0
3             7            1                 2                  0         0
4             1            0                 0                  0         1
```

```
[45]: numerical_cols
```

```
[45]: ['CustomerID',
'Age',
'Subscription_Length_Months',
'Monthly_Bill',
'Total_Usage_GB',
'Support_Calls',
'Has_Additional_Services',
'Churn']
```

```
[46]: df3.columns
```

```
[46]: Index(['CustomerID', 'Age', 'Subscription_Length_Months', 'Monthly_Bill',
       'Total_Usage_GB', 'Support_Calls', 'Contract_Type', 'Payment_Method',
       'Has_Additional_Services', 'Churn'],
       dtype='object')
```

```
[47]: # Standardize the numerical variable (Age variable)
```

```
from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler()

df3 = df3.copy();
scale_cols = [x for x in numerical_cols if x != 'CustomerID']
df3.loc[:,scale_cols] = scaler.fit_transform(df3[scale_cols]); # we define ↴
                   # list of numerical variable in line [22]

# df3 is the new dataframe after standardizing the numerical variable
```

```
/var/folders/9p/02j0tmhs5w35zmgt4_pc9src0000gn/T/ipykernel_4362/3217076613.py:9:
```

```
FutureWarning: Setting an item of incompatible dtype is deprecated and will
raise in a future error of pandas. Value '[0.63793103 0.46551724 0.81034483 ...
0.29310345 0.75862069 0.48275862]' has dtype incompatible with int64, please
explicitly cast to a compatible dtype first.
```

```
df3.loc[:,scale_cols] = scaler.fit_transform(df3[scale_cols]); # we define
list of numerical variable in line [22]
/var/folders/9p/02j0tmhs5w35zmgt4_pc9src0000gn/T/ipykernel_4362/3217076613.py:9:
FutureWarning: Setting an item of incompatible dtype is deprecated and will
raise in a future error of pandas. Value '[0.88888889 0.66666667 0.11111111 ...
0.33333333 0.77777778 0.44444444]' has dtype incompatible with int64, please
explicitly cast to a compatible dtype first.
df3.loc[:,scale_cols] = scaler.fit_transform(df3[scale_cols]); # we define
list of numerical variable in line [22]
```

```
[48]: df3
```

```
[48]:      CustomerID      Age  Subscription_Length_Months  Monthly_Bill \
0              1  0.750000                  0.637931    0.539425
1              2  0.930556                  0.465517    0.865933
2              3  0.611111                  0.810345    0.132172
3              4  0.416667                  0.637931    0.595536
4              5  0.805556                  0.517241    0.449416
...
9995         ...        ...                  ...        ...
9996         9996  0.527778                  0.120690    0.858301
9996         9997  0.777778                  0.551724    0.870934
9997         9998  0.763889                  0.293103    0.020002
9998         9999  0.611111                  0.758621    0.055743
9999        10000  0.430556                  0.482759    0.918412

      Total_Usage_GB  Support_Calls Contract_Type Payment_Method \
0          0.122848      0.888889          2            1
1          0.620926      0.666667          2            1
2          0.619664      0.111111          0            3
3          0.869745      0.777778          1            2
4          0.572656      0.111111          0            0
...
9995        ...        ...                  ...        ...
9996        0.254084      0.222222          2            0
9996        0.770842      0.111111          2            1
9997        0.229359      0.333333          0            3
9998        0.521242      0.777778          1            2
9999        0.110675      0.444444          0            3

      Has_Additional_Services  Churn
0                      1      0
1                      1      0
2                      1      0
3                      0      0
```

```

4          0      1
...
9995      ...    ...
9996          0      0
9997          1      0
9998          1      0
9999          1      0

```

[9998 rows x 10 columns]

[ ]:

Define Features (X) and Response/Label (y)

```
[50]: X= df3.drop(columns=['Churn'])
y=df3['Churn']
```

```
[51]: # We saw earlier that the data is not balanced. We apply SMOTE sampling
      ↴technique to our imbalanced data
```

```
smote = SMOTE()

X_res, y_res = smote.fit_resample(X,y)
```

```
[52]: from collections import Counter

# Count occurrences in resampled data
class_counts = Counter(y_res) # Get counts of Churn (1) and No Churn (0)

l = ['Churn', 'No Churn']

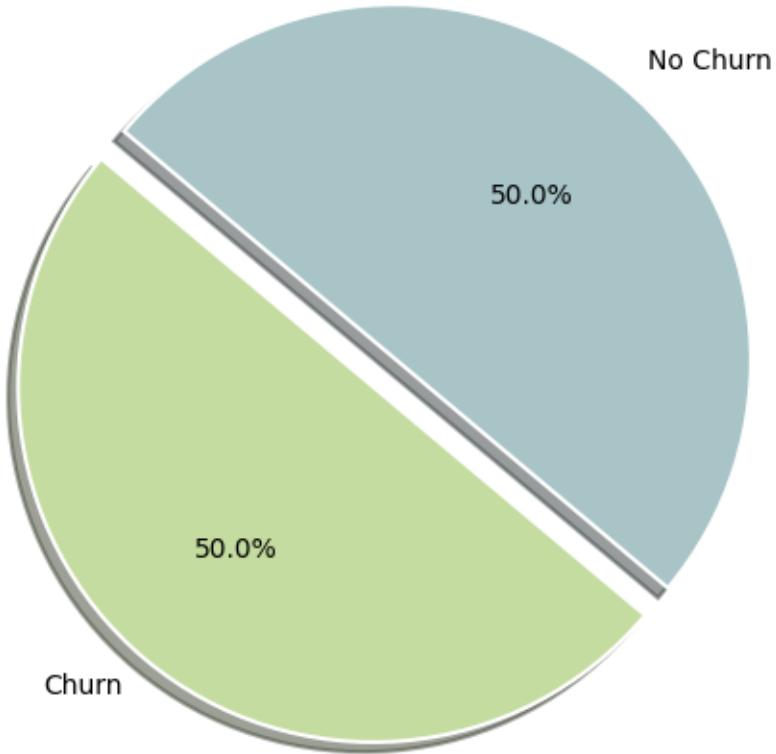
# Values for pie chart
s = [class_counts[1], class_counts[0]] # Count of churned and non-churned
      ↴customers

c = ['#C5DCA0', '#A8C4C7']

# Plot pie chart
plt.figure(figsize=(6,6))
plt.pie(s, labels=l, colors=c, explode=(0.1, 0), shadow=True, autopct='%1.1f%%',
        startangle=140, wedgeprops={'edgecolor': 'white', 'linewidth': 1.5})

plt.title("Class Distribution After SMOTE")
plt.show()
```

Class Distribution After SMOTE



```
[53]: y_res
```

```
[53]: 0      0
      1      0
      2      0
      3      0
      4      1
      ..
15779    1
15780    1
15781    1
15782    1
15783    1
Name: Churn, Length: 15784, dtype: int64
```

```
[54]: y_res.value_counts()  # the response/ label has been balanced
```

```
[54]: Churn  
0    7892  
1    7892  
Name: count, dtype: int64
```

```
[55]: # split into training and testing test  
  
X_train, X_test, y_train, y_test = train_test_split(X_res, y_res, test_size=0.  
↪3, random_state=8)
```

```
[ ]:
```

Generating Models

```
[57]: # Naive Bayes Model  
  
nb_model = GaussianNB()  
nb_model.fit(X_train, y_train)  
y_pred_nb = nb_model.predict(X_test)
```

```
[58]: # Decision Tree Model
```

```
dt_model = DecisionTreeClassifier()  
dt_model.fit(X_train, y_train)  
y_pred_dt = dt_model.predict(X_test)
```

```
[59]: # Random Forest Model
```

```
rf_model = RandomForestClassifier()  
rf_model.fit(X_train, y_train)  
y_pred_rf=rf_model.predict(X_test)
```

```
[60]: # KNN model
```

```
knn_model =KNeighborsClassifier(n_neighbors=3)  
# For n_neighbors =5, accuracy is 0.72  
knn_model.fit(X_train, y_train)  
y_pred_knn = knn_model.predict(X_test)
```

```
[ ]:
```

Evaluation: Confusion Matrix and Other Metrics

```
[62]: def evaluate_model(y_test, y_pred, model_name):  
    print(f"Model: {model_name}")  
    print("Confusion Matrix:")  
    print(confusion_matrix(y_test, y_pred))  
    print(f"Accuracy: {accuracy_score(y_test, y_pred):.2f}")
```

```
print(f"Precision: {precision_score(y_test, y_pred):.2f}")
print(f"Recall: {recall_score(y_test, y_pred):.2f}")
print(f"F1 Score: {f1_score(y_test, y_pred):.2f}\n")
```

[63]: # Evaluate all models

```
evaluate_model(y_test, y_pred_nb, "Naive Bayes")
evaluate_model(y_test, y_pred_dt, "Decision Tree")
evaluate_model(y_test, y_pred_rf, "Random Forest")
evaluate_model(y_test, y_pred_knn, "KNN")
```

Model: Naive Bayes

Confusion Matrix:

```
[[1348 1008]
 [ 788 1592]]
```

Accuracy: 0.62

Precision: 0.61

Recall: 0.67

F1 Score: 0.64

Model: Decision Tree

Confusion Matrix:

```
[[1763 593]
 [ 552 1828]]
```

Accuracy: 0.76

Precision: 0.76

Recall: 0.77

F1 Score: 0.76

Model: Random Forest

Confusion Matrix:

```
[[2277 79]
 [ 676 1704]]
```

Accuracy: 0.84

Precision: 0.96

Recall: 0.72

F1 Score: 0.82

Model: KNN

Confusion Matrix:

```
[[1282 1074]
 [ 593 1787]]
```

Accuracy: 0.65

Precision: 0.62

Recall: 0.75

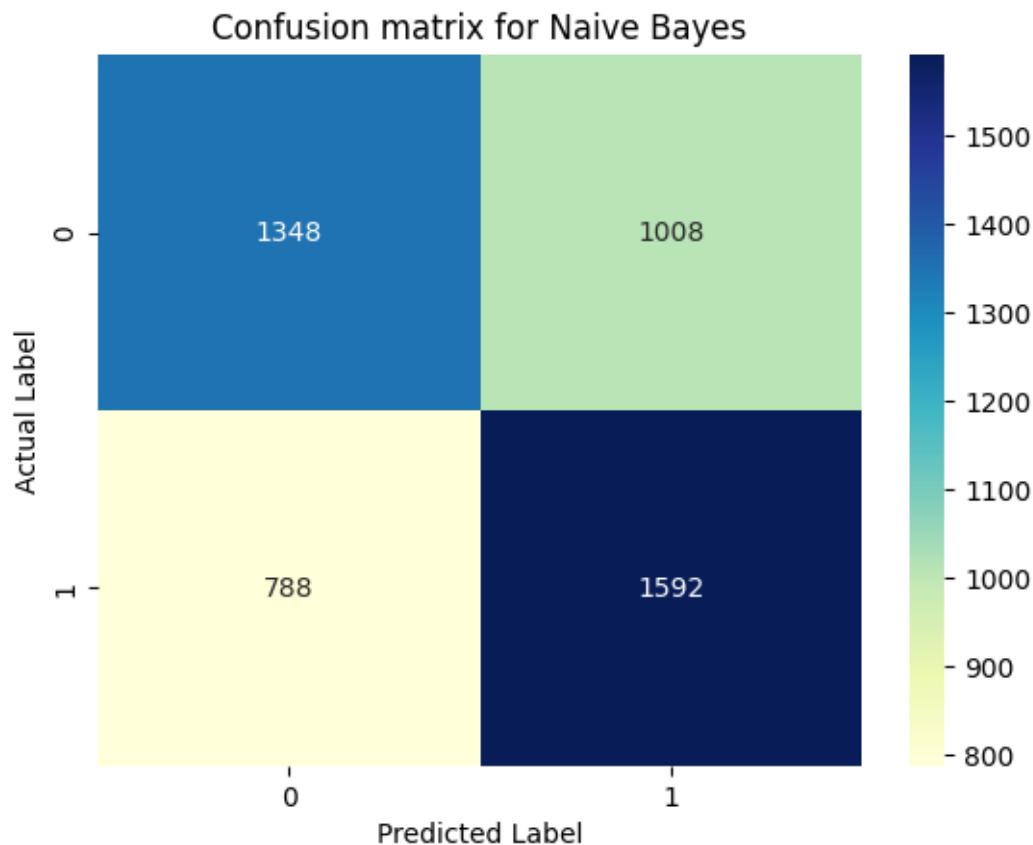
F1 Score: 0.68

```
[64]: # plot confusion matrix using seaborn library
```

```
cm1=confusion_matrix(y_test, y_pred_nb)
cm2=confusion_matrix(y_test, y_pred_dt)
cm3=confusion_matrix(y_test, y_pred_rf)
cm4=confusion_matrix(y_test, y_pred_knn)
```

```
[65]: # confusion matrix for Naive Bayes
```

```
sns.heatmap(cm1, annot=True, cmap='YlGnBu', fmt='g')
plt.ylabel('Actual Label')
plt.xlabel('Predicted Label')
plt.title('Confusion matrix for Naive Bayes');
```



Does overfitting happen for Naive Bayes? We need to compare the accuracy of testing against accuracy of training

```
[67]: print ("Accuracy of Naive Bayes for training database is:", nb_model.
    .score(X_train, y_train))
```

Accuracy of Naive Bayes for training database is: 0.6366763215061549

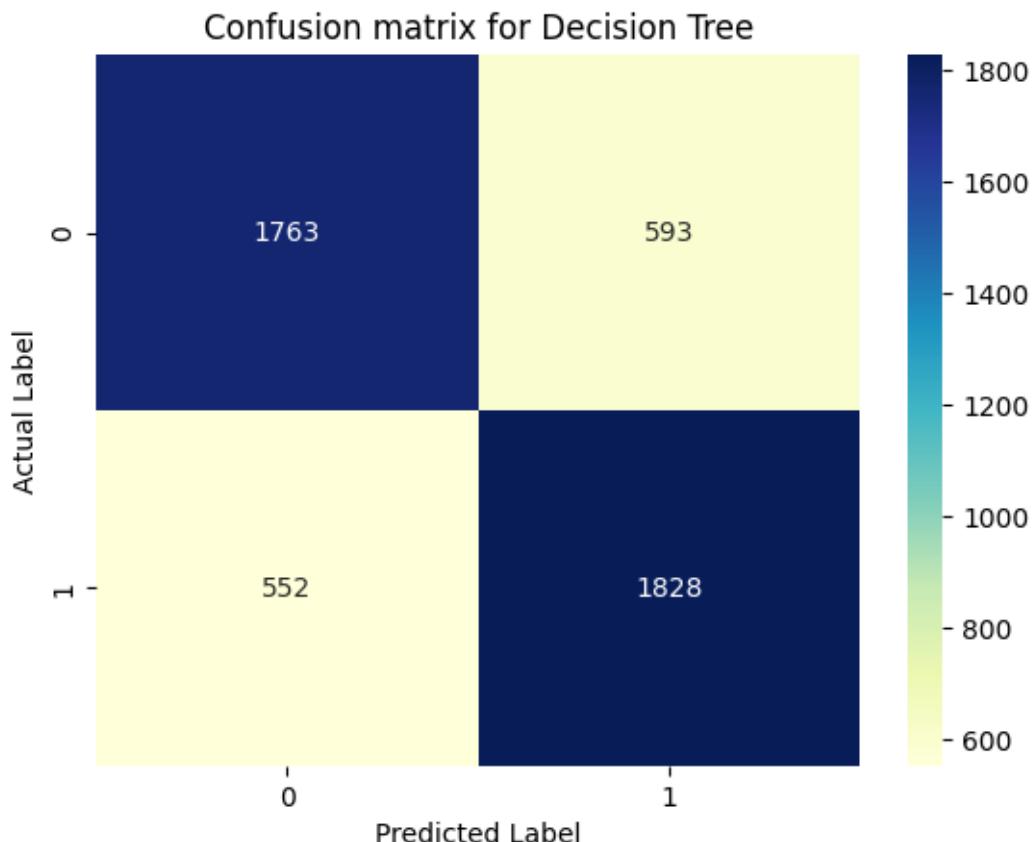
```
[68]: print ("Accuracy of Naive Bayes for testing datatse is:", nb_model.score(X_test, y_test))
```

Accuracy of Naive Bayes for testing datatse is: 0.620777027027027

Naive Bayes thus doesn't overfit. Underfits actually and thus has poor accuracy for both train and test datasets.

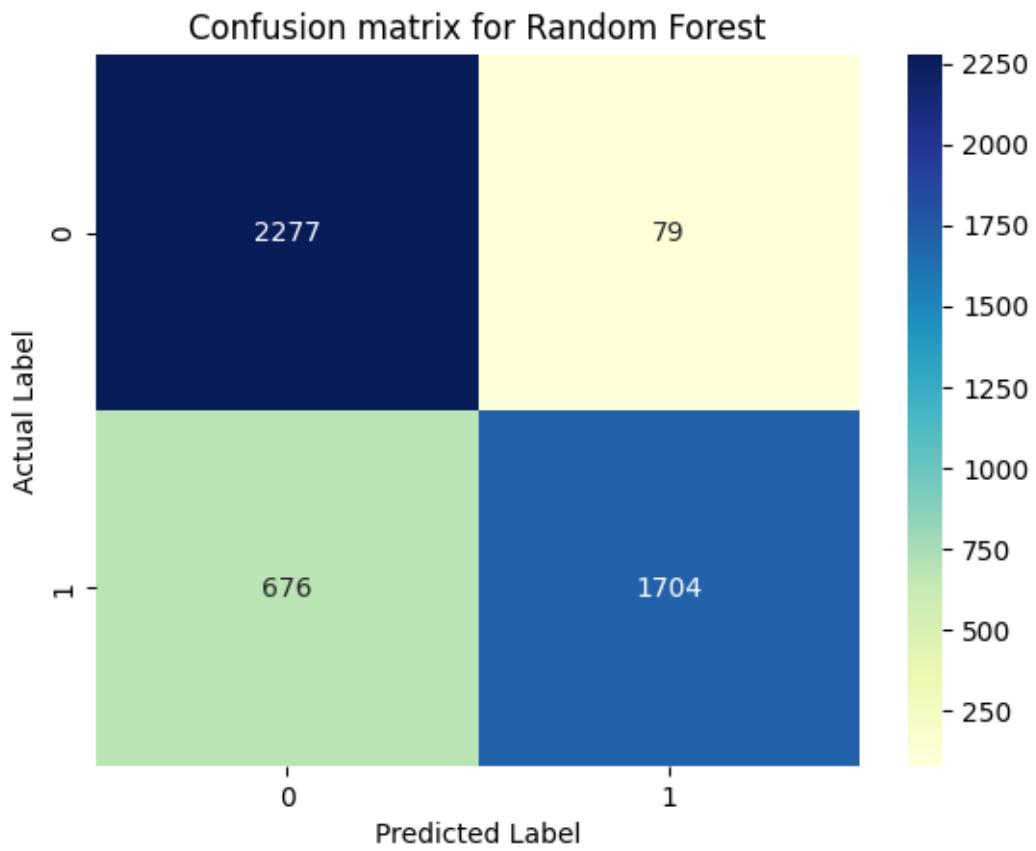
```
[70]: # confusion matrix for Decision Tree
```

```
sns.heatmap(cm2, annot=True, cmap='YlGnBu', fmt='g')
plt.ylabel('Actual Label')
plt.xlabel('Predicted Label')
plt.title('Confusion matrix for Decision Tree');
```



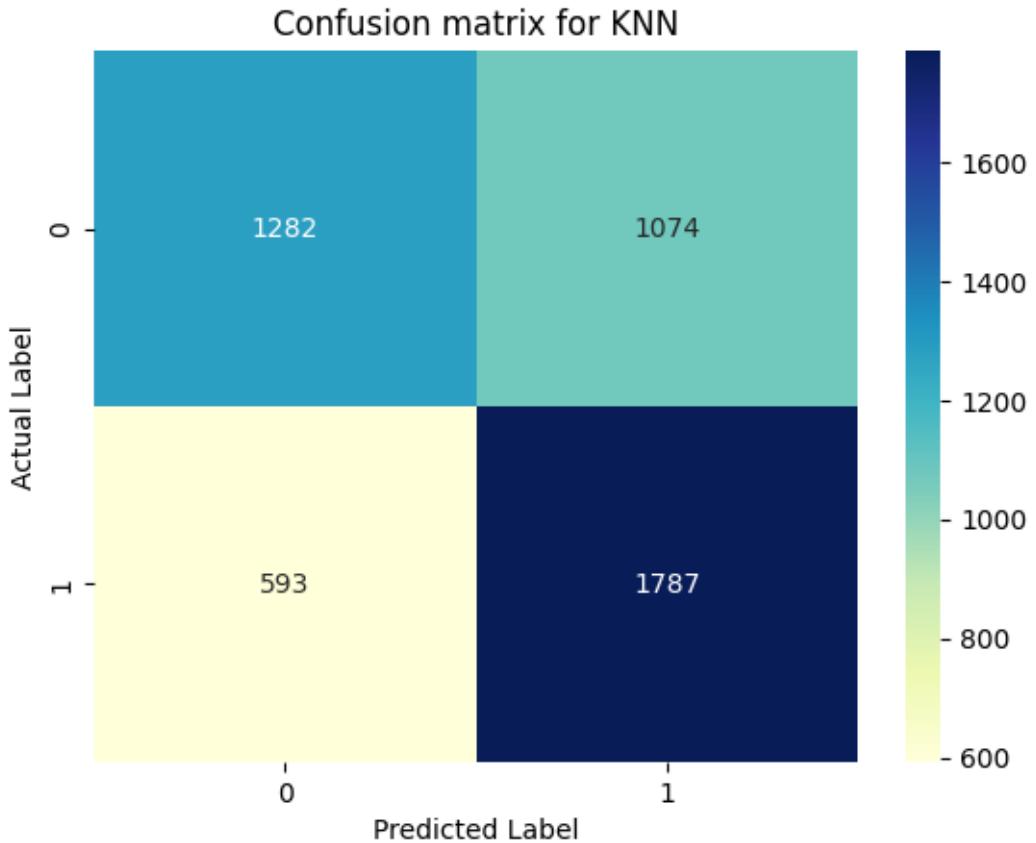
```
[71]: # confusion matrix for Random Forest
```

```
sns.heatmap(cm3, annot=True, cmap='YlGnBu', fmt='g')
plt.ylabel('Actual Label')
plt.xlabel('Predicted Label')
plt.title('Confusion matrix for Random Forest');
```



```
[72]: # confusion matrix for KNN

sns.heatmap(cm4, annot=True, cmap='YlGnBu', fmt='g')
plt.ylabel('Actual Label')
plt.xlabel('Predicted Label')
plt.title('Confusion matrix for KNN');
```



```
[73]: # Random forest has the best performance. let's put all its measures in a report
```

```
print(classification_report(y_test, y_pred_rf))
```

	precision	recall	f1-score	support
0	0.77	0.97	0.86	2356
1	0.96	0.72	0.82	2380
accuracy			0.84	4736
macro avg	0.86	0.84	0.84	4736
weighted avg	0.86	0.84	0.84	4736

```
[ ]:
```

Plot ROC curve for Xtrain and Xtest

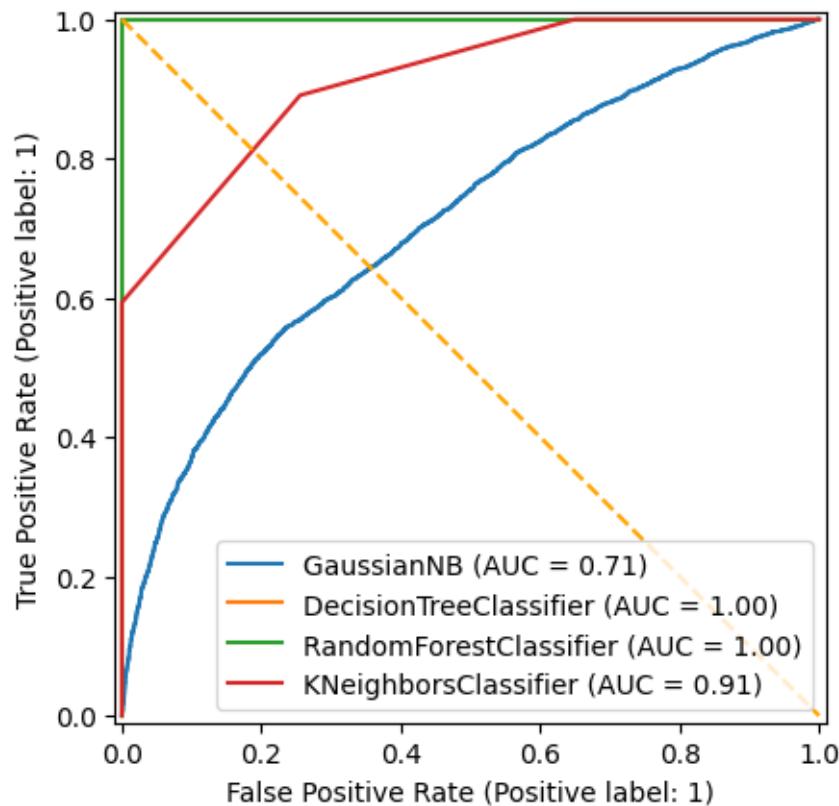
```
[75]: # ROC Curve for Training dataset
```

```
from sklearn.metrics import RocCurveDisplay
```

```

disp =RocCurveDisplay.from_estimator(nb_model, X_train, y_train)
RocCurveDisplay.from_estimator(dt_model, X_train, y_train, ax=disp.ax_)
RocCurveDisplay.from_estimator(rf_model, X_train, y_train, ax=disp.ax_)
RocCurveDisplay.from_estimator(knn_model, X_train, y_train, ax=disp.ax_)
plt.plot([0,1], [1,0], color='orange', linestyle='--');

```



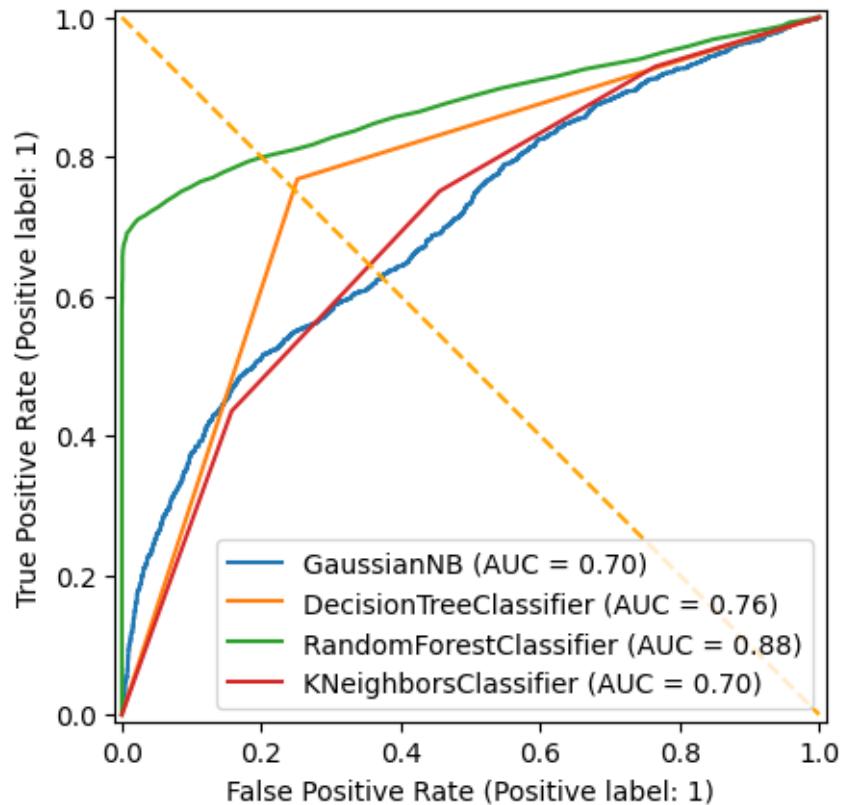
[ ]:

[76]: # ROC Curve for Testing dataset

```

disp =RocCurveDisplay.from_estimator(nb_model, X_test, y_test)
RocCurveDisplay.from_estimator(dt_model, X_test, y_test, ax=disp.ax_)
RocCurveDisplay.from_estimator(rf_model, X_test, y_test, ax=disp.ax_)
RocCurveDisplay.from_estimator(knn_model, X_test, y_test, ax=disp.ax_)
plt.plot([0,1], [1,0], color='orange', linestyle='--');

```



[ ]:

Evaluate model for new dataset

```
[78]: # Given 3 new customers
customer_df = pd.DataFrame({
    'CustomerID': [10001, 10002, 10003],
    'Age': [43.296, 55.704, 35.928],
    'Subscription_Length_Months': [33.66, 48.45, 27.08],
    'Monthly_Bill': [107.8, 144.6, 71.65],
    'Total_Usage_GB': [230.95, 398.05, 131.65],
    'Support_Calls': [1, 9, 1],
    'Contract_Type': ['Month-to-Month', 'One Year', 'Month-to-Month'],
    'Payment_Method': ['Debit Card', 'Credit Card', 'PayPal'],
    'Has_Additional_Services': [1, 1, 0]
})
```

```
[79]: customer_df
```

	CustomerID	Age	Subscription_Length_Months	Monthly_Bill	\
0	10001	43.296		33.66	107.80

```

1      10002  55.704          48.45      144.60
2      10003  35.928          27.08      71.65

  Total_Usage_GB  Support_Calls  Contract_Type Payment_Method \
0        230.95           1 Month-to-Month     Debit Card
1        398.05           9 One Year       Credit Card
2        131.65           1 Month-to-Month     PayPal

Has_Additional_Services
0            1
1            1
2            0

```

[80]: # Make a new prediction for new customers (customer\_df)

```

# load new observations (new customers data normalized and encoded)
new_data = pd.DataFrame({
    'CustomerID': [10001, 10002, 10003],
    'Age': [-0.142, 0.892, -0.756],
    'Subscription_Length_Months': [-0.234, 1.245, -0.892],
    'Monthly_Bill': [0.156, 0.892, -0.567],
    'Total_Usage_GB': [-0.127, 0.987, -0.789],
    'Support_Calls': [0, 1, 0],
    'Contract_Type': [0, 1, 0],
    'Payment_Method': [1, 0, 2],
    'Has_Additional_Services': [1, 1, 0]
})

```

[81]: new\_data.head(5)

```

[81]: CustomerID  Age  Subscription_Length_Months  Monthly_Bill \
0      10001 -0.142                  -0.234      0.156
1      10002  0.892                  1.245      0.892
2      10003 -0.756                  -0.892     -0.567

  Total_Usage_GB  Support_Calls  Contract_Type Payment_Method \
0        -0.127           0           0           1
1         0.987           1           1           0
2        -0.789           0           0           2

Has_Additional_Services
0            1
1            1
2            0

```

[82]: # Create new data with exact same structure as X\_train  
new\_data\_standardized = pd.DataFrame({

```

'CustomerID': [10001, 10002, 10003],
'Age': [-0.142, 0.892, -0.756],
'Subscription_Length_Months': [-0.234, 1.245, -0.892],
'Monthly_Bill': [0.156, 0.892, -0.567],
'Total_Usage_GB': [-0.127, 0.987, -0.789],
'Support_Calls': [0, 1, 0],
'Contract_Type': [0, 1, 0],
'Payment_Method': [1, 0, 2],
'Has_Additional_Services': [1, 1, 0]
})

# Make predictions
y_pred_new = rf_model.predict(new_data_standardized)
print("Predictions for new customers:", y_pred_new)

# Get prediction probabilities (optional)
y_pred_prob = rf_model.predict_proba(new_data_standardized)
print("\nPrediction probabilities:")
for i, prob in enumerate(y_pred_prob):
    print(f"Customer {i+1} - No Churn: {prob[0]:.2f}, Churn: {prob[1]:.2f}")

```

Predictions for new customers: [0 0 0]

Prediction probabilities:

Customer 1 - No Churn: 0.56, Churn: 0.44

Customer 2 - No Churn: 0.68, Churn: 0.32

Customer 3 - No Churn: 0.65, Churn: 0.35

[83]: y\_pred\_new\_observation=rf\_model.predict(new\_data)

[84]: y\_pred\_new\_observation

[84]: array([0, 0, 0])

[ ]:

Find the Importance of Features

[86]: # ----- Features importance using Random Forest

```
rf=RandomForestClassifier(random_state=1, n_estimators=100)
```

[87]: rf.fit(X,y)

[87]: RandomForestClassifier(random\_state=1)

[88]: I = rf.feature\_importances\_

[89]: I

```
[89]: array([0.18267379, 0.13463516, 0.13316602, 0.18010421, 0.18317085,
   0.07943338, 0.03809709, 0.04785989, 0.02085962])
```

```
[90]: df3.columns
```

```
[90]: Index(['CustomerID', 'Age', 'Subscription_Length_Months', 'Monthly_Bill',
   'Total_Usage_GB', 'Support_Calls', 'Contract_Type', 'Payment_Method',
   'Has_Additional_Services', 'Churn'],
  dtype='object')
```

```
[91]: # we need the name of features columns. df3.columns include Churn name. we need
   ↵to drop Churn

# Define a variable that contains only the feature column names
# First drop CustomerID since it adds no value to target prediction
#df_no_id = df3.drop(['CustomerID'], axis=1)

# Then drop Churn from df_no_id
name_features = df3.drop(df3['Churn']).columns.tolist()
name_features
```

```
[91]: ['CustomerID',
   'Age',
   'Subscription_Length_Months',
   'Monthly_Bill',
   'Total_Usage_GB',
   'Support_Calls',
   'Contract_Type',
   'Payment_Method',
   'Has_Additional_Services',
   'Churn']
```

```
[92]: len(name_features)
```

```
[92]: 10
```

```
[93]: len(I)
```

```
[93]: 9
```

```
[94]: # generate a new dataframe call it df_I

df_I=pd.DataFrame([I, name_features])
```

```
[95]: df_I
```

```
[95]:          0           1           2           3  \
0    0.182674  0.134635  0.133166  0.180104
```

```

1 CustomerID      Age   Subscription_Length_Months  Monthly_Bill
                4           5           6           7   \
0          0.183171    0.079433    0.038097    0.04786
1 Total_Usage_GB  Support_Calls  Contract_Type  Payment_Method
                8           9
0             0.02086  None
1 Has_Additional_Services  Churn

```

```
[96]: # Remove 'Churn' and 'CustomerID' from name_features
name_features = name_features[:-1] # Remove the last element (Churn)
name_features_filtered = [feature for feature in name_features[:-1] if feature != 'CustomerID'] # Remove CustomerID as it isn't a target determining attribute

# Remove the first element in I (corresponding to CustomerID)
I_filtered = I[1:-1]

# Create the DataFrame
feature_importance_df = pd.DataFrame({
    'Feature': name_features_filtered,
    'Importance': I_filtered
})

# Sort by importance value
feature_importance_df = feature_importance_df.sort_values('Importance', ascending=False)

# Create the barplot with different colors
plt.figure(figsize=(10, 6))

# Option 1: Using a colorful seaborn palette
colors = sns.color_palette("husl", len(feature_importance_df)) # Create a colorful palette

# Create the plot with different colors for each bar
ax = sns.barplot(
    data=feature_importance_df,
    x='Importance',
    y='Feature',
    palette=colors
)

# Add value labels to the bars for better readability
for i, v in enumerate(feature_importance_df['Importance']):
    ax.text(v + 0.001, i, f"{v:.4f}", va='center')

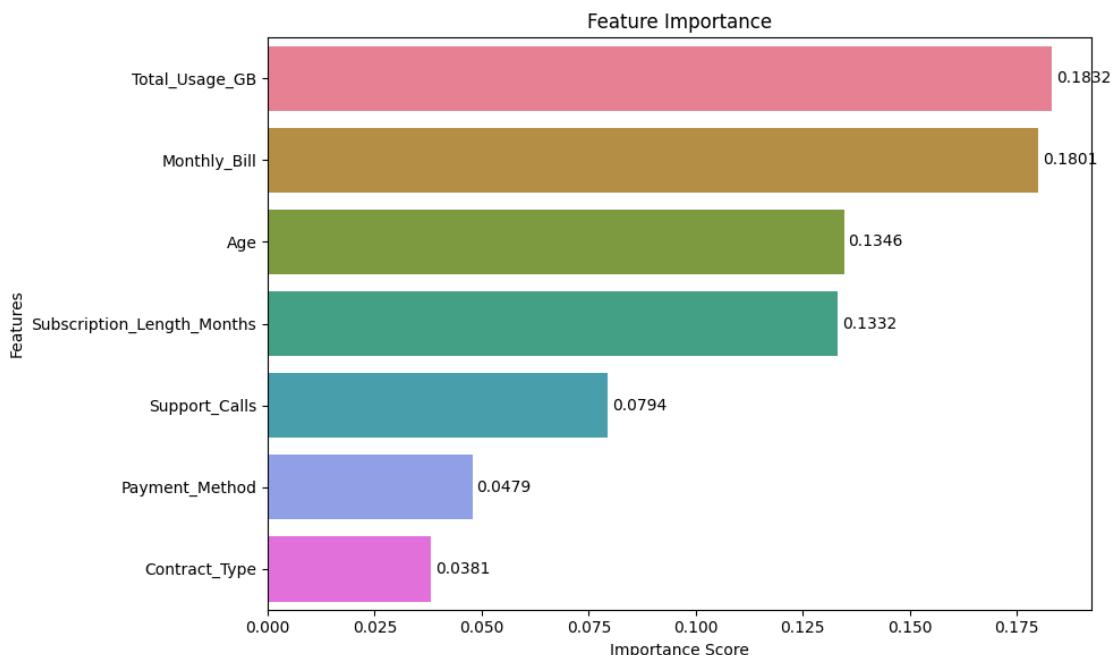
```

```
plt.title('Feature Importance')
plt.xlabel('Importance Score')
plt.ylabel('Features')
plt.tight_layout()
plt.show()
```

```
/var/folders/9p/02j0tmhs5w35zmgt4_pc9src0000gn/T/ipykernel_4362/2523790702.py:24
: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
ax = sns.barplot(
```



```
[ ]:
```

Find the best value of K in KNN

```
[98]: from sklearn.metrics import balanced_accuracy_score
```

```
[99]: error=[]
```

```
#calculating the error for k values between 1 and 39
```

```

for i in range(1,21): # i=1,2,...,20
    knn = KNeighborsClassifier(n_neighbors=i)
    knn.fit(X_train, y_train)
    pred_i = knn.predict(X_test)
    accuracy = balanced_accuracy_score(y_test, pred_i)
    error.append(1-accuracy)

error;

```

```

[100]: import matplotlib.pyplot as plt

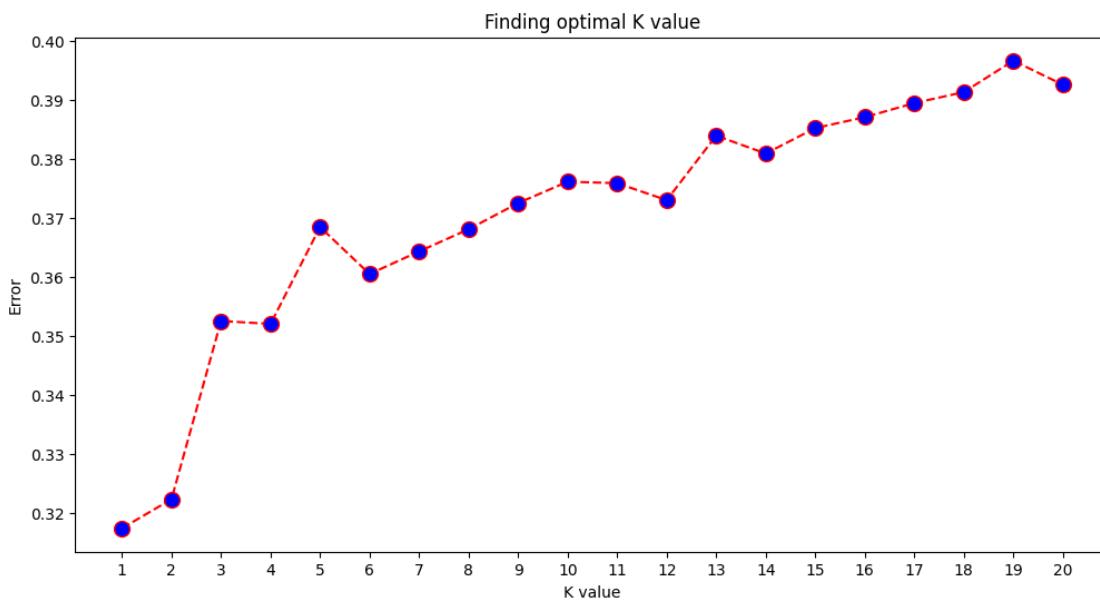
plt.figure(figsize=(12,6))

plt.plot(range(1,21), error, color='red', linestyle='dashed', marker='o', markerfacecolor='blue', markersize=10)

plt.title ('Finding optimal K value')
plt.xlabel ('K value')
plt.ylabel('Error');

# Set the ticks on the X-axis to integer values
plt.xticks(ticks=range(1, 21)); # Set the ticks to integers 1 to 20

```



Remodeling KNN using the optimal K value

```
[102]: # KNN model with n=1
```

```

knn_model =KNeighborsClassifier(n_neighbors=1)
# For n_neighbors =3, accuracy is 0.74
knn_model.fit(X_train, y_train)
y_pred_knn = knn_model.predict(X_test)
evaluate_model(y_test, y_pred_knn, "KNN")

```

Model: KNN  
 Confusion Matrix:  
 [[1313 1043]  
 [ 457 1923]]  
 Accuracy: 0.68  
 Precision: 0.65  
 PRecall: 0.81  
 F1 Score: 0.72

[103]: # KNN model with n=2

```

knn_model =KNeighborsClassifier(n_neighbors=2)
# For n_neighbors =3, accuracy is 0.74
knn_model.fit(X_train, y_train)
y_pred_knn = knn_model.predict(X_test)
evaluate_model(y_test, y_pred_knn, "KNN")

```

Model: KNN  
 Confusion Matrix:  
 [[1758 598]  
 [ 930 1450]]  
 Accuracy: 0.68  
 Precision: 0.71  
 PRecall: 0.61  
 F1 Score: 0.65

[104]: # KNN model with n=4

```

knn_model =KNeighborsClassifier(n_neighbors=4)
# For n_neighbors =3, accuracy is 0.74
knn_model.fit(X_train, y_train)
y_pred_knn = knn_model.predict(X_test)
evaluate_model(y_test, y_pred_knn, "KNN")

```

Model: KNN  
 Confusion Matrix:  
 [[1611 745]  
 [ 923 1457]]  
 Accuracy: 0.65  
 Precision: 0.66  
 PRecall: 0.61

F1 Score: 0.64

As **n** increases in KNN, the model's performance varies. At **n=1**, the model achieves the highest recall (0.95) but at the cost of lower precision (0.74), indicating a tendency to favor positive predictions. While accuracy remains relatively stable (~0.74–0.81) across different values of **n**, a trade-off between precision and recall becomes evident. **n=2** offers the best balance (Precision: 0.78, Recall: 0.83, F1: 0.80). However, performance declines for **n=3** and **n=4**, suggesting that increasing the number of neighbors beyond **n=2** diminishes the model's effectiveness.

[ ]:

[ ]:

Conditional Probabilities

[107]: # 1. What is the probability of churn for customers with a Two Year contract and subscription length over 36 months?  
#  $P(\text{Churn} = 1 | \text{Contract\_Type} = \text{"Two Year"}, \text{Subscription\_Length\_Months} > 36)$   
 $((\text{df}[\text{'Contract\_Type'}] == \text{"Two Year"}) \& (\text{df}[\text{'Subscription\_Length\_Months'}] > 36)) \cdot \frac{(\text{df}[\text{'Churn'}] == 1).sum()}{((\text{df}[\text{'Contract\_Type'}] == \text{"Two Year"}) \& (\text{df}[\text{'Subscription\_Length\_Months'}] > 36)).sum()}$

[107]: 0.2153846153846154

[ ]:

[108]: # 2. What is the probability of churn for customers with usage over 300GB and monthly bill over \$150?  
#  $P(\text{Churn} = 1 | \text{Total\_Usage\_GB} > 300, \text{Monthly\_Bill} > 150)$   
 $((\text{df}[\text{'Total\_Usage\_GB'}] > 300) \& (\text{df}[\text{'Monthly\_Bill'}] > 150) \& (\text{df}[\text{'Churn'}] == 1)).sum() / ((\text{df}[\text{'Total\_Usage\_GB'}] > 300) \& (\text{df}[\text{'Monthly\_Bill'}] > 150)).sum()$

[108]: 0.20408163265306123

[ ]:

[109]: # 3. What is the probability of churn for customers with more than 5 support calls?  
#  $P(\text{Churn} = 1 | \text{Support\_Calls} > 5)$   
 $((\text{df}[\text{'Support\_Calls'}] > 5) \& (\text{df}[\text{'Churn'}] == 1)).sum() / (\text{df}[\text{'Support\_Calls'}] > 5).sum()$

[109]: 0.20586062546170894

[ ]:

[110]: # 4. What is the probability of churn for customers aged 60+ with low usage (<100GB)?

```
# P(Churn = 1 / Age >= 60, Total_Usage_GB < 100)
((df['Age'] >= 60) & (df['Total_Usage_GB'] < 100) & (df['Churn'] == 1)).sum() / 
    ((df['Age'] >= 60) & (df['Total_Usage_GB'] < 100)).sum()
```

[110]: 0.18495934959349594

[ ]:

[111]: # 5. What is the probability of churn for customers with monthly bill over \$90?  
# P(Churn = 1 / Monthly\_Bill > 90)  
((df['Monthly\_Bill'] > 90) & (df['Churn'] == 1)).sum() / (df['Monthly\_Bill'] > 90).sum()

[111]: 0.20373277516134658

[ ]:

[112]: # 6. What is the probability that a heavy usage customer (>300GB) has a  
 ↪Month-to-Month contract?  
# P(Contract\_Type = "Month-to-Month" / Total\_Usage\_GB > 300)  
((df['Total\_Usage\_GB'] > 300) & (df['Contract\_Type'] == "Month-to-Month")).  
 ↪sum() / (df['Total\_Usage\_GB'] > 300).sum()

[112]: 0.33398821218074654

[ ]:

[113]: # 7. What is the probability of churn for customers with more than 5 support  
 ↪calls, a Month-to-Month contract, and monthly bill over \$100?  
# P(Churn = 1 / Support\_Calls > 5, Contract\_Type = "Month-to-Month",  
 ↪Monthly\_Bill > 100)  
((df['Support\_Calls'] > 5) & (df['Contract\_Type'] == "Month-to-Month") &  
 ↪(df['Monthly\_Bill'] > 100) & (df['Churn'] == 1)).sum() /  
 ↪((df['Support\_Calls'] > 5) & (df['Contract\_Type'] == "Month-to-Month") &  
 ↪(df['Monthly\_Bill'] > 100)).sum()

[113]: 0.20084865629420084

[ ]:

Feature Importance using Naive Bayes

[115]: # Determining the significance of Monthly Bill when predicting customer churn  
# (Churn = 1) as the target variable.  
  
# Conditional probabilities for high vs low monthly bill  
median\_bill = df['Monthly\_Bill'].median()

```

prob_high_bill = ((df['Monthly_Bill'] > median_bill) & (df['Churn'] == 1)) .
    ↪sum() / (df['Monthly_Bill'] > median_bill).sum()
prob_low_bill = ((df['Monthly_Bill'] <= median_bill) & (df['Churn'] == 1)) .
    ↪sum() / (df['Monthly_Bill'] <= median_bill).sum()

print("Conditional probabilities:", "\nHigh_Bill:", prob_high_bill, "\nLow_Bill:
    ↪", prob_low_bill)

# Calculating difference to determine significance
bill_difference = abs(prob_high_bill - prob_low_bill)
print("Feature importance of Monthly Bill is:")
print("Difference:", bill_difference)

```

Conditional probabilities:  
 High\_Bill: 0.2002  
 Low\_Bill: 0.2212  
 Feature importance of Monthly Bill is:  
 Difference: 0.021000000000000002

[ ]:

```

[116]: # Determining the significance of Total Usage when predicting customer churn
median_usage = df['Total_Usage_GB'].median()
prob_high_usage = ((df['Total_Usage_GB'] > median_usage) & (df['Churn'] == 1)) .
    ↪sum() / (df['Total_Usage_GB'] > median_usage).sum()
prob_low_usage = ((df['Total_Usage_GB'] <= median_usage) & (df['Churn'] == 1)) .
    ↪sum() / (df['Total_Usage_GB'] <= median_usage).sum()

print("\nConditional probabilities:", "High_Usage:", prob_high_usage, "Low_Usage:",
    ↪prob_low_usage)

# Calculating difference to determine significance
usage_difference = abs(prob_high_usage - prob_low_usage)
print("Feature importance of Total Usage is:")
print("Difference:", usage_difference)

```

Conditional probabilities: High\_Usage: 0.21 Low\_Usage: 0.2114  
 Feature importance of Total Usage is:  
 Difference: 0.001400000000000123

[ ]:

```

[117]: # Determining the significance of Age > 30 when predicting customer churn
prob_age_over_30 = ((df['Age'] > 30) & (df['Churn'] == 1)).sum() / (df['Age'] >
    ↪30).sum()
prob_age_under_30 = ((df['Age'] <= 30) & (df['Churn'] == 1)).sum() / (df['Age'] <
    ↪30).sum()

```

```

print("\nConditional probabilities:", "Age>30:", prob_age_over_30, "Age<=30:", prob_age_under_30)

# Calculating difference to determine significance
age_difference = abs(prob_age_over_30 - prob_age_under_30)
print("Feature importance of Age > 30 is:")
print("Difference:", age_difference)

```

Conditional probabilities: Age>30: 0.21192648759799512 Age<=30:  
0.206399278954484  
Feature importance of Age > 30 is:  
Difference: 0.005527208643511111

[ ]:

```

[118]: # Determining the significance of Subscription Length when predicting customer churn
median_sub_length = df['Subscription_Length_Months'].median()
prob_long_sub = ((df['Subscription_Length_Months'] > median_sub_length) & (df['Churn'] == 1)).sum() / (df['Subscription_Length_Months'] > median_sub_length).sum()
prob_short_sub = ((df['Subscription_Length_Months'] <= median_sub_length) & (df['Churn'] == 1)).sum() / (df['Subscription_Length_Months'] <= median_sub_length).sum()

print("\nConditional probabilities:", "Long_Subscription:", prob_long_sub, "Short_Subscription:", prob_short_sub)

# Calculating difference to determine significance
sub_difference = abs(prob_long_sub - prob_short_sub)
print("Feature importance of Subscription Length is:")
print("Difference:", sub_difference)

```

Conditional probabilities: Long\_Subscription: 0.2132132132132132  
Short\_Subscription: 0.2081918081918082  
Feature importance of Subscription Length is:  
Difference: 0.005021405021405023

[ ]:

```

[119]: # Comparing feature importances
features = ['Monthly_Bill', 'Total_Usage_GB', 'Age > 30', 'Subscription_Length']
importances = [bill_difference, usage_difference, age_difference, sub_difference]

```

```
# Print ranked list of features by importance
for feature, importance in sorted(zip(features, importances), key=lambda x:x[1], reverse=True):
    print(f"{feature}: {importance:.4f}")
```

```
Monthly_Bill: 0.0210
Age > 30: 0.0055
Subscription_Length: 0.0050
Total_Usage_GB: 0.0014
```

```
[ ]:
```

# Regression Analysis

March 3, 2025

## 1 Regression Analysis

```
[2]: # Import necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures
from sklearn.metrics import r2_score, mean_squared_error
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler
import statsmodels.api as sm
from statsmodels.stats.outliers_influence import variance_inflation_factor
from lazypredict.Supervised import LazyRegressor
```

DATASET: House Rent Dataset FROM Kaggle.com

```
[3]: # -----
# Load and preprocess dataset
# -----
df = pd.read_csv(r'House_Rent_Dataset.csv')
```

Data Preprocessing

```
[4]: # -----
# Checking for missing values
# -----
missing_values = df.isnull().sum()
print("Missing values:\n", missing_values)
```

Missing values:

Posted On	0
BHK	0
Rent	0
Size	0
Floor	0

```
Area Type          0
Area Locality     0
City              0
Furnishing Status 0
Tenant Preferred   0
Bathroom          0
Point of Contact   0
dtype: int64
```

```
[5]: def convert_to_floor_number(floor_str):
    """
    Convert floor string to numeric value
    -2: Lower Basement
    -1: Upper Basement
    0: Ground Floor
    1+: Regular floors
    """
    if isinstance(floor_str, str):
        if 'Upper Basement' in floor_str:
            return -1
        elif 'Lower Basement' in floor_str:
            return -2
        elif 'Ground' in floor_str:
            return 0
        else:
            try:
                return int(floor_str.split()[0])
            except:
                return None
    return None

# Add floor_number feature
df['floor_number'] = df['Floor'].apply(convert_to_floor_number)

# Display first few rows to verify
print(df['floor_number'].head())

# Show value counts to verify distribution
print("\nFloor number distribution:")
print(df['floor_number'].value_counts().sort_index())
```

```
0      0
1      1
2      1
3      1
4      1
Name: floor_number, dtype: int64
```

Floor number distribution:

floor\_number

-2	11
-1	23
0	927
1	1161
2	945
3	512
4	272
5	164
6	93
7	74
8	66
9	65
10	67
11	43
12	47
13	15
14	34
15	41
16	21
17	22
18	26
19	16
20	12
21	6
22	3
23	9
24	6
25	12
26	3
27	3
28	4
29	1
30	5
32	3
33	1
34	4
35	3
36	2
37	1
39	1
40	2
41	1
43	1
44	1
45	2
46	1

```

47      1
48      2
49      1
50      1
53      2
60      3
65      3
76      1
Name: count, dtype: int64

```

```
[7]: # Drop unnecessary columns
df_cleaned = df.drop(columns=['Posted On', 'Floor', 'Area Locality',
                             'Tenant Preferred', 'Point of Contact'])
```

```
[12]: # Convert categorical variables into numeric format
furnishing_mapping = {'Unfurnished': 0, 'Semi-Furnished': 1, 'Furnished': 2}
df_cleaned['Furnishing Status'] = df_cleaned['Furnishing Status'].
    ↪map(furnishing_mapping)

df_cleaned = pd.get_dummies(df_cleaned, columns=['Area Type', 'City'],
                           ↪drop_first=True)

# Ensure all columns are numeric
df_cleaned = df_cleaned.apply(pd.to_numeric, errors='coerce')

df_cleaned.head()
```

```
[12]:   BHK    Rent   Size  Furnishing Status  Bathroom  floor_number \
0      2   10000   1100                  0          2            0
1      2   20000    800                  1          1            1
2      2   17000   1000                  1          1            1
3      2   10000    800                  0          1            1
4      2    7500    850                  0          1            1

           Area Type_Carpet Area  Area Type_Super Area  City_Chennai  City_Delhi \
0                False          True        False       False
1                False          True        False       False
2                False          True        False       False
3                False          True        False       False
4                 True          False       False       False

   City_Hyderabad  City_Kolkata  City_Mumbai
0             False         True        False
1             False         True        False
2             False         True        False
3             False         True        False
4             False         True        False
```

```
[14]: # Convert boolean columns to integers (for OLS)
for col in df_cleaned.columns:
    if df_cleaned[col].dtype == 'bool':
        df_cleaned[col] = df_cleaned[col].astype(int)
```

```
[16]: # -----
# Define dependent and independent variables
# -----
X = df_cleaned.drop('Rent', axis=1)
y = df_cleaned['Rent']
```

## Generating Models

```
[18]: # -----
# Split the data into training and testing sets
# -----
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
                                                    random_state=42)
```

```
[20]: # -----
# Multiple Linear Regression using sklearn
# -----
lr = LinearRegression()
lr.fit(X_train, y_train)
y_pred = lr.predict(X_test)

print("== Multiple Linear Regression ==")
r2=r2_score(y_test, y_pred)
print("R2:", r2)
print("MSE:", mean_squared_error(y_test, y_pred))
# Calculating Adjusted R2
n = X_test.shape[0] # Number of samples
p = X_test.shape[1] # Number of predictors
adjusted_r2 = 1 - ((1 - r2) * (n - 1) / (n - p - 1))
print("Adjusted R2:", adjusted_r2)

# -----
# Polynomial Regression (e.g., degree=2)
# -----
poly = PolynomialFeatures(degree=2, include_bias=False)
X_poly_train = poly.fit_transform(X_train)
X_poly_test = poly.transform(X_test)

lr_poly = LinearRegression()
lr_poly.fit(X_poly_train, y_train)
y_poly_pred = lr_poly.predict(X_poly_test)

r2=r2_score(y_test, y_poly_pred)
```

```

print("==== Polynomial Regression (degree=2) ====")
print("R2:", r2)
print("MSE:", mean_squared_error(y_test, y_poly_pred))

# Calculating Adjusted R2
n = X_poly_test.shape[0] # Number of samples
p = X_poly_test.shape[1] # Number of predictors
adjusted_r2 = 1 - ((1 - r2) * (n - 1) / (n - p - 1))
print("Adjusted R2:", adjusted_r2)

print("Interpretation: Since the R2 score is higher for polynomial regression than for multiple linear regression, the model follows a polynomial curve.")
# Extract coefficients and intercept for polynomial regression
coefficients_poly = lr_poly.coef_
intercept_poly = lr_poly.intercept_

# Get feature names
feature_names = poly.get_feature_names_out(X_train.columns)
#print(feature_names)

# Construct the equation
equation = f"Rent = {intercept_poly:.2f}"
for i, coef in enumerate(coefficients_poly):
    equation += f" + {coef:.2f} * {feature_names[i]}"

print("==== Equation of the Model ====")
print(equation)

# -----
# LazyPredict: Evaluate multiple regression models
# -----
reg = LazyRegressor(verbose=0, ignore_warnings=True, custom_metric=None)
models, predictions = reg.fit(X_train, X_test, y_train, y_test)
print("==== LazyPredict Model Comparison ====")
print(models)

```

```

==== Multiple Linear Regression ====
R2: 0.5228843329265807
MSE: 1901487865.6850293
Adjusted R2: 0.5167739935403682
==== Polynomial Regression (degree=2) ====
R2: 0.8137733813128744
MSE: 742184086.8760061
Adjusted R2: 0.7942618613107308
Interpretation: Since the R2 score is higher for polynomial regression than for multiple linear regression, the model follows a polynomial curve.
==== Equation of the Model ====
Rent = 236693.67 + -40788374731.34 * BHK + 81575716.39 * Size + 174315.13 *

```

Furnishing Status + 157807.45 \* Bathroom + 26404.82 \* floor\_number + -112617.09  
 \* Area Type\_Carpet Area + -114163.93 \* Area Type\_Super Area + 52728.54 \*  
 City\_Chennai + -4197.02 \* City\_Delhi + -4768.28 \* City\_Hyderabad + 509.40 \*  
 City\_Kolkata + -8347.23 \* City\_Mumbai + -382.45 \* BHK^2 + -4.24 \* BHK Size +  
 -1763.72 \* BHK Furnishing Status + 6712.43 \* BHK Bathroom + 891.13 \* BHK  
 floor\_number + 40788361229.53 \* BHK Area Type\_Carpet Area + 40788362939.56 \* BHK  
 Area Type\_Super Area + 8132.09 \* BHK City\_Chennai + 3667.74 \* BHK City\_Delhi +  
 7881.57 \* BHK City\_Hyderabad + 10132.94 \* BHK City\_Kolkata + -20261.69 \* BHK  
 City\_Mumbai + 0.00 \* Size^2 + 7.13 \* Size Furnishing Status + 4.03 \* Size  
 Bathroom + -1.07 \* Size floor\_number + -81575669.32 \* Size Area Type\_Carpet Area  
 + -81575685.43 \* Size Area Type\_Super Area + -33.46 \* Size City\_Chennai + -34.51  
 \* Size City\_Delhi + -38.79 \* Size City\_Hyderabad + -42.21 \* Size City\_Kolkata +  
 109.30 \* Size City\_Mumbai + -682.73 \* Furnishing Status^2 + -736.49 \* Furnishing  
 Status Bathroom + 648.78 \* Furnishing Status floor\_number + -175857.84 \*  
 Furnishing Status Area Type\_Carpet Area + -175173.56 \* Furnishing Status Area  
 Type\_Super Area + 1054.59 \* Furnishing Status City\_Chennai + 1934.58 \*  
 Furnishing Status City\_Delhi + -679.30 \* Furnishing Status City\_Hyderabad +  
 1298.05 \* Furnishing Status City\_Kolkata + 6327.87 \* Furnishing Status  
 City\_Mumbai + -3317.53 \* Bathroom^2 + 141.87 \* Bathroom floor\_number +  
 -165789.41 \* Bathroom Area Type\_Carpet Area + -165857.30 \* Bathroom Area  
 Type\_Super Area + 4721.69 \* Bathroom City\_Chennai + 22330.81 \* Bathroom  
 City\_Delhi + 1478.58 \* Bathroom City\_Hyderabad + 5507.58 \* Bathroom City\_Kolkata  
 + 4141.27 \* Bathroom City\_Mumbai + -14.38 \* floor\_number^2 + -28583.44 \*  
 floor\_number Area Type\_Carpet Area + -28223.98 \* floor\_number Area Type\_Super  
 Area + 768.83 \* floor\_number City\_Chennai + -360.44 \* floor\_number City\_Delhi +  
 1869.63 \* floor\_number City\_Hyderabad + -113.98 \* floor\_number City\_Kolkata +  
 1164.90 \* floor\_number City\_Mumbai + -112617.06 \* Area Type\_Carpet Area^2 + 0.00  
 \* Area Type\_Carpet Area Area Type\_Super Area + -109708.23 \* Area Type\_Carpet  
 Area City\_Chennai + -8605.85 \* Area Type\_Carpet Area City\_Delhi + 7524.22 \* Area  
 Type\_Carpet Area City\_Hyderabad + -4002.08 \* Area Type\_Carpet Area City\_Kolkata  
 + 1563.71 \* Area Type\_Carpet Area City\_Mumbai + -114163.93 \* Area Type\_Super  
 Area^2 + -100236.42 \* Area Type\_Super Area City\_Chennai + 4408.78 \* Area  
 Type\_Super Area City\_Delhi + 23599.69 \* Area Type\_Super Area City\_Hyderabad +  
 4511.47 \* Area Type\_Super Area City\_Kolkata + -9910.94 \* Area Type\_Super Area  
 City\_Mumbai + 52728.53 \* City\_Chennai^2 + 0.00 \* City\_Chennai City\_Delhi + 0.00 \*  
 \* City\_Chennai City\_Hyderabad + 0.00 \* City\_Chennai City\_Kolkata + 0.00 \*  
 City\_Chennai City\_Mumbai + -4197.07 \* City\_Delhi^2 + 0.00 \* City\_Delhi  
 City\_Hyderabad + 0.00 \* City\_Delhi City\_Kolkata + 0.00 \* City\_Delhi City\_Mumbai  
 + -4768.27 \* City\_Hyderabad^2 + 0.00 \* City\_Hyderabad City\_Kolkata + 0.00 \*  
 City\_Hyderabad City\_Mumbai + 509.39 \* City\_Kolkata^2 + 0.00 \* City\_Kolkata  
 City\_Mumbai + -8347.24 \* City\_Mumbai^2

100% | 42/42 [00:05<00:00, 7.81it/s]

[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000177 seconds.

You can set `force\_row\_wise=true` to remove the overhead.

And if memory is not enough, you can set `force\_col\_wise=true`.

[LightGBM] [Info] Total Bins 329

[LightGBM] [Info] Number of data points in the train set: 3796, number of used features: 12

[LightGBM] [Info] Start training from score 35151.516333

==== LazyPredict Model Comparison ===

Model	Adjusted R-Squared	R-Squared	RMSE	\
PoissonRegressor	0.79	0.79	28664.06	
GradientBoostingRegressor	0.68	0.69	35210.52	
RandomForestRegressor	0.67	0.68	35880.73	
HistGradientBoostingRegressor	0.66	0.67	36347.19	
LGBMRegressor	0.64	0.64	37843.00	
BaggingRegressor	0.62	0.63	38616.00	
AdaBoostRegressor	0.62	0.63	38634.73	
XGBRegressor	0.55	0.56	42090.64	
GammaRegressor	0.54	0.55	42381.71	
LassoLarsIC	0.52	0.52	43604.87	
LassoLarsCV	0.52	0.52	43605.25	
LarsCV	0.52	0.52	43605.28	
LassoCV	0.52	0.52	43605.80	
Lasso	0.52	0.52	43606.01	
LassoLars	0.52	0.52	43606.02	
TransformedTargetRegressor	0.52	0.52	43606.05	
LinearRegression	0.52	0.52	43606.05	
Ridge	0.52	0.52	43607.47	
Lars	0.52	0.52	43610.93	
RidgeCV	0.52	0.52	43620.32	
BayesianRidge	0.52	0.52	43674.06	
SGDRegressor	0.51	0.52	43697.38	
OrthogonalMatchingPursuitCV	0.51	0.52	43943.77	
ElasticNet	0.47	0.48	45517.64	
KNeighborsRegressor	0.45	0.46	46593.13	
TweedieRegressor	0.44	0.45	46890.23	
DecisionTreeRegressor	0.40	0.40	48719.24	
ExtraTreesRegressor	0.38	0.39	49243.16	
HuberRegressor	0.36	0.37	50114.02	
PassiveAggressiveRegressor	0.34	0.35	50970.66	
ExtraTreeRegressor	0.32	0.33	51665.41	
OrthogonalMatchingPursuit	0.30	0.31	52344.68	
KernelRidge	0.22	0.23	55289.87	
ElasticNetCV	0.06	0.07	60886.67	
RANSACRegressor	0.05	0.06	61287.57	
DummyRegressor	-0.01	-0.00	63134.82	
MLPRegressor	-0.03	-0.02	63755.27	
NuSVR	-0.05	-0.04	64310.21	
SVR	-0.10	-0.08	65658.21	
QuantileRegressor	-0.10	-0.08	65746.02	
LinearSVR	-0.25	-0.23	70051.32	
GaussianProcessRegressor	-43292.31	-42744.87	13052148.71	

Model	Time Taken
PoissonRegressor	0.01
GradientBoostingRegressor	0.15
RandomForestRegressor	0.43
HistGradientBoostingRegressor	0.26
LGBMRegressor	0.04
BaggingRegressor	0.04
AdaBoostRegressor	0.07
XGBRegressor	0.27
GammaRegressor	0.02
LassoLarsIC	0.01
LassoLarsCV	0.02
LarsCV	0.01
LassoCV	0.03
Lasso	0.05
LassoLars	0.01
TransformedTargetRegressor	0.00
LinearRegression	0.01
Ridge	0.00
Lars	0.01
RidgeCV	0.00
BayesianRidge	0.01
SGDRegressor	0.00
OrthogonalMatchingPursuitCV	0.01
ElasticNet	0.00
KNeighborsRegressor	0.01
TweedieRegressor	0.00
DecisionTreeRegressor	0.01
ExtraTreesRegressor	0.36
HuberRegressor	0.05
PassiveAggressiveRegressor	0.01
ExtraTreeRegressor	0.01
OrthogonalMatchingPursuit	0.00
KernelRidge	0.33
ElasticNetCV	0.04
RANSACRegressor	0.05
DummyRegressor	0.00
MLPRegressor	1.22
NuSVR	0.26
SVR	0.35
QuantileRegressor	0.31
LinearSVR	0.00
GaussianProcessRegressor	0.88

```
[21]: # -----
# Ensure all data is numeric
# -----
print("Final Data Types of X_train:\n", X_train.dtypes)
```

Final Data Types of X\_train:

BHK	int64
Size	int64
Furnishing Status	int64
Bathroom	int64
floor_number	int64
Area Type_Carpet Area	int64
Area Type_Super Area	int64
City_Chennai	int64
City_Delhi	int64
City_Hyderabad	int64
City_Kolkata	int64
City_Mumbai	int64
	dtype: object

```
[22]: # -----
# Checking assumptions using statsmodels (for Multiple Linear Regression)
# -----
X_train_sm = sm.add_constant(X_train) # Add constant for intercept
model_sm = sm.OLS(y_train, X_train_sm).fit()

# Fit the model using statsmodels to get a detailed summary (including p-values)
print("== Statsmodels Regression Summary ==")
print(model_sm.summary())

# P-values for the coefficients
print("P-values for coefficients:")
print(model_sm.pvalues)
print()
```

==== Statsmodels Regression Summary ===

OLS Regression Results

---

Dep. Variable:	Rent	R-squared:	0.288
Model:	OLS	Adj. R-squared:	0.286
Method:	Least Squares	F-statistic:	127.7
Date:	Sun, 02 Mar 2025	Prob (F-statistic):	2.07e-268
Time:	23:57:11	Log-Likelihood:	-47664.
No. Observations:	3796	AIC:	9.535e+04
Df Residuals:	3783	BIC:	9.543e+04
Df Model:	12		
Covariance Type:	nonrobust		

---

	coef	std err	t	P> t	[0.025
0.975]					
const	-3.186e+04	4.88e+04	-0.653	0.514	-1.28e+05
6.38e+04					
BHK	2859.5472	2387.188	1.198	0.231	-1820.752
7539.847					
Size	37.2556	2.920	12.758	0.000	31.530
42.981					
Furnishing Status	2829.7975	1686.370	1.678	0.093	-476.485
6136.080					
Bathroom	1.021e+04	2510.059	4.069	0.000	5292.992
1.51e+04					
floor_number	1369.5544	243.030	5.635	0.000	893.071
1846.038					
Area_Type_Carpet Area	-5617.1468	4.87e+04	-0.115	0.908	-1.01e+05
8.99e+04					
Area_Type_Super Area	-1.058e+04	4.87e+04	-0.217	0.828	-1.06e+05
8.49e+04					
City_Chennai	-7493.7050	3660.268	-2.047	0.041	-1.47e+04
-317.416					
City_Delhi	1.093e+04	4133.658	2.643	0.008	2821.728
1.9e+04					
City_Hyderabad	-1.676e+04	3725.922	-4.498	0.000	-2.41e+04
-9454.794					
City_Kolkata	-353.8969	4411.010	-0.080	0.936	-9002.084
8294.290					
City_Mumbai	4.607e+04	4147.924	11.108	0.000	3.79e+04
5.42e+04					
Omnibus:	10361.874	Durbin-Watson:			2.013
Prob(Omnibus):	0.000	Jarque-Bera (JB):			397954117.987
Skew:	33.316	Prob(JB):			0.00
Kurtosis:	1587.804	Cond. No.			8.74e+04

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 8.74e+04. This might indicate that there are strong multicollinearity or other numerical problems.

P-values for coefficients:

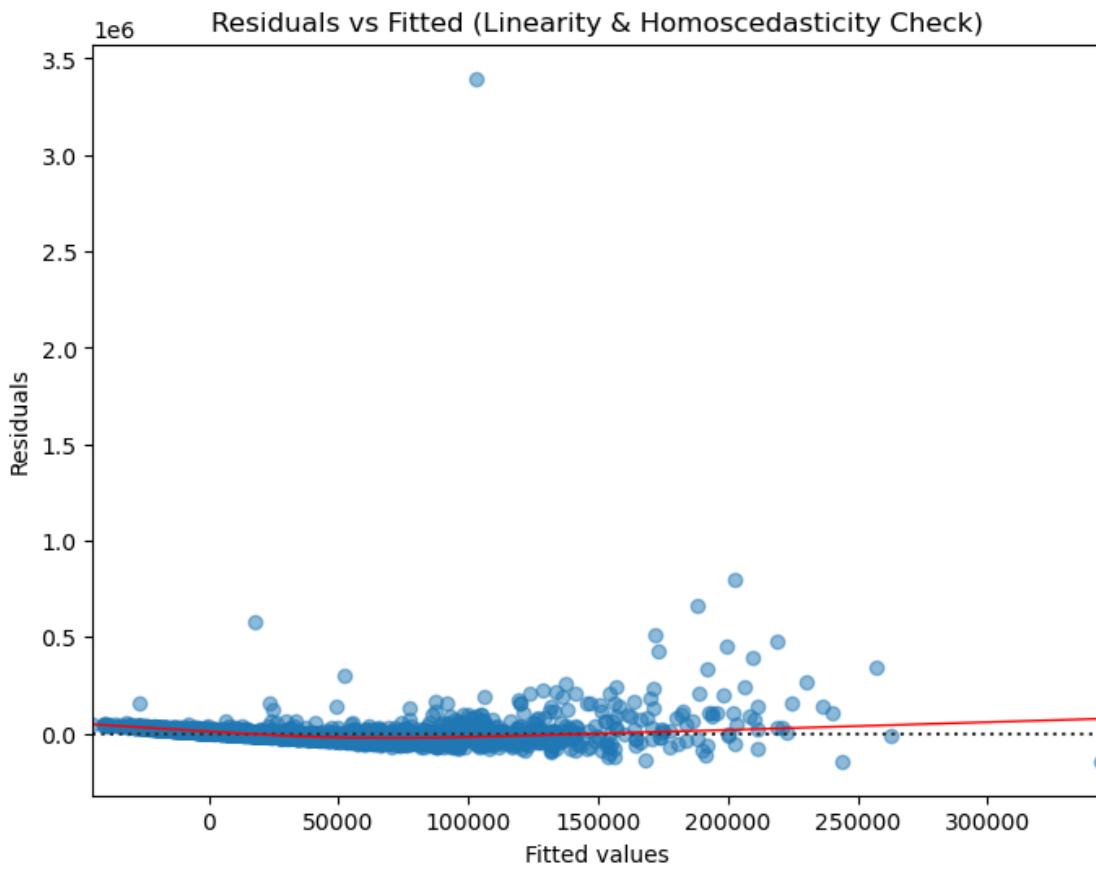
const	0.51
BHK	0.23
Size	0.00
Furnishing Status	0.09

```
Bathroom          0.00
floor_number      0.00
Area_Type_Carpet Area  0.91
Area_Type_Super Area  0.83
City_Chennai       0.04
City_Delhi         0.01
City_Hyderabad     0.00
City_Kolkata        0.94
City_Mumbai         0.00
dtype: float64
```

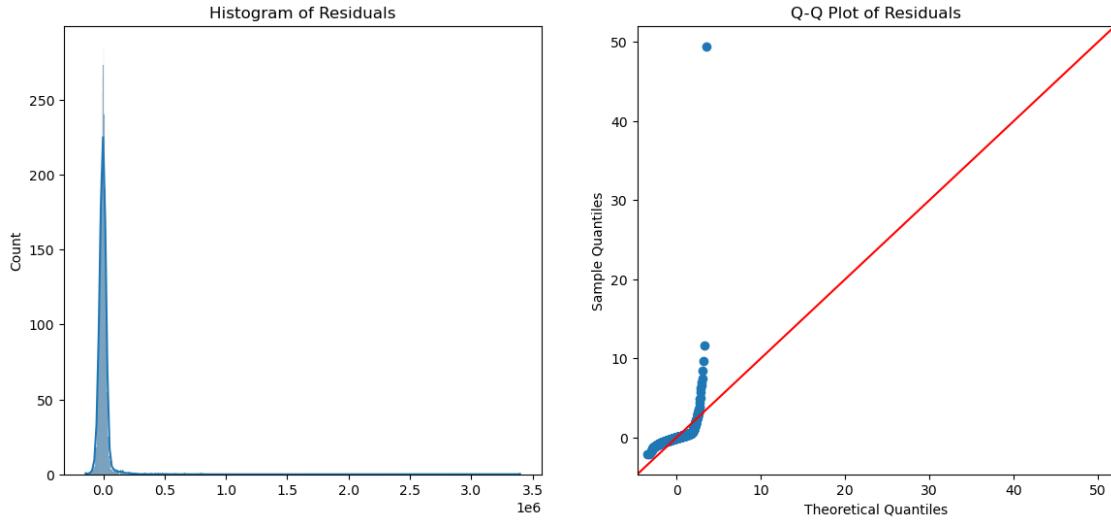
## Performance Analysis

```
[23]: # -----
# Diagnostic Plots and Assumptions Check
# -----
# Get residuals and fitted values from the statsmodels fit
residuals = model_sm.resid
fitted = model_sm.fittedvalues
```

```
[24]: # Linearity & Constant Variance (Homoscedasticity):
plt.figure(figsize=(8, 6))
sns.residplot(x=fitted, y=y_train, lowess=True,
               line_kws={'color': 'red', 'lw': 1}, scatter_kws={'alpha':0.5})
plt.xlabel("Fitted values")
plt.ylabel("Residuals")
plt.title("Residuals vs Fitted (Linearity & Homoscedasticity Check)")
plt.show()
```



```
[43]: # Normality of Errors: Histogram and Q-Q plot
plt.figure(figsize=(14, 6))
plt.subplot(1, 2, 1)
sns.histplot(residuals, kde=True)
plt.title("Histogram of Residuals")
ax=plt.subplot(1, 2, 2)
sm.qqplot(residuals, line='45', fit=True, ax=ax)
plt.title("Q-Q Plot of Residuals")
plt.show()
```



## Feature Dependence

```
[26]: # Independence of Errors: Durbin-Watson test statistic
dw_stat = sm.stats.stattools.durbin_watson(residuals)
print("Durbin-Watson statistic:", dw_stat)
# (A value around 2 suggests no autocorrelation.)  
  

# Multicollinearity Analysis: Compute Variance Inflation Factor (VIF)
vif_data = pd.DataFrame()
vif_data[\"Feature\"] = X_train.columns
vif_data[\"VIF\"] = [variance_inflation_factor(X_train.values, i) for i in range(X_train.shape[1])]
print("\n==== Variance Inflation Factor (VIF) ===")
print(vif_data)
```

Durbin-Watson statistic: 2.0126348081090955

```
==== Variance Inflation Factor (VIF) ===
      Feature    VIF
0          BHK  23.18
1          Size   9.15
2 Furnishing Status   2.37
3          Bathroom 23.54
4      floor_number  1.99
5 Area Type_Carpet Area   7.17
6 Area Type_Super Area   6.13
7          City_Chennai  2.02
8          City_Delhi    1.76
9          City_Hyderabad 2.00
```

```

10          City_Kolkata  1.67
11          City_Mumbai   2.81

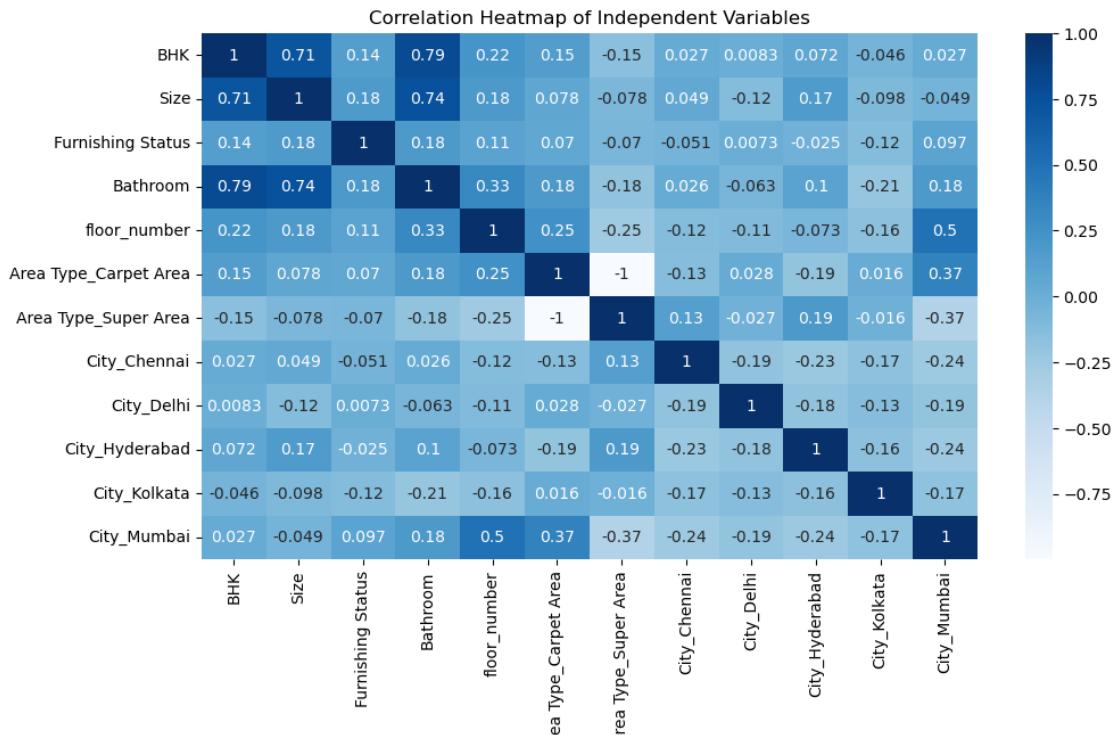
```

[27]: # Also, a correlation heatmap among independent variables:

```

plt.figure(figsize=(11, 6))
sns.heatmap(X_train.corr(), annot=True, cmap="Blues")
plt.title("Correlation Heatmap of Independent Variables")
plt.show()

```

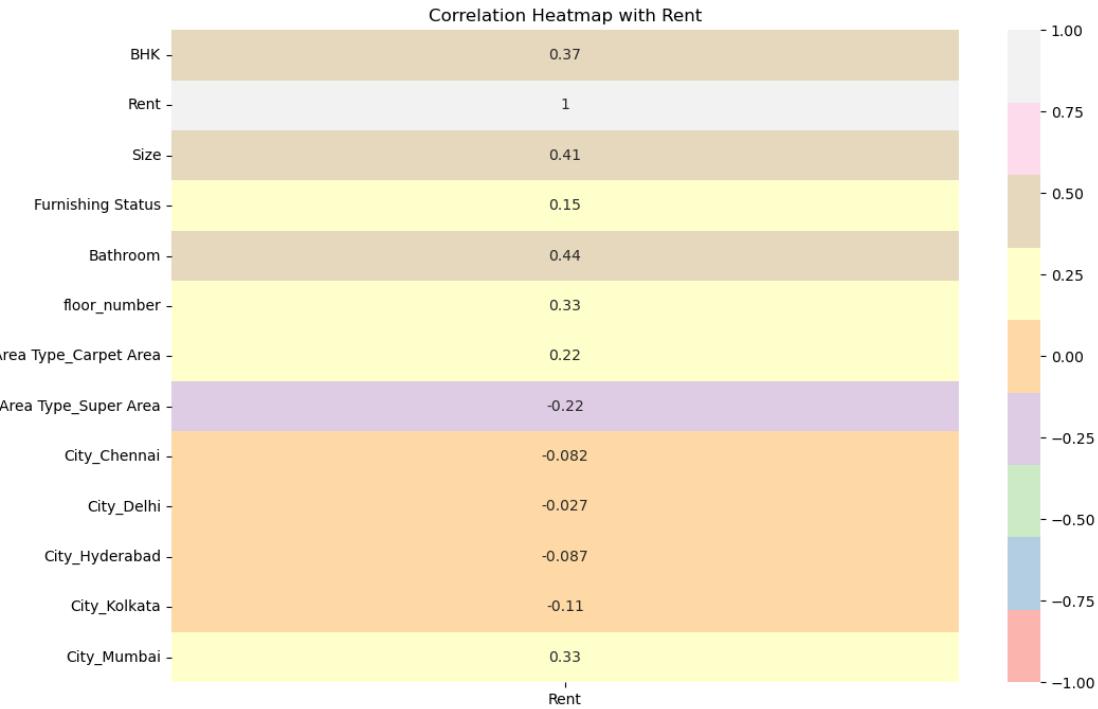


[28]: #HeatMap (Correlation Heatmap with Rent)

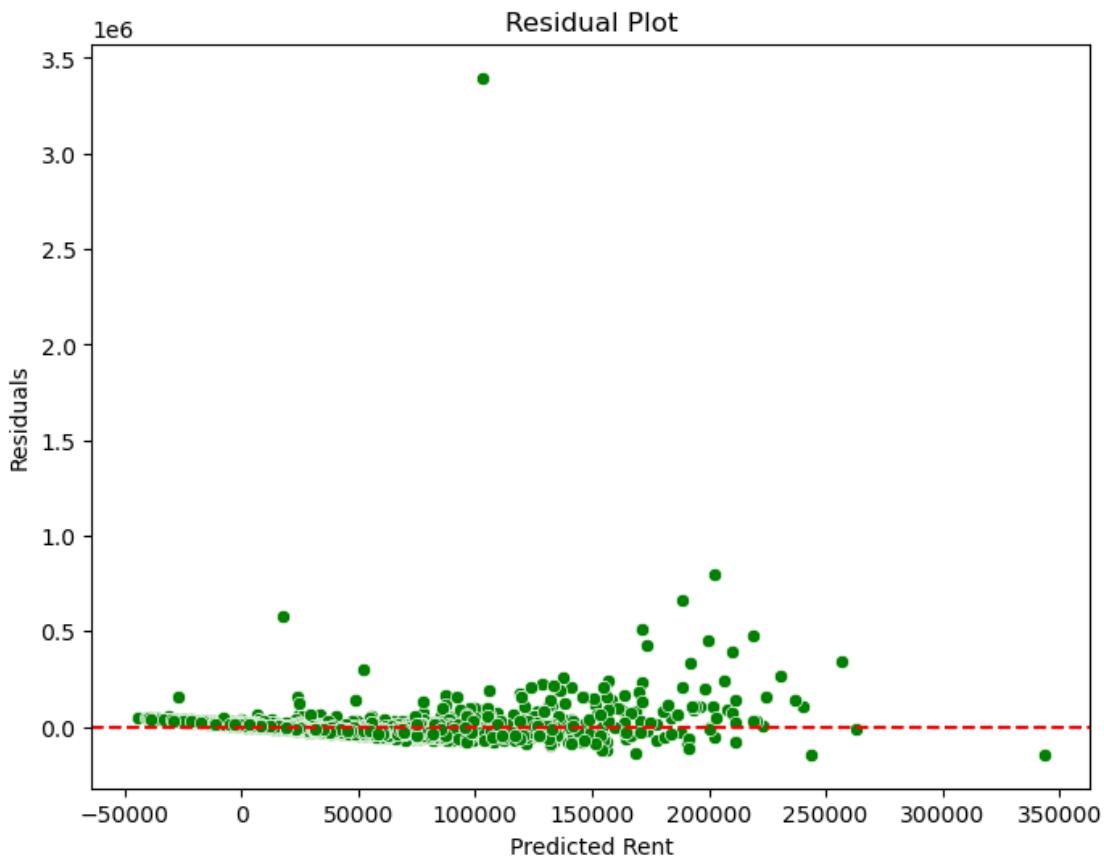
```

plt.figure(figsize=(12, 8))
corr = df_cleaned.corr()
sns.heatmap(corr[['Rent']], annot=True, cmap='Pastel1', vmin=-1, vmax=1)
plt.title('Correlation Heatmap with Rent')
plt.show()

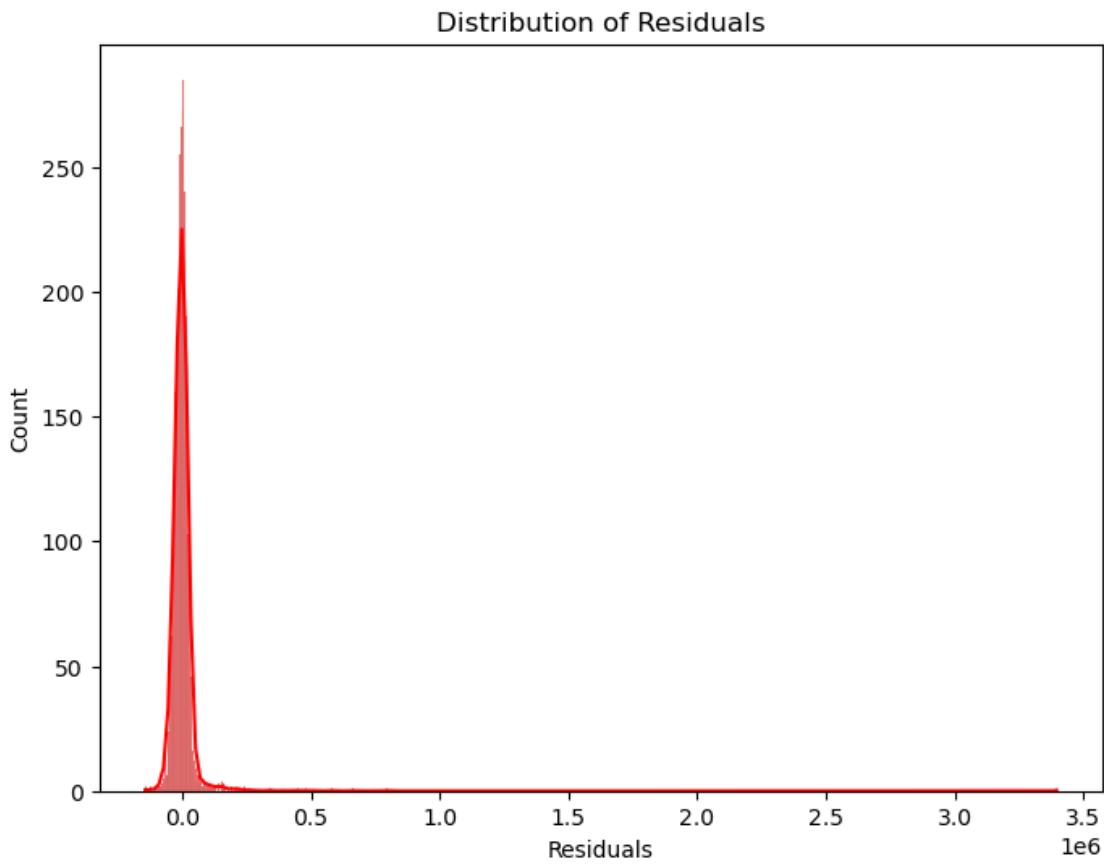
```



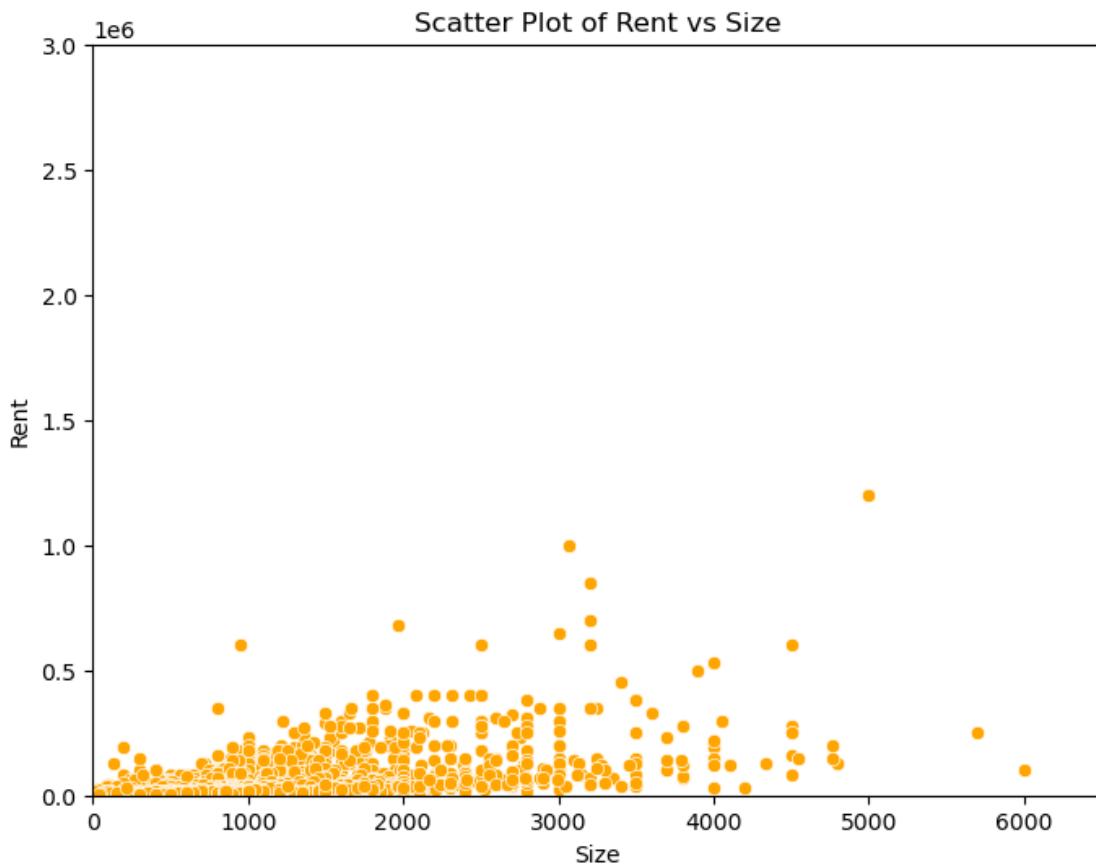
```
[95]: #ResidualPlot
residuals = y_train - y_pred
plt.figure(figsize=(8, 6))
sns.scatterplot(x=y_pred, y=residuals, color='green')
plt.axhline(y=0, color='red', linestyle='--')
plt.xlabel('Predicted Rent')
plt.ylabel('Residuals')
plt.title('Residual Plot')
plt.show()
```



```
[97]: plt.figure(figsize=(8, 6))
sns.histplot(residuals, kde=True, color='red')
plt.xlabel('Residuals')
plt.title('Distribution of Residuals')
plt.show()
```



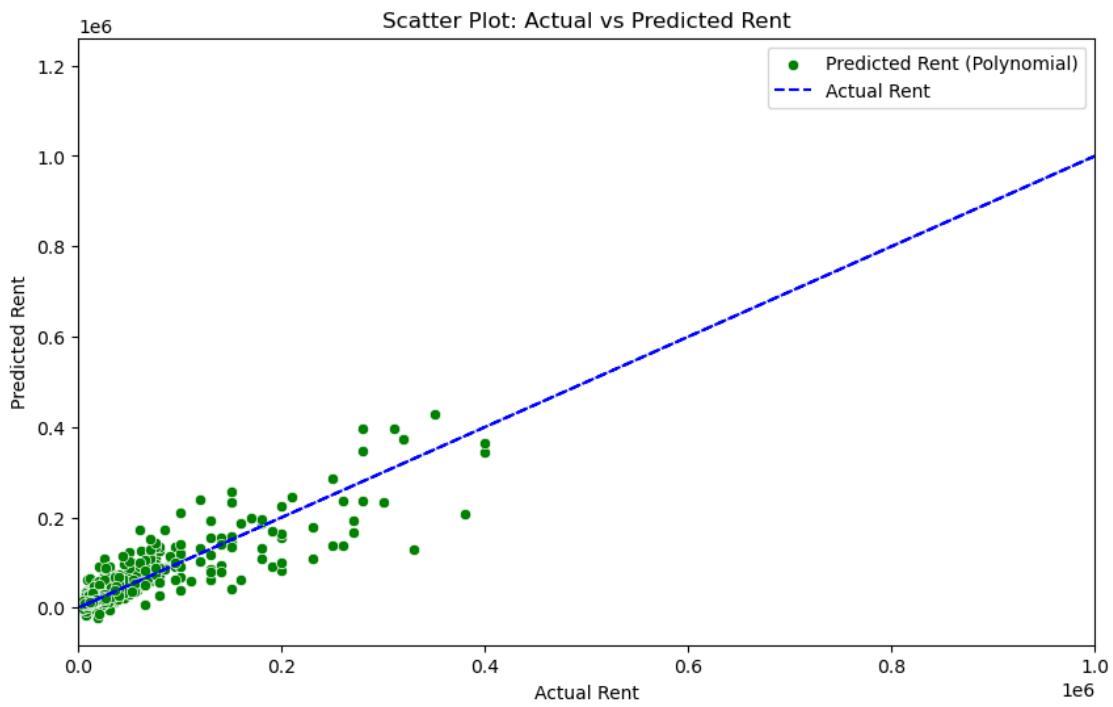
```
[39]: #scatter plot for rent vs size
plt.figure(figsize=(8, 6))
sns.scatterplot(x='Size', y='Rent', data=df, color='orange')
plt.xlim(0,6500)
plt.ylim(0,3000000)
plt.title('Scatter Plot of Rent vs Size')
plt.xlabel('Size')
plt.ylabel('Rent')
plt.show()
```



```
[41]: plt.figure(figsize=(10, 6))

# Scatter plot for Polynomial Regression
sns.scatterplot(x=y_test, y=y_poly_pred, color="green", label="Predicted Rent ↴(Polynomial)")

# Perfect Fit Line (y = x)
plt.plot(y_test, y_test, color="blue", linestyle="dashed", label="Actual Rent ")
plt.xlim(0,1000000)
# Labels and Title
plt.xlabel("Actual Rent")
plt.ylabel("Predicted Rent")
plt.title("Scatter Plot: Actual vs Predicted Rent")
plt.legend()
plt.show()
```



[ ]:

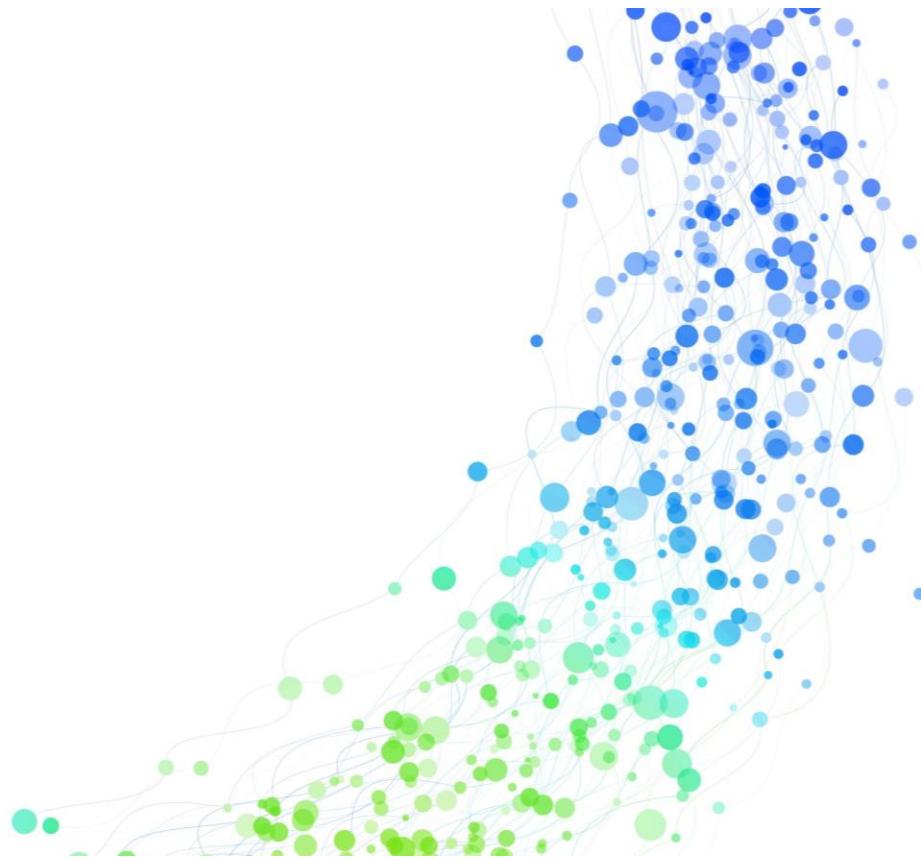
[ ]:

# Final Project

## SCM 516

by **Synergy Sparks**

- Akanksha Manohar
- Meenakshi Rajeev Nair
- Neha Tiwari
- Niketha Venkatesh
- Shriya Karthikeyan



# **Link of our Presentation:**

- [SCM-516\\_Final\\_Project](#)

# Objective

---

To harness predictive modeling to extract meaningful insights from data using classification and regression algorithms.

---

**Classification dataset:** Naïve Bayes, Decision Tree, Random Forest, KNN, feature importance and conditional probability.

---

**Regression dataset:** linear and polynomial regression models to predict the values of a response (target) variable.

---

By the end, we aim to identify the most effective model and assess its generalization ability.



# About Classification Dataset

Customer Churn Prediction Dataset:

## Dataset Overview:

- We will be predicting customer churn where definition of the target variable is:
- Churn: 1 = Customer left,
- 0 = Customer stayed.

## Key Features:

- Demographics: Age of the customer.
- Subscription Details: Subscription length, contract type (one-year, two-year, month-to-month).
- Usage Behavior: Total data usage (GB), number of support calls.
- Billing & Payment: Payment method, additional services subscribed.

## Objective:

- Use machine learning for binary classification to predict churn.
- Help businesses identify at-risk customers and reduce churn.



# Preprocessing steps:

- **Dataset Overview:** Collected real-world data with multiple features for classification.
- **Handling Missing Data:** Implemented imputation strategies to ensure data completeness.
- **Label Encoding:** Converted categorical labels into numerical values.
- **Outlier Removal:** Used boxplot analysis to detect and remove outliers.
- **SMOTE Analysis:** Applied SMOTE to balance the dataset.
- **Data Normalization:** Standardized feature values to enhance model efficiency.
- **Splitting Strategy:** Divided data into training and testing sets for evaluation.



# Models

Naïve Bayes

Decision  
Trees

Random  
Forest

KNN

# Goodness of Models

Confusion  
Matrix

Accuracy

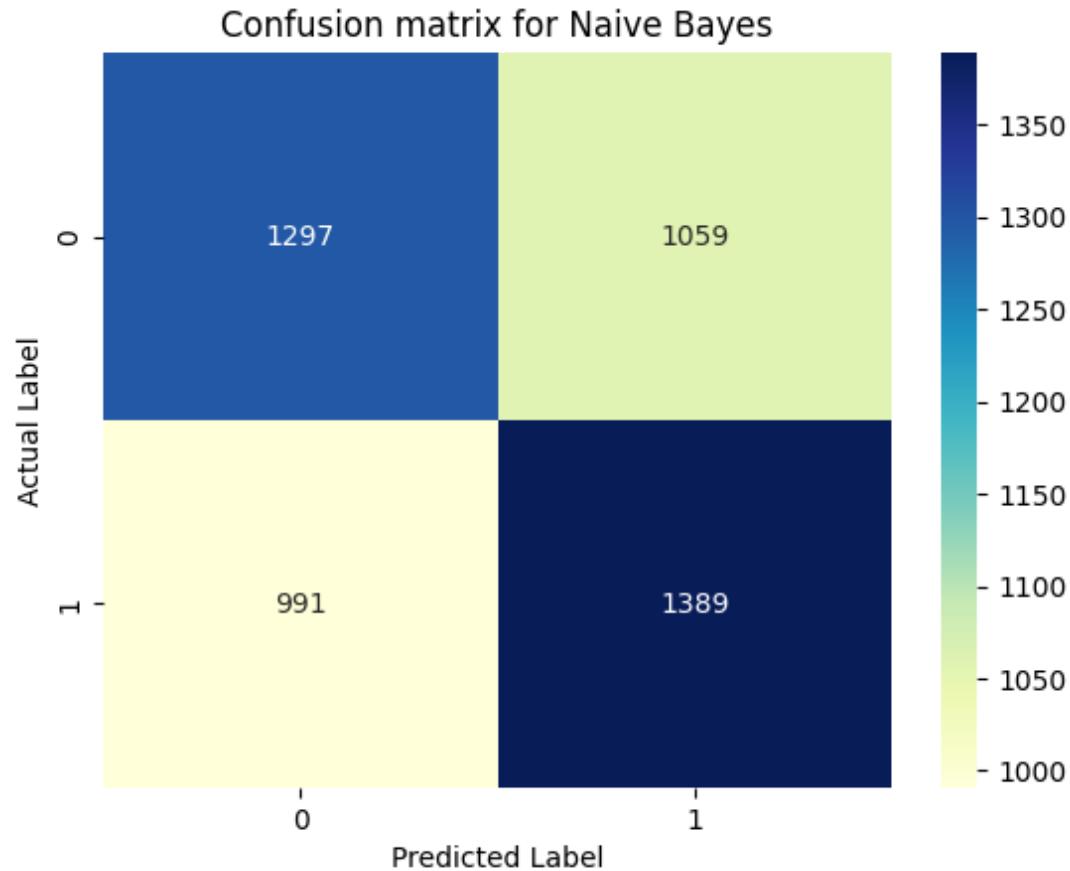
Precision

Recall

F1  
Score

ROC  
Curve

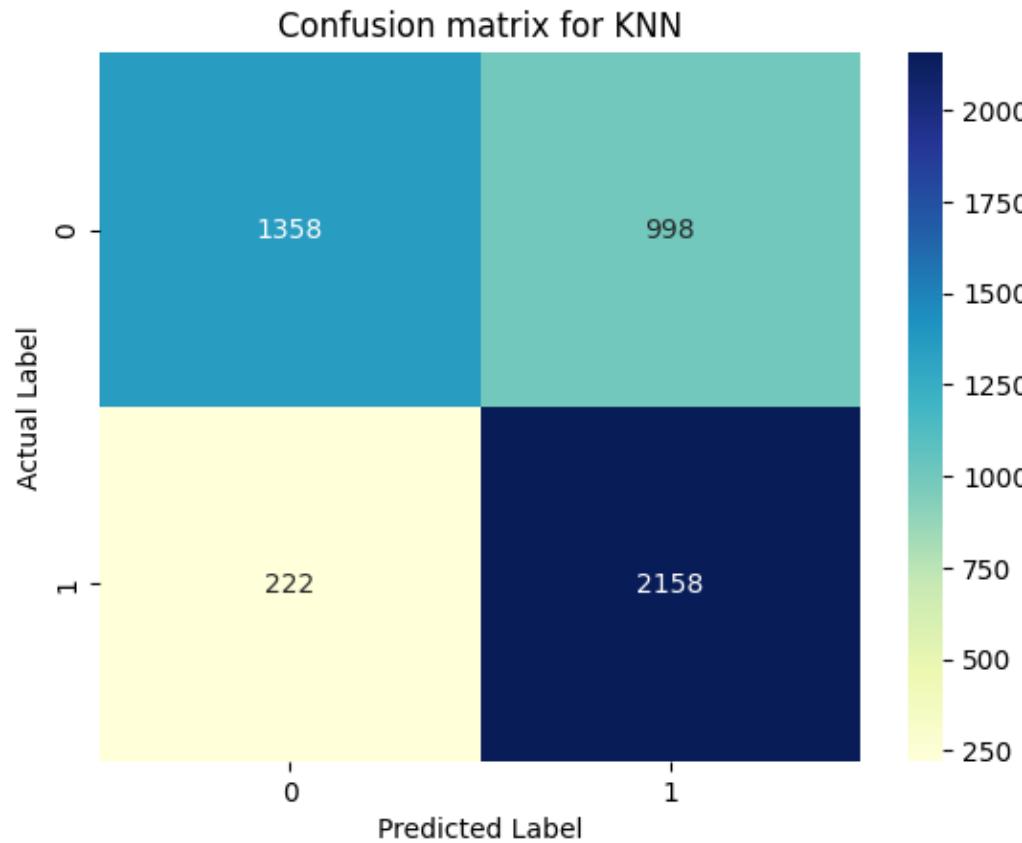
# Confusion Matrix for Naïve Bayes



For Naïve Bayes:

**Accuracy:** 63%  
**Precision:** 62%  
**Recall:** 68%  
**F1 Score:** 65%

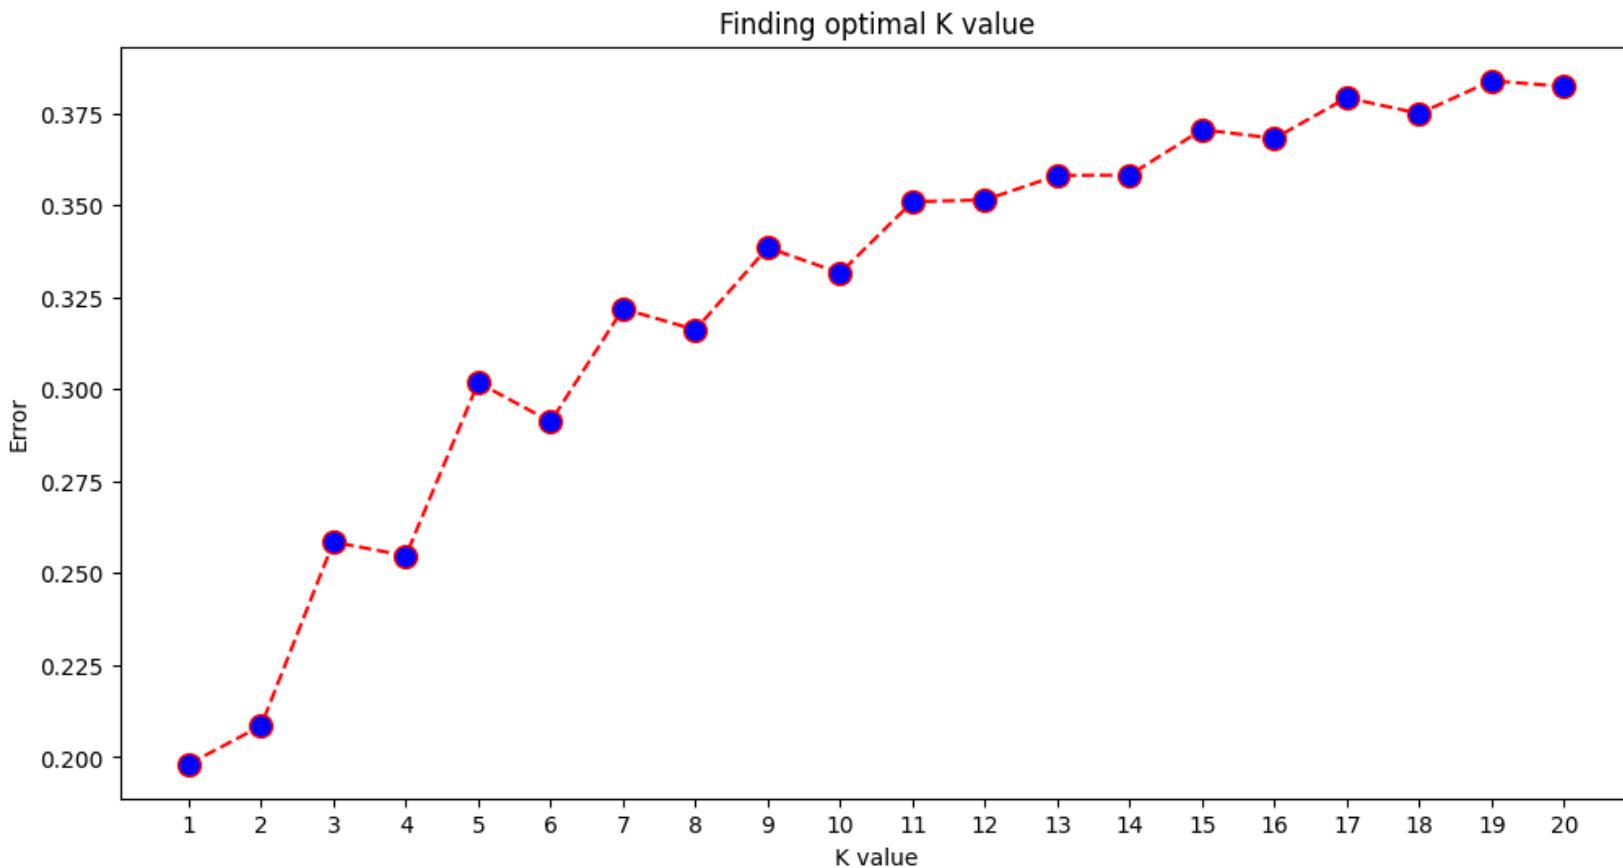
# Confusion Matrix for KNN



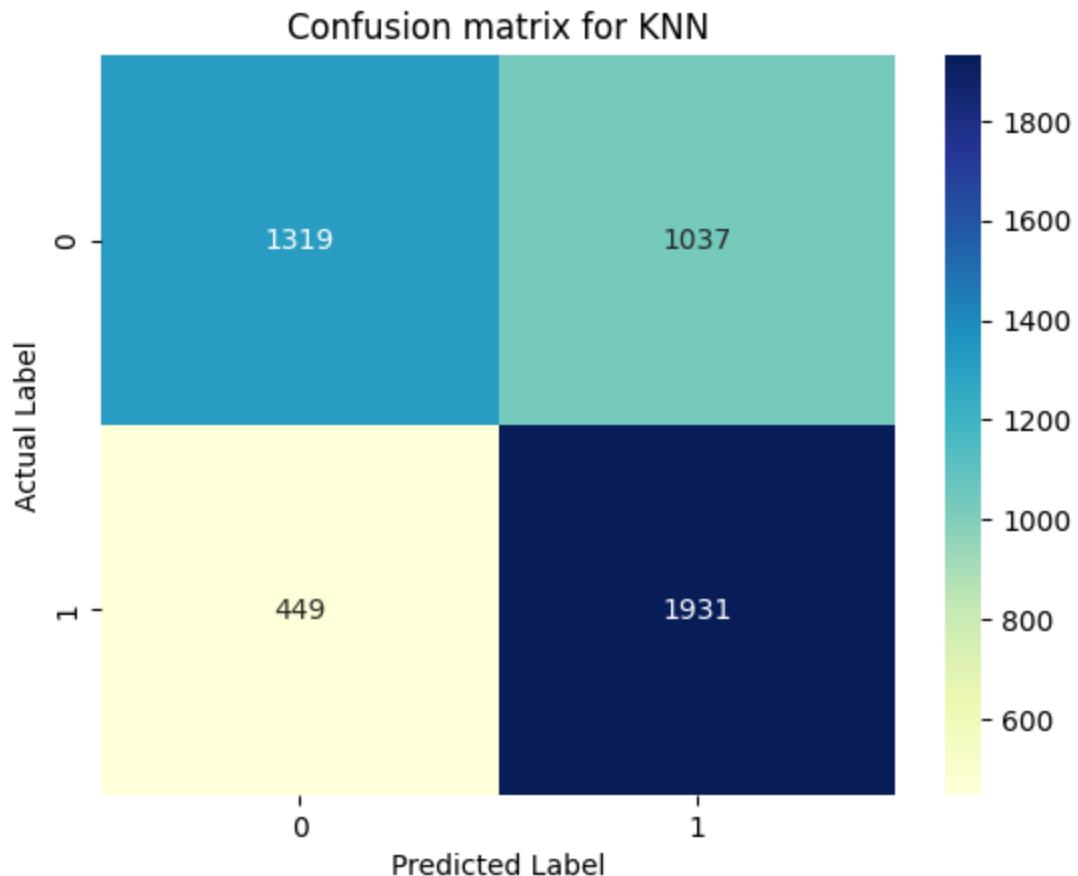
For KNN:

**Accuracy:** 68%  
**Precision:** 65%  
**Recall:** 78%  
**F1 Score:** 71%

# Finding the Optimal K:



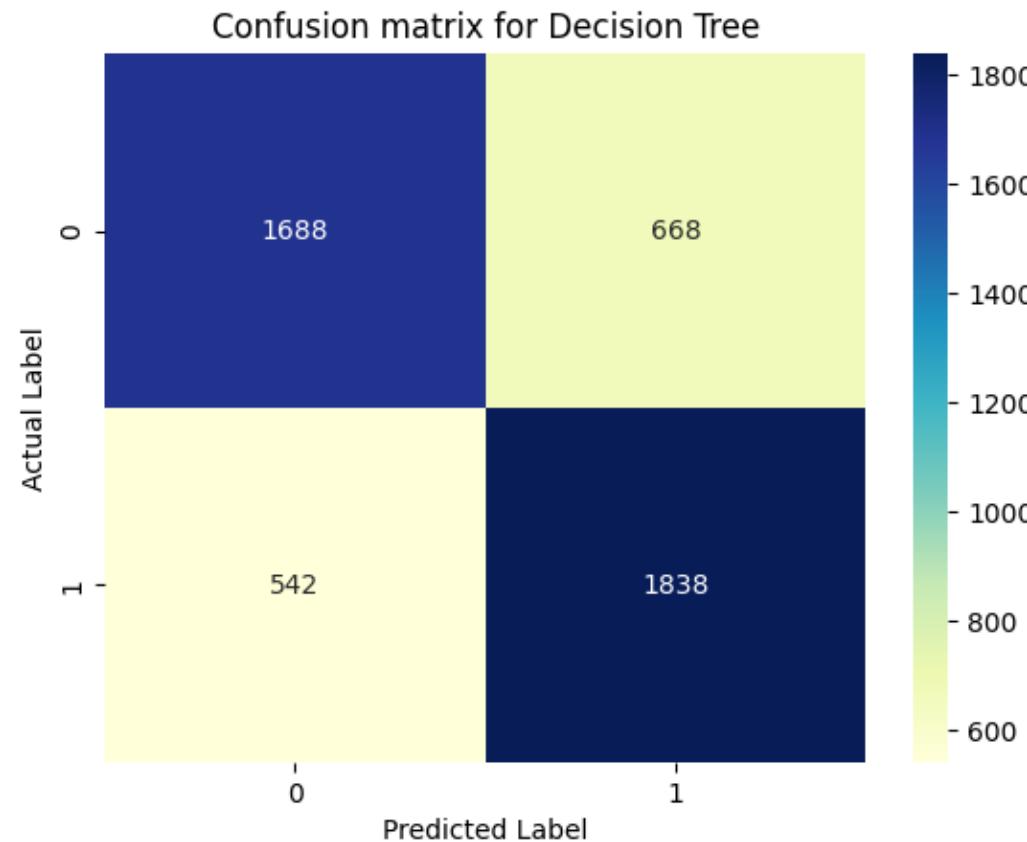
# Confusion Matrix for KNN for optimal k=1



For KNN:

**Accuracy:** 69%  
**Precision:** 65%  
**Recall:** 81%  
**F1 Score:** 72%

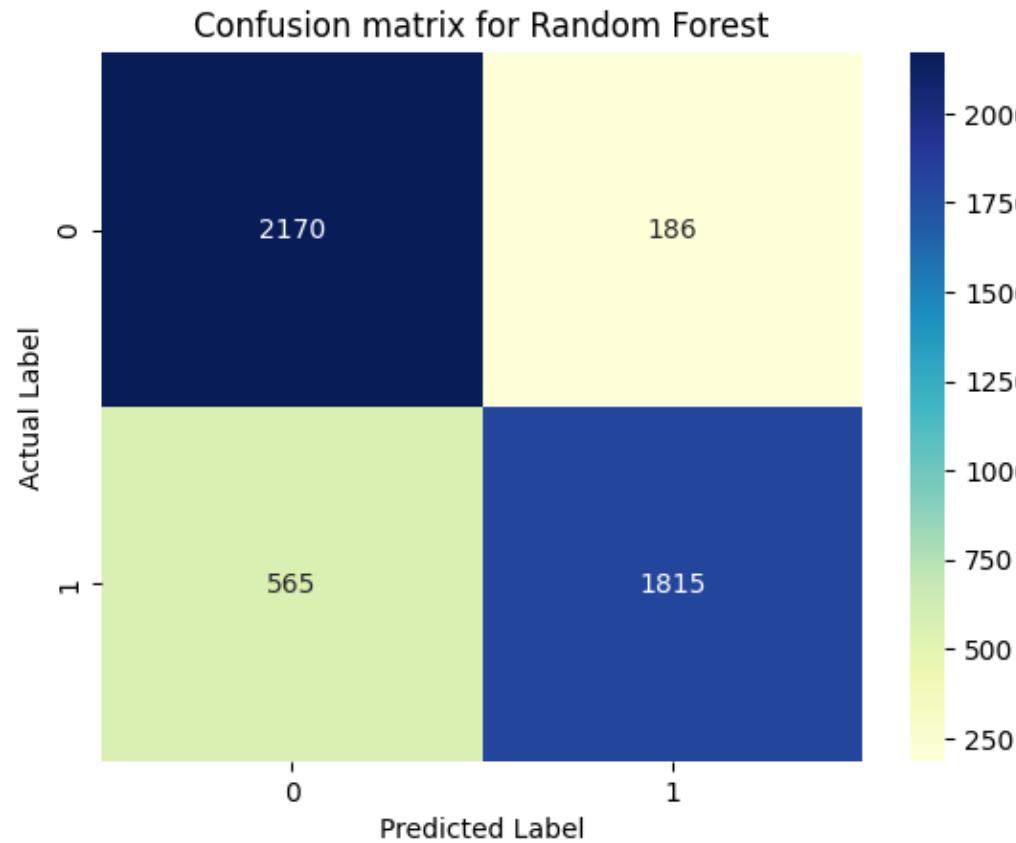
# Confusion Matrix for Decision Tree



For Decision Tree:

**Accuracy:** 76%  
**Precision:** 75%  
**Recall:** 77%  
**F1 Score:** 76%

# Confusion Matrix for Random Forest

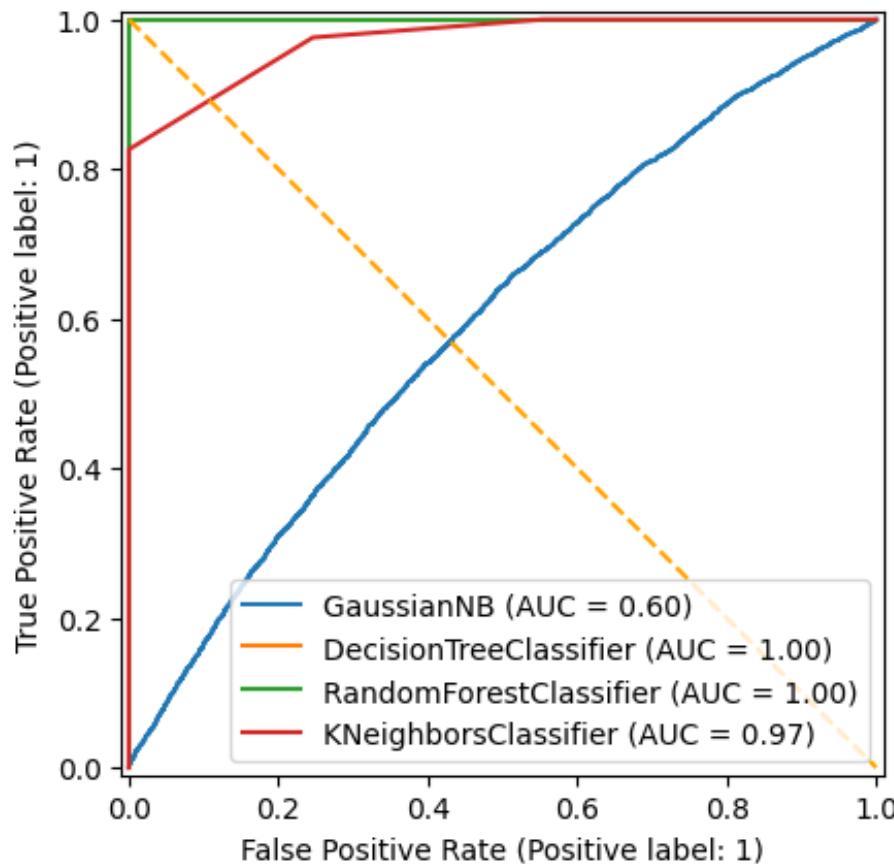


For Random Forest:

**Accuracy:** 85%  
**Precision:** 97%  
**Recall:** 72%  
**F1 Score:** 83%

# ROC Curve Analysis

(for training dataset):



On TRAINING Data-

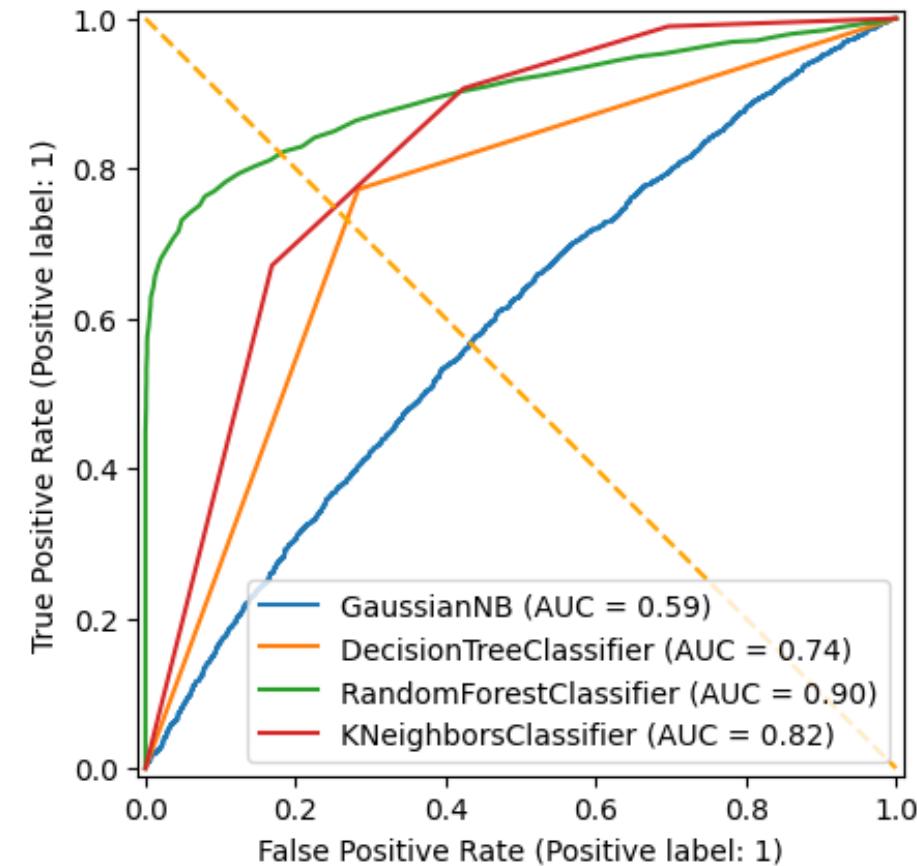
- 1. Gaussian NB :** Its curve is close to the diagonal line, indicating it's barely better than random guessing. The AUC is also relatively low.
- 2. K neighbors Classifier :** This model performs very well, with its curve close to the top-left corner.
- 3. Decision Tree Classifier :** This model has perfect performance. However, a perfect AUC often suggests potential overfitting.
- 4. Random Forest Classifier :** Similar to the Decision Tree, the same caution about potential overfitting applies.

# ROC Curve Analysis

(for testing dataset):

On TESTING DATA-

1. **GaussianNB (AUC = 0.59)**: Very poor. The curve is close to the diagonal line.
2. **Kneighbors Classifier (AUC = 0.82)**: Reasonably good performance. The curve is above the diagonal line. This model also performs well on the testing data, though not as strongly as Random Forest.
3. **Decision Tree Classifier (AUC = 0.74)**: Moderate performance. The curve is above the diagonal line, but not significantly.
4. **Random Forest Classifier (AUC = 0.90)**: Good performance. The curve is well above the diagonal line, indicating a strong ability to distinguish between classes. **This model generalizes well to the testing data and is a good choice for this classification task.**



# Performance Comparison

## Evaluating Model Effectiveness

- **Accuracy:** Random Forest achieved the highest accuracy, outperforming other models.
- **Precision & Recall:** Trade-off observed: Decision Trees had high precision but lower recall, while KNN struggled with both.
- **F1-Score:** Random Forest maintained the best balance.
- **Class Imbalance:** Some models were impacted by imbalanced data, affecting recall and precision.

**Overall :**

**Random Forest** demonstrated the best overall performance across all metrics.

# Classification report of Random Forest

	Precision	Recall	F1- score	support
Class 0	0.79	0.92	0.85	2356
Class 1	0.91	0.76	0.83	2380
Accuracy			0.84	4736
Macro avg	0.85	0.84	0.84	4736
Weighted avg	0.85	0.84	0.84	4736

**Conclusion: The model has an overall accuracy of 84% (0.84).**

This means it correctly predicted the class labels for 84% of the samples in the test set.

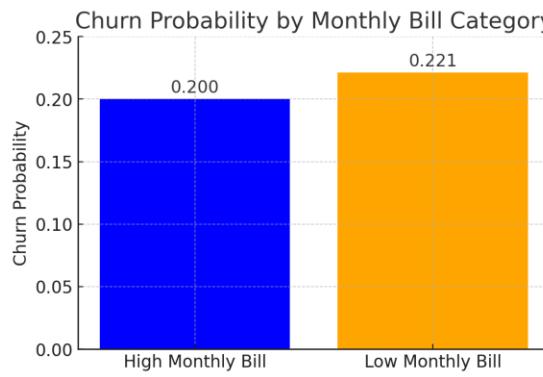
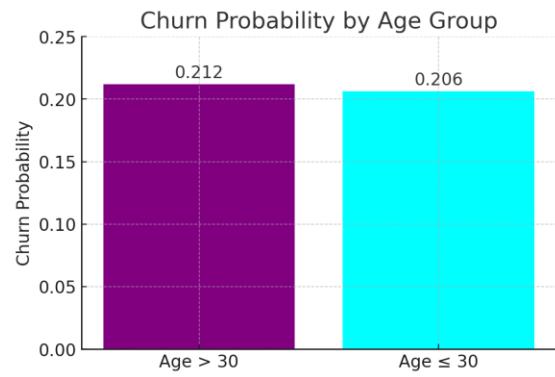
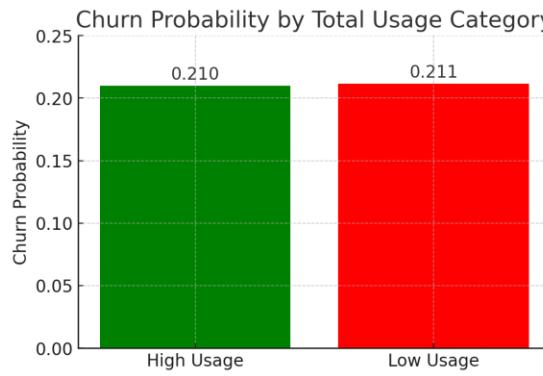
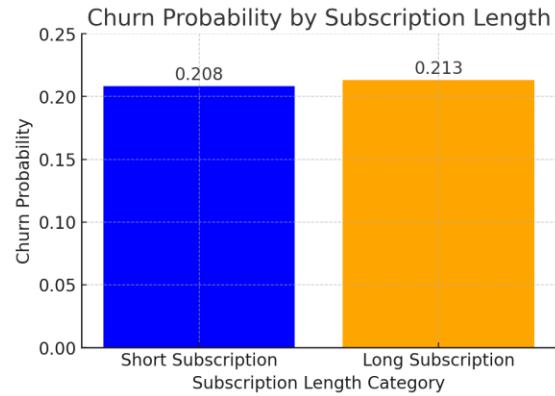
**Let's see how  
well our model  
performs on  
unseen data?**



# New dataset:

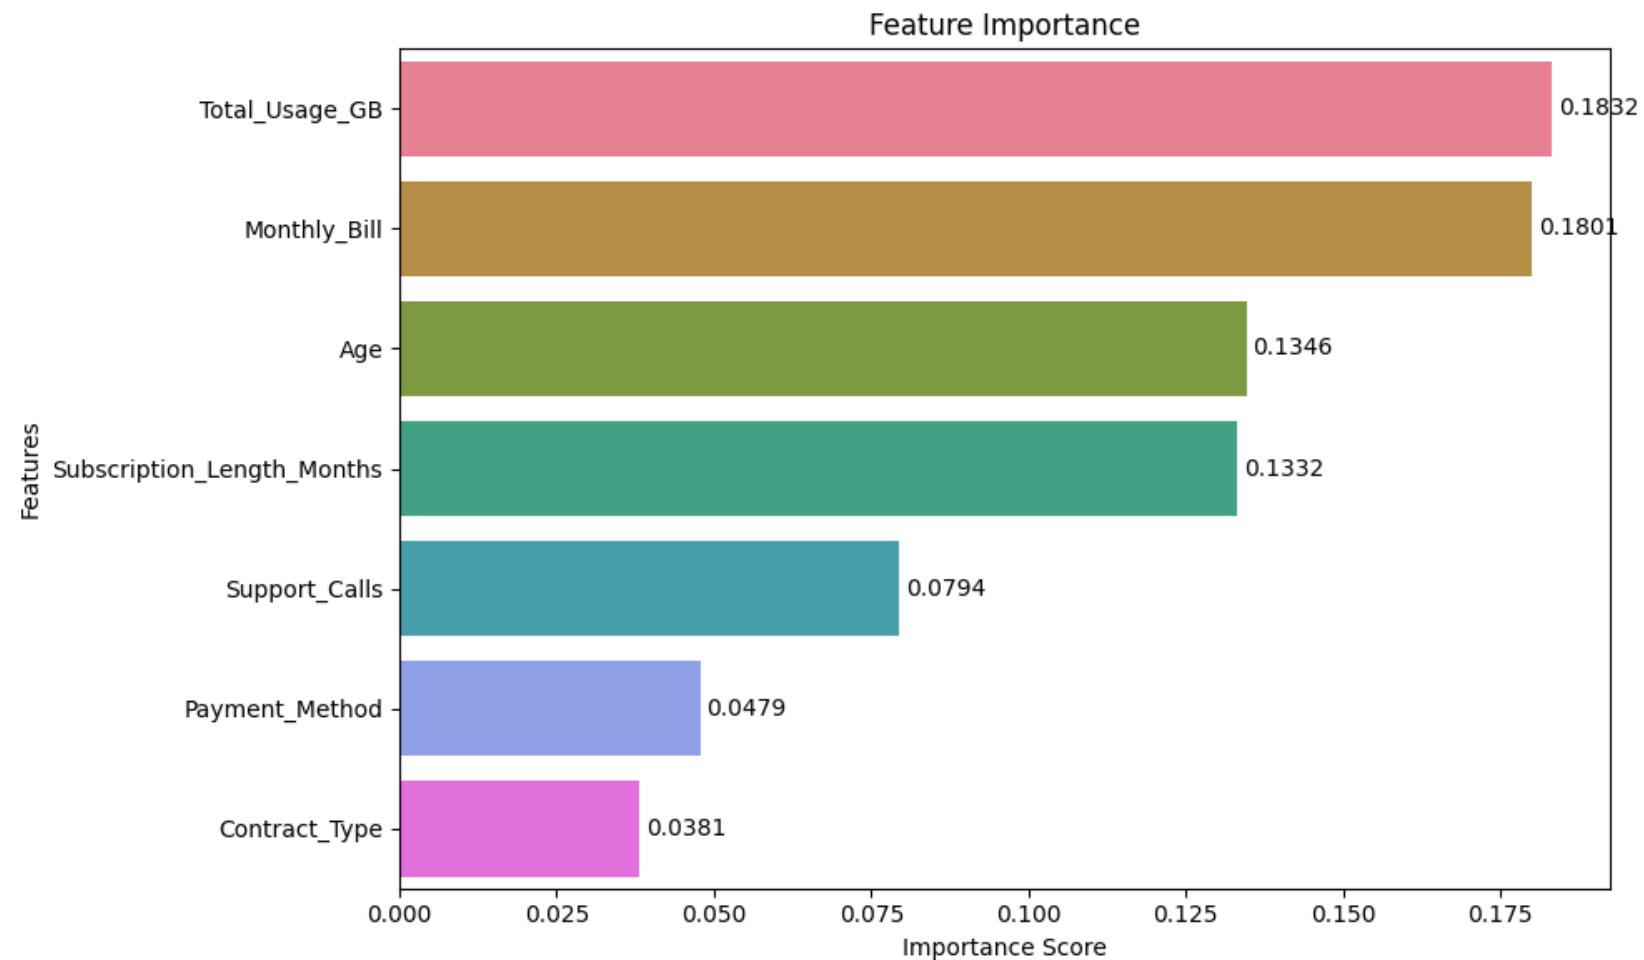
	Customer ID	Age	Subscription Length (months)	Monthly Bill (\$)	Total Usage (GB)	Support Calls	Contract Type	Payment Method	Has Additional Services
1	10001	43.296	33.66	107.8	230.95	1	Month-to-Month	Debit Card	1
2	10002	55.704	48.45	144.6	398.05	9	One Year	Credit Card	1
3	10003	35.928	27.08	71.65	131.65	1	Month-to-Month	PayPal	0

# Feature Importance using Naïve Bayes



# Importance of Feature

(determined using Random Forest):



# Which model performed best?

- **Random Forest** emerged as the **top-performing model** due to its ensemble learning approach. Unlike a single decision tree, it aggregates predictions from multiple trees, significantly improving accuracy.
- One of its key strengths is robustness—Random Forest is less sensitive to noisy data and missing values compared to other classifiers.
- It also provides feature importance rankings, helping us understand which variables influence predictions the most.
- Moreover, it strikes an excellent balance between bias and variance, avoiding overfitting while maintaining high predictive power.
- Our results confirm that **Random Forest** consistently outperformed other models across all key performance metrics.

# New dataset with churn values:

	Customer ID	Age	Subscription Length (months)	Monthly Bill (\$)	Total Usage (GB)	Support Calls	Contract Type	Payment Method	Has Additional Services	No Churn	Churn = 1
1	10001	43.296	33.66	107.8	230.95	1	Month-to-Month	Debit Card	1	62%	38%
2	10002	55.704	48.45	144.6	398.05	9	One Year	Credit Card	1	72%	28%
3	10003	35.928	27.08	71.65	131.65	1	Month-to-Month	PayPal	0	61%	39%

# Are there any signs of overfitting?

While machine learning models offer powerful predictive capabilities, they come with challenges and limitations.

- **Data quality** —no model can outperform the quality of its input data.
- **Overfitting** – where a model performs exceptionally well on training data but struggles on new data.
- **Computational cost** –models like Random Forest require significant processing power, making them less feasible for real-time applications.
- **Hyperparameter tuning** – finding the right balance between bias and variance.
- **Interpretability** – making it difficult to explain predictions, especially in fields requiring transparency, such as healthcare and finance.

# **Key Insights from Conditional Probability Analysis of Customer Churn:**

## **Churn Risk Factors:**

- Customers with **high monthly bills (\$90+)**, **heavy data usage (300GB+)**, or **multiple support calls (5+ calls)** → (~20%)
- Month-to-Month contract customers (especially heavy users) are at **higher churn risk (33%)**
- Customers with **long-term contracts (2+ years, 36+ months)** have slightly lower churn (~21%)
- **Older customers (60+)** with **low usage (<100GB)** have the lowest churn probability (~18%)

# Regression Analysis:



# About Regression Dataset

House Rent Prediction Dataset:2

## Dataset Overview:

- Target variable is rent.
- Aims to predict rent prices based on other independent variables.

## Key Features:

- Includes numerical attributes like BHK, size, bathroom.
- Includes categorical attributes like floor level, city, furnishing status and tenant preferences.

## Objective:

- Apply Multiple Linear and Polynomial Regression to predict rental prices based on relevant features.
- Analyze the impact of different property attributes on rent to improve decision-making for renters and property owners.

# Independent variables:

	BHK	Size	Furnishing status	Bathroom	Floor number	Carpet Area	Super Area	Chennai	Delhi	Hyderabad	Kolkata	Mumbai
0	2	1100	0	2	0	False	True	False	False	False	True	False
1	2	800	1	1	1	False	True	False	False	False	True	False
2	2	1000	1	1	1	False	True	False	False	False	True	False
3	2	800	0	1	1	False	True	False	False	False	True	False
4	2	850	0	1	1	True	False	False	False	False	True	False

# Preprocessing steps:

- **Dataset Overview:** Real-world housing rental data with features like rent, size, location, and furnishing status.
- **Handling Missing Data:** Used imputation techniques to ensure data completeness.
- **Categorical to Numerical Conversion:** Transformed categorical features for model compatibility.
- **One-Hot Encoding:** Applied one-hot Encoding to multi-category variables to retain information.
- **Boolean Conversion:** Converted Boolean columns to integers (0/1).
- **Splitting Strategy:** Split data into training and testing sets for evaluation.

# Equation of the Model:

```
== Equation of the Model ==
Rent = 236693.67 + -40788374731.34 * BHK + 81575716.39 * Size + 174315.13 * Furnishing Status + 157807.45 * Bathroom + 26404.82 * floor_number
+ -112617.09 * Area Type_Carpet Area + -114163.93 * Area Type_Super Area + 52728.54 * City_Chennai + -4197.02 * City_Delhi + -4768.28 * City_Hyderabad
+ 509.40 * City_Kolkata + -8347.23 * City_Mumbai + -382.45 * BHK^2 + -4.24 * BHK Size + -1763.72 * BHK Furnishing Status + 6712.43 * BHK Bathroom
+ 891.13 * BHK floor_number + 40788361229.53 * BHK Area Type_Carpet Area + 40788362939.56 * BHK Area Type_Super Area + 8132.09 * BHK City_Chennai
+ 3667.74 * BHK City_Delhi + 7881.57 * BHK City_Hyderabad + 10132.94 * BHK City_Kolkata + -20261.69 * BHK City_Mumbai + 0.00 * Size^2 + 7.13 * Size Furnishing Status
+ 4.03 * Size Bathroom + -1.07 * Size floor_number + -81575669.32 * Size Area Type_Carpet Area + -81575685.43 * Size Area Type_Super Area
+ -33.46 * Size City_Chennai + -34.51 * Size City_Delhi + -38.79 * Size City_Hyderabad + -42.21 * Size City_Kolkata + 109.30 * Size City_Mumbai
+ -682.73 * Furnishing Status^2 + -736.49 * Furnishing Status Bathroom + 648.78 * Furnishing Status floor_number + -175857.84 * Furnishing Status Area Type_Carpet Area
+ -175173.56 * Furnishing Status Area Type_Super Area + 1054.59 * Furnishing Status City_Chennai + 1934.58 * Furnishing Status City_Delhi
+ -679.30 * Furnishing Status City_Hyderabad + 1298.05 * Furnishing Status City_Kolkata + 6327.87 * Furnishing Status City_Mumbai
+ -3317.53 * Bathroom^2 + 141.87 * Bathroom floor_number + -165789.41 * Bathroom Area Type_Carpet Area + -165857.30 * Bathroom Area Type_Super Area
+ 4721.69 * Bathroom City_Chennai + 22330.81 * Bathroom City_Delhi + 1478.58 * Bathroom City_Hyderabad + 5507.58 * Bathroom City_Kolkata
+ 4141.27 * Bathroom City_Mumbai + -14.38 * floor_number^2 + -28583.44 * floor_number Area Type_Carpet Area + -28223.98 * floor_number Area Type_Super Area
+ 768.83 * floor_number City_Chennai + -360.44 * floor_number City_Delhi + 1869.63 * floor_number City_Hyderabad + -113.98 * floor_number City_Kolkata
+ 1164.90 * floor_number City_Mumbai + -112617.06 * Area Type_Carpet Area^2 + 0.00 * Area Type_Carpet Area Area Type_Super Area
+ -109708.23 * Area Type_Carpet Area City_Chennai + -8605.85 * Area Type_Carpet Area City_Delhi + 7524.22 * Area Type_Carpet Area City_Hyderabad
+ -4002.08 * Area Type_Carpet Area City_Kolkata + 1563.71 * Area Type_Carpet Area City_Mumbai + -114163.93 * Area Type_Super Area^2
+ -100236.42 * Area Type_Super Area City_Chennai + 4408.78 * Area Type_Super Area City_Delhi + 23599.69 * Area Type_Super Area City_Hyderabad
+ 4511.47 * Area Type_Super Area City_Kolkata + -9910.94 * Area Type_Super Area City_Mumbai + 52728.53 * City_Chennai^2 + 0.00 * City_Chennai City_Delhi
+ 0.00 * City_Chennai City_Hyderabad + 0.00 * City_Chennai City_Kolkata + 0.00 * City_Chennai City_Mumbai + -4197.07 * City_Delhi^2 + 0.00 * City_Delhi City_Hyderabad
+ 0.00 * City_Delhi City_Kolkata + -4768.27 * City_Hyderabad^2 + 0.00 * City_Hyderabad City_Kolkata + 0.00 * City_Hyderabad City_Mumbai + 509.39 * City_Kolkata^2
+ 0.00 * City_Kolkata City_Mumbai + -8347.24 * City_Mumbai^2
```

# Interpretation of the Regression Equation

Rent is influenced by multiple factors, including:

- **Size:** Larger properties generally have higher rents.
- **BHK:** Unexpected large negative impact suggests possible data issues.
- **Furnishing Status:** Furnished properties command higher rents.
- **Bathrooms:** More bathrooms increase rent.
- **Floor Number:** Higher floors slightly increase rent.



# Regression model methods:

Multiple  
Linear  
Regression

Polynomial  
Regression

# Interpretation of R<sub>2</sub> and adjusted R<sub>2</sub> (using Linear & Polynomial Regression):

- For Multiple Linear Regression:
  - **R<sub>2</sub>:** 0.5228843329265807
  - **Adjusted R<sub>2</sub>:** 0.5167739935403682

For Polynomial Regression  
(degree=2):

**R<sub>2</sub>:** 0.8137733813128744

**Adjusted R<sub>2</sub>:** 0.7942618613107308

# Interpretation of R<sup>2</sup> & R<sup>2</sup> adjusted values:

- **Multiple Linear Regression:**

- **R<sup>2</sup> = 0.52** → The model explains 52.29% of the variance in rent prices can be explained by the independent features.
- **Adjusted R<sup>2</sup>** → After accounting for the number of predictors, the model still explains **51.7%** of the variance, confirming that adding more features did not drastically improve the fit.

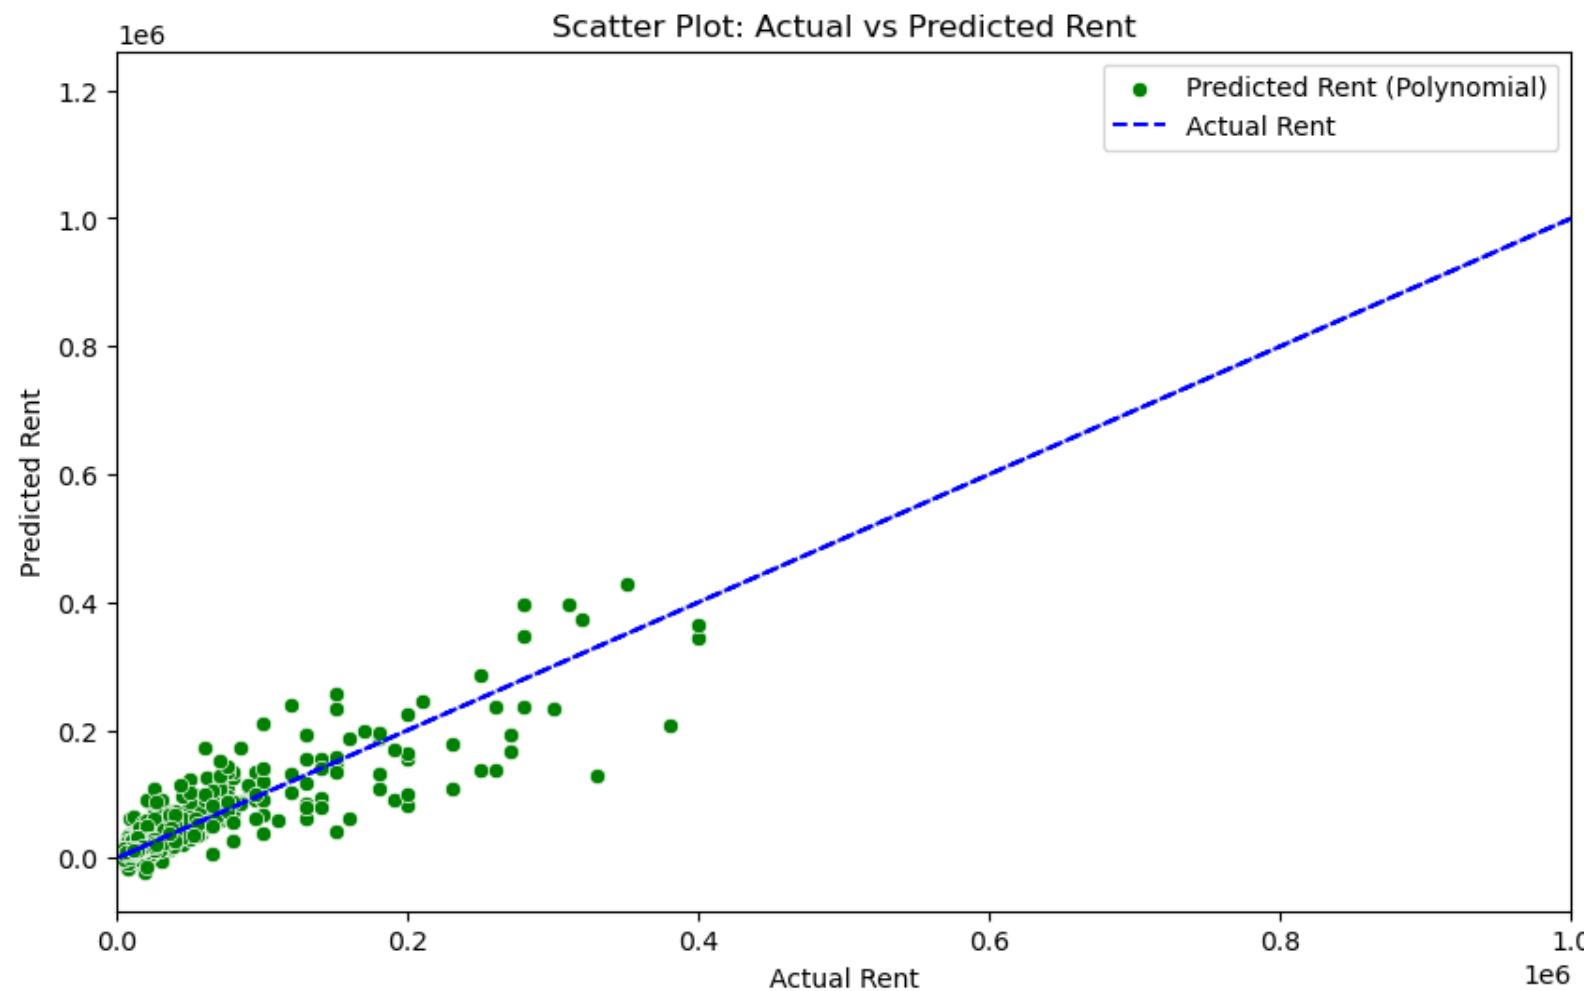
- **Polynomial Regression:**

- **R<sup>2</sup> = 0.81** → The model explains 81.38% of the variance in rent prices can be explained bu the independent variable , which is significantly higher than the multiple linear regression model.
- **Adjusted R<sup>2</sup> = 0.51** →After adjusting for additional predictors, the model still explains **79.4%** of the variance.

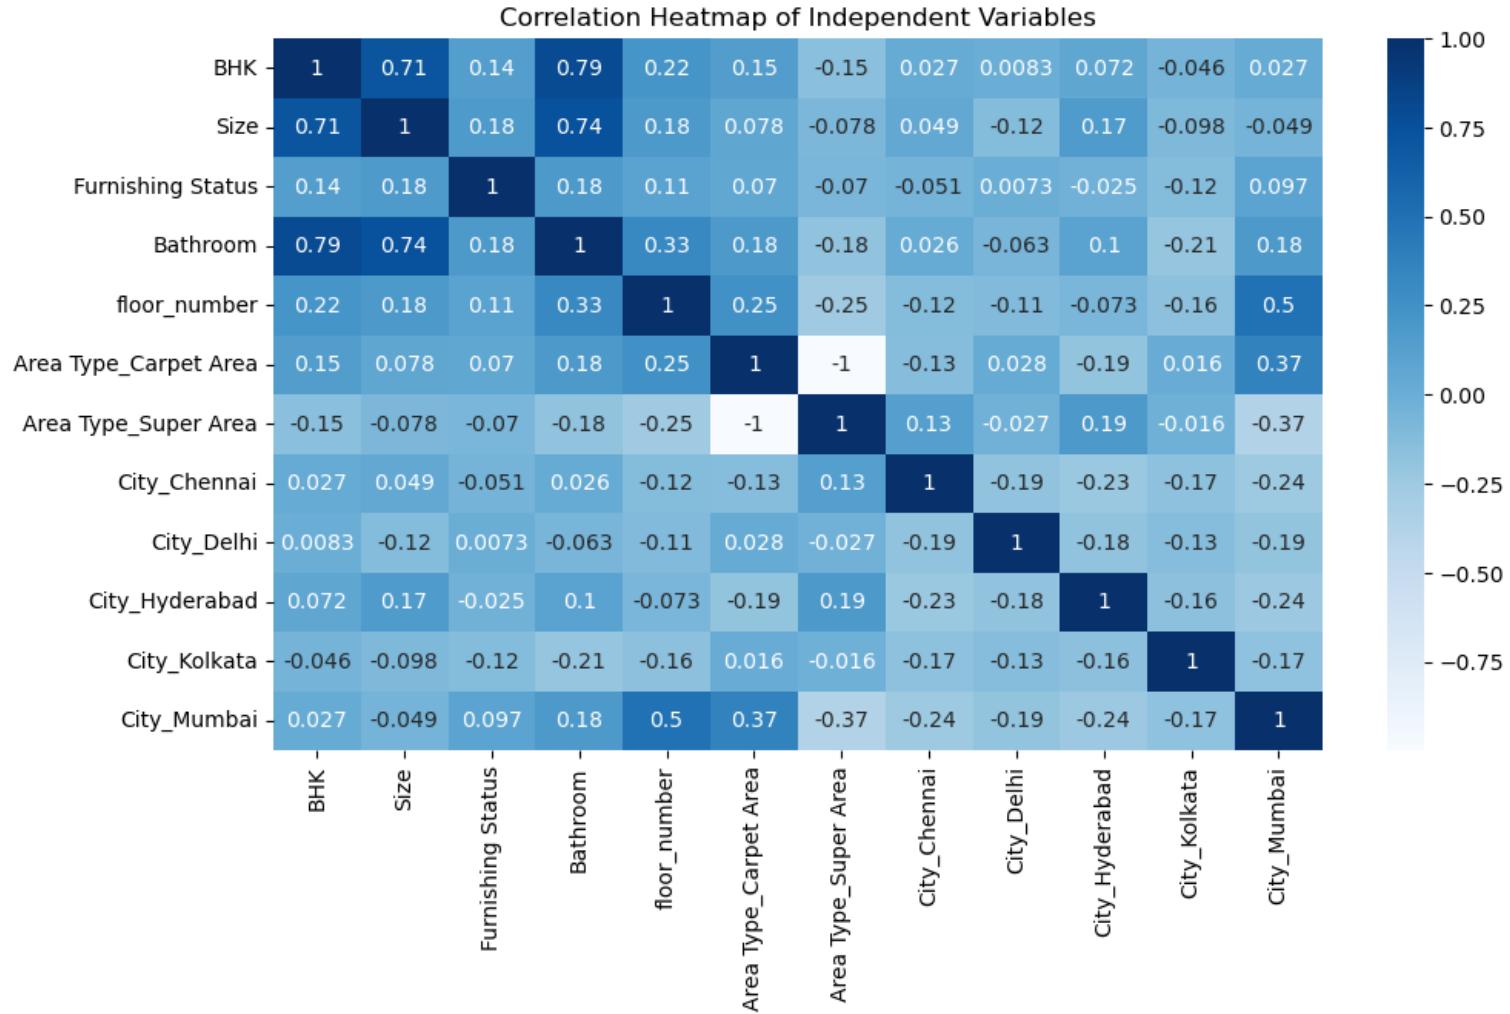
## INTERPRETATION:

**Polynomial regression outperforms** multiple linear regression in explaining the data, indicating that a nonlinear model is more suitable for our dataset.

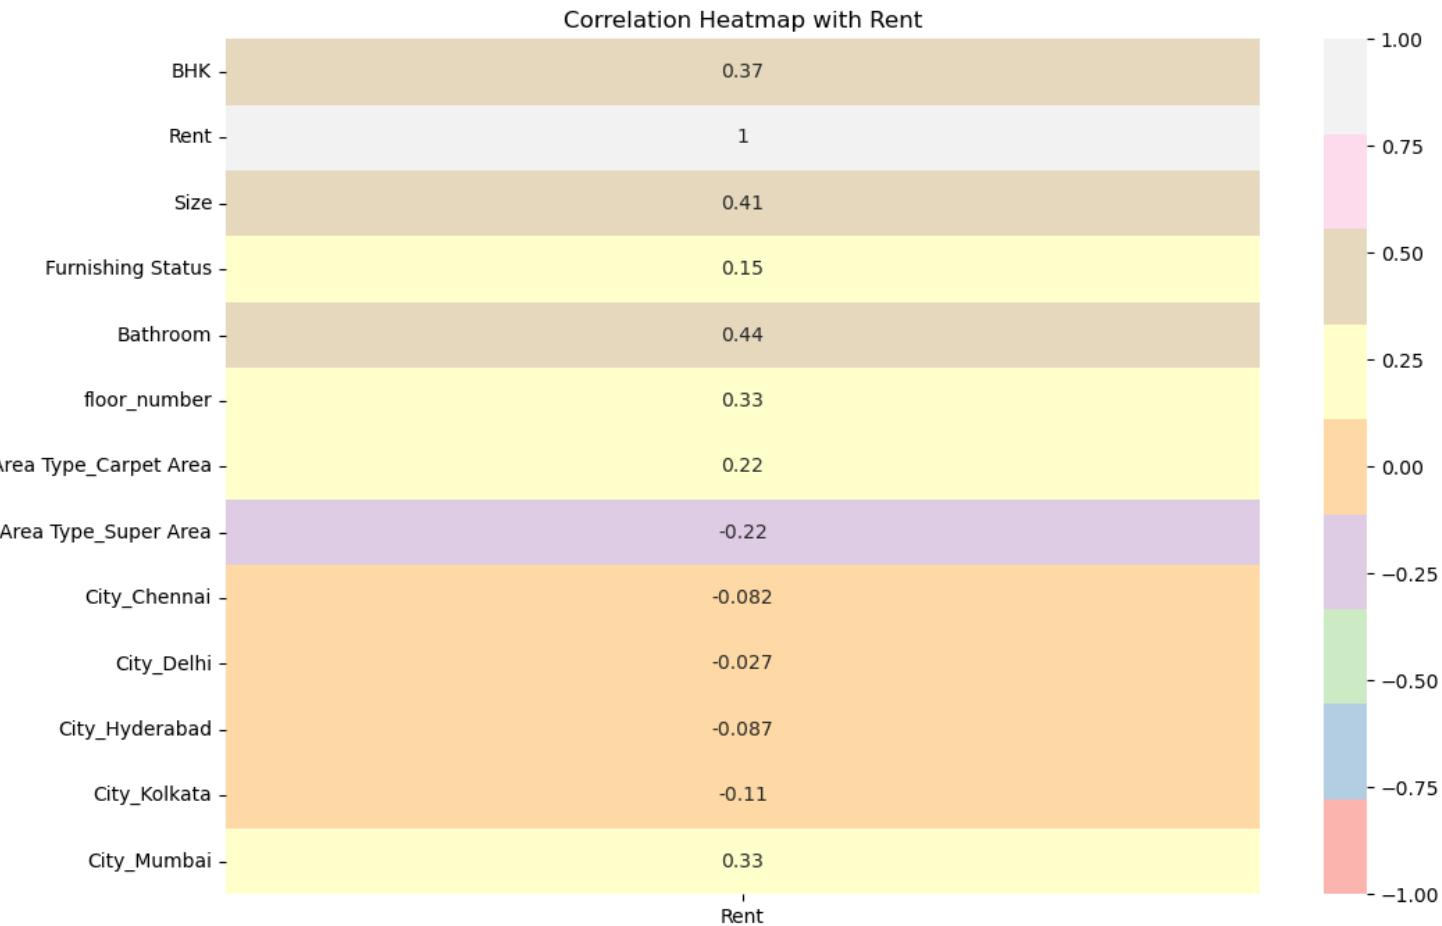
# Actual vs Predicted Rent:



# Correlation between Input variables:

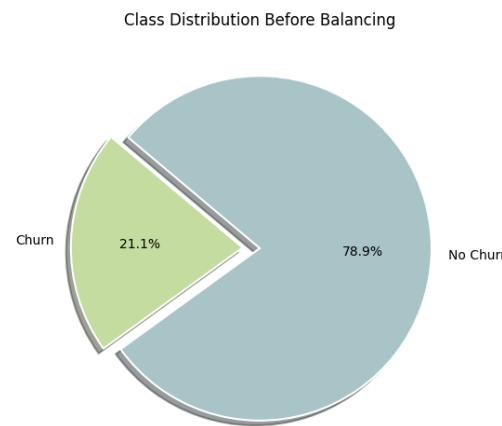
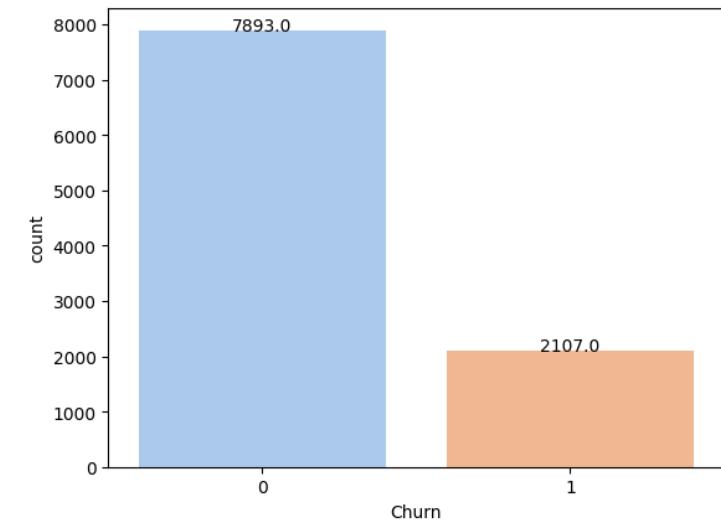


# Correlation between Input & Output variable:

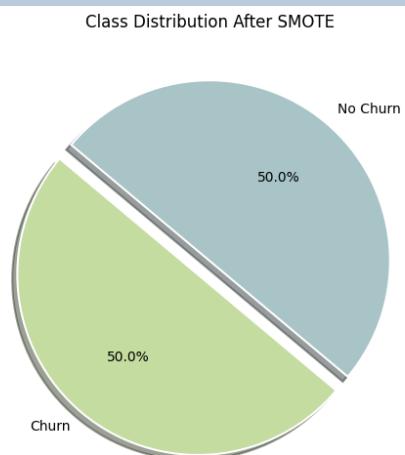


# Visualization

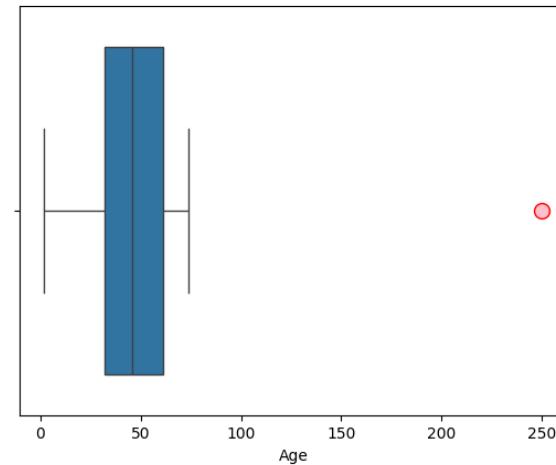
# Visualizations for Binary Classification Analysis:



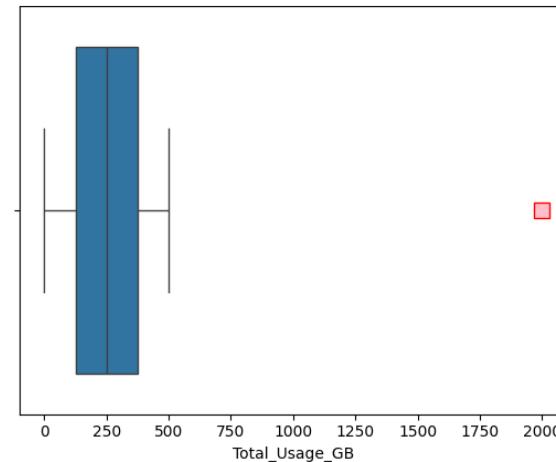
Class Distribution Before  
Balancing



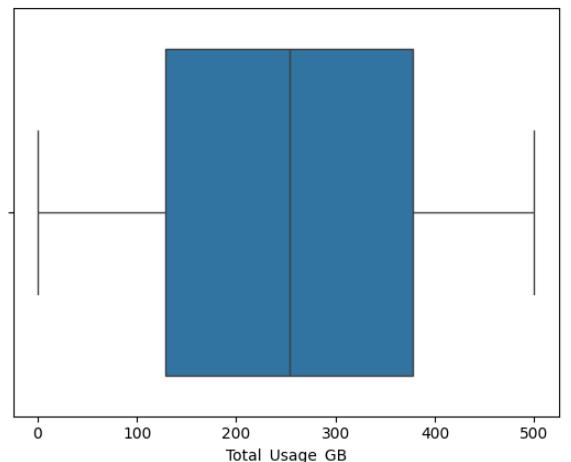
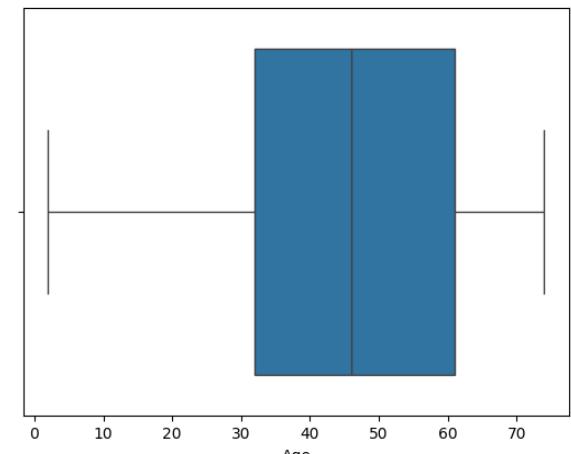
Class Distribution After  
SMOTE



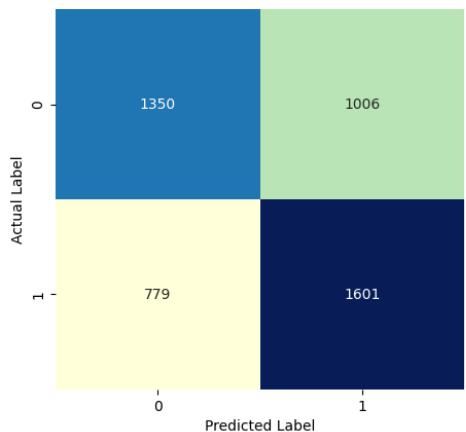
Outliers detected and removed for "Age" column



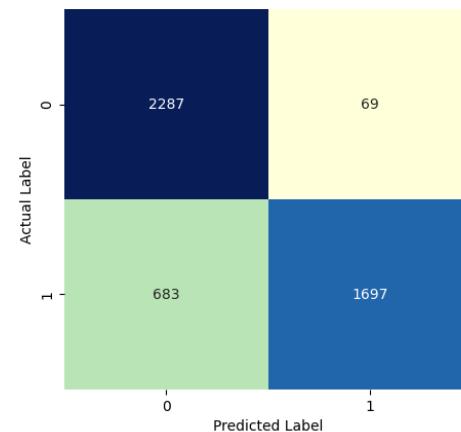
Outliers detected and removed for "Monthly Bill" column



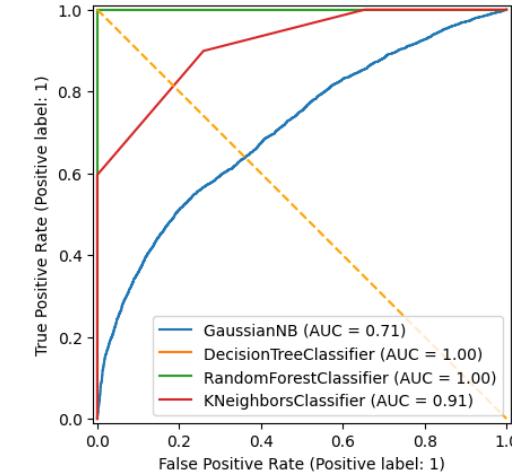
# Visualizations for Binary Classification Analysis:



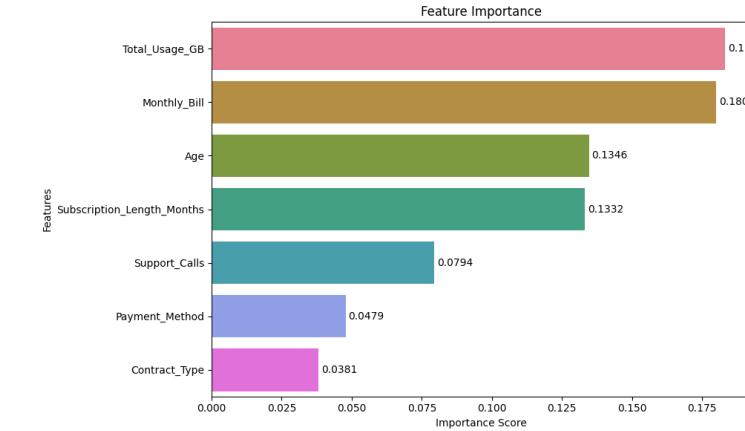
confusion matrix for Naive Bayes



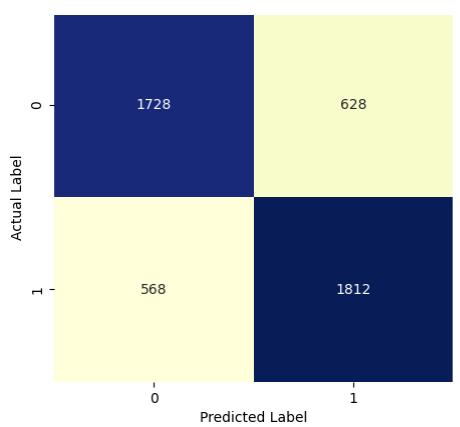
confusion matrix for Random Forest



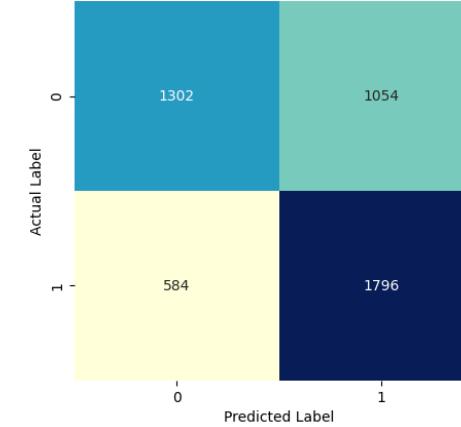
ROC Curve for Training dataset



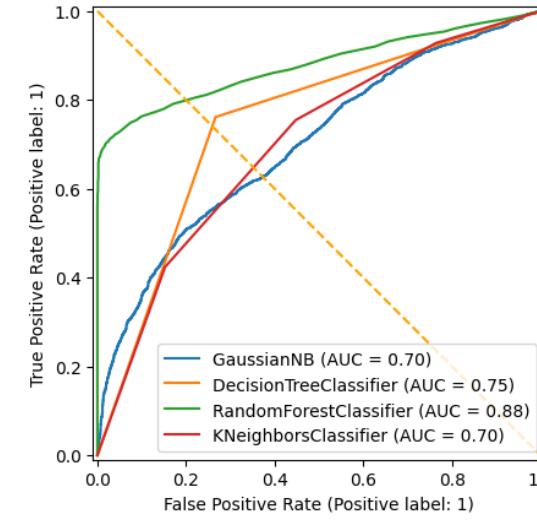
Features importance using Random Forest



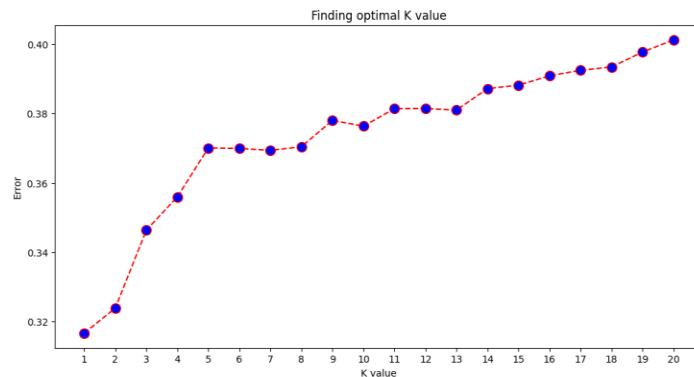
confusion matrix for Decision Tree



confusion matrix for KNN

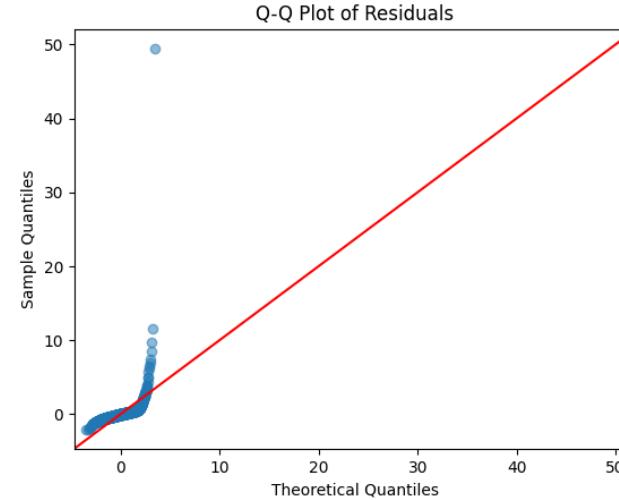
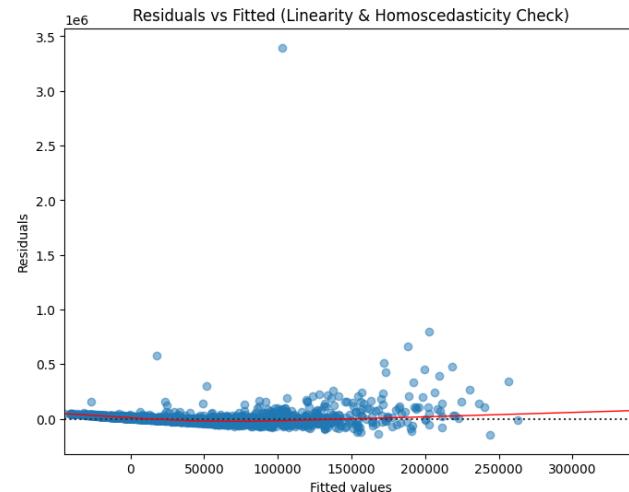


ROC Curve for Testing dataset

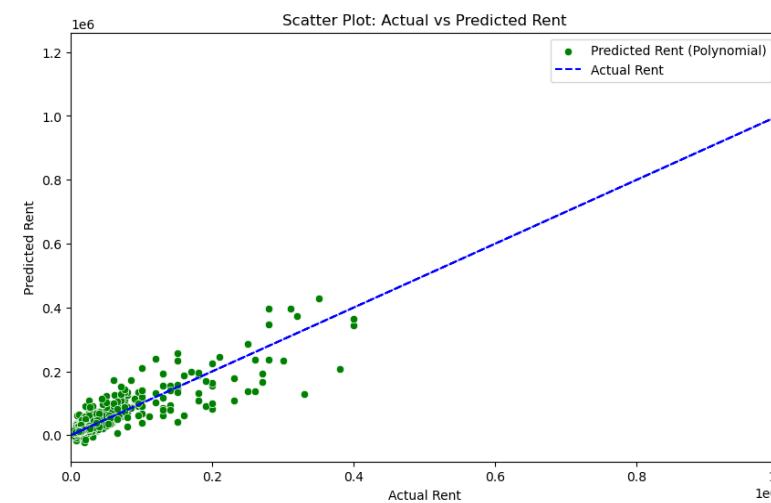


Finding best value of K in KNN

# Visualizations for Regression Analysis:

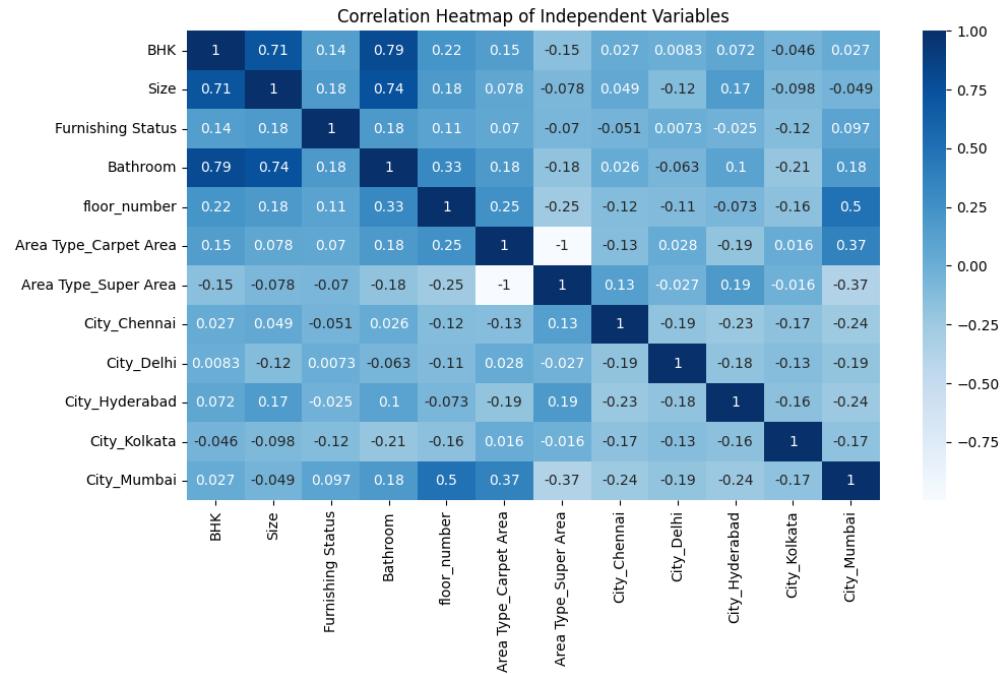
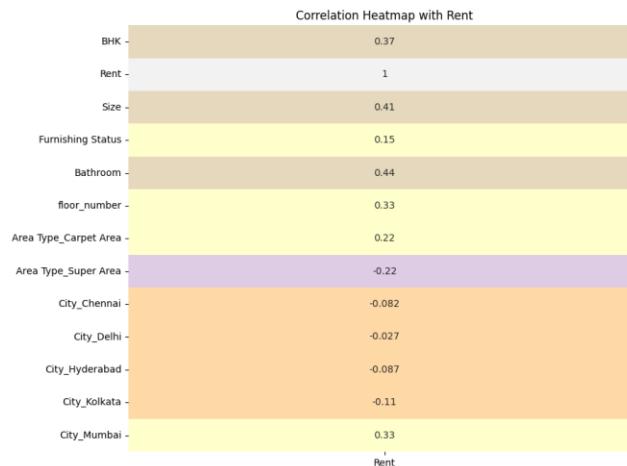


Residuals vs Fitted (Homoscedasticity)



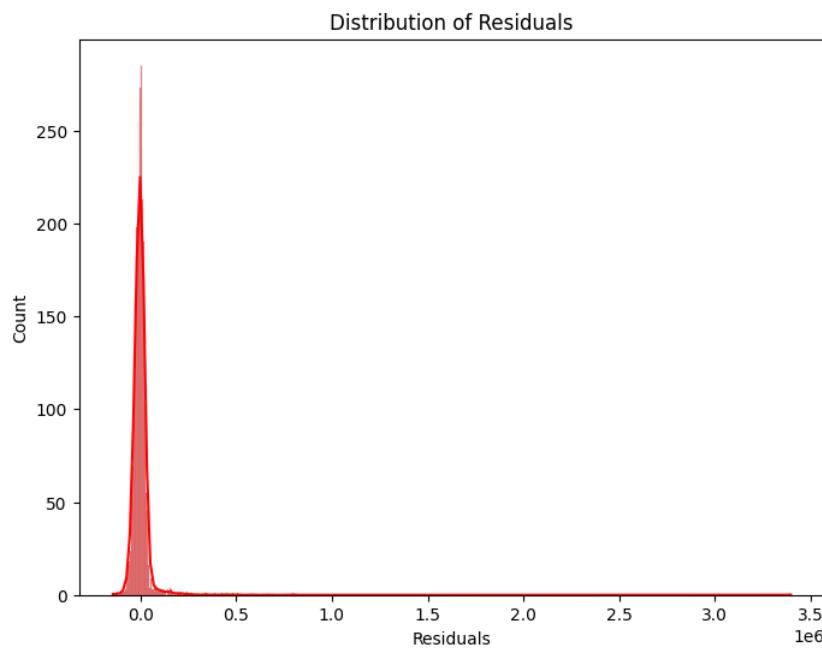
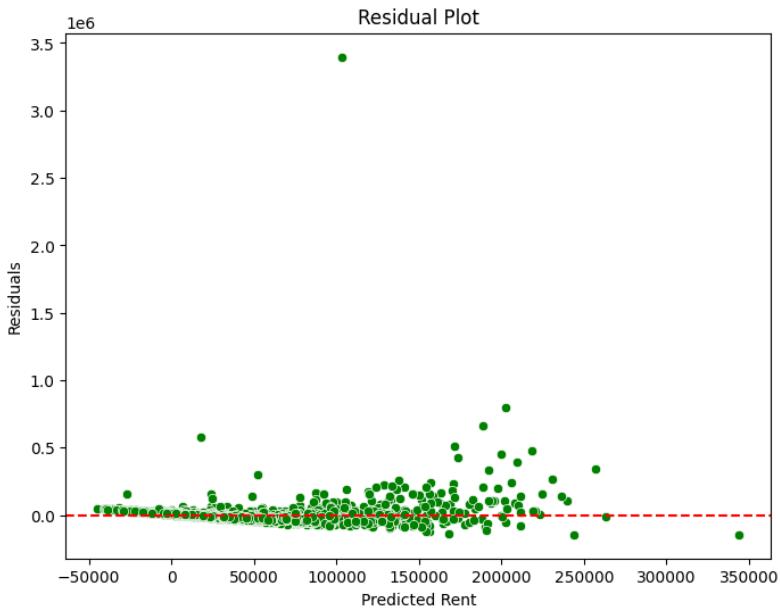
Scatter Plot: Actual vs Predicted Rent

Correlation Heatmap with Rent



Correlation Heatmap of Independent Variables

# Visualizations for Regression Analysis:



Residual Plot

Distribution of Residuals

Scatter Plot of Rent vs Size

# Thank you!

